RIGOROUS INTEGER FACTORISATION IN EXPECTED TIME $L_N[1/2, c]$ FOR EVERY $c > 1/\sqrt{2}$

Abstract. We describe a Las Vegas algorithm which, given any integer $N \geq 2$, outputs the complete prime factorisation of N, and we give a proof that for every $c > 1/\sqrt{2}$ its parameters can be fixed so that its expected running time is $O(L_N[1/2, c])$, where $L_N[1/2, c] = \exp\big(c\sqrt{\log N \log\log N}\big)$. The proof uses no unproved hypothesis. The previous best rigorous bound for factoring every integer is $L_N[1/2, 1 + o(1)]$, due to Lenstra and Pomerance (1992).

For each fixed choice of parameters (an odd degree d, a number k of large primes, a base exponent μ and a slack η_0) the expected time is $L_N[1/2, (1+\theta)/\sqrt{2}+o(1)]$, where $\theta = \theta(d, k, \mu, \eta_0)$ is an explicit exponent. For $d = 9$ and $k = 13$ this is below 1 ($c \approx 0.9972$). The infimum over fixed parameters at $d = 9$ is $1.3/\sqrt{2} \approx 0.9192$, and over all d it is $1/\sqrt{2}$. These are infima: we do not prove that a single algorithm of fixed degree attains $(1 + 3/(d + 1))/\sqrt{2} + o(1)$.

The algorithm combines four ingredients. The first is a "Dixon descent": a uniformly random square $r^2 \bmod N$ is written as a ratio u/v of two integers of size $\sqrt{N}$, and each prime of uv is descended by a special-q relation in a randomised number field sieve of fixed degree. Because the root r is never used, every dependency is fruitful with probability $1 - 2^{1-\omega(N)}$. Second, the halves u, v are completely factored with the Lenstra–Pomerance algorithm at cost $L_{\sqrt{N}}[1/2, 1] = L_N[1/2, 1/\sqrt{2}]$, rather than being combined. Third, the supply of relations is reduced to a single averaged estimate, which we prove by the Paley–Zygmund inequality from a first moment (a restricted Bombieri–Vinogradov theorem for the moduli $\ell \cdot$ (smooth), in the spirit of Lee–Venkatesan) and a second moment. The second moment is made tractable by randomising the rational side in two dimensions and by a random prime multiplier w of the resultant, $F(g_0, g_1) = Nw$, which spreads a forced residue class. It is then proved with divisor switching and the large sieve. Fourth, the number field sieve post-processing of Buhler–Lenstra–Pomerance and Lee–Venkatesan is carried out rigorously at fixed degree, and batch smoothness detection removes any dependence on the hyperelliptic curve method.

As further consequences of the same proof (Sections 14 and 15) we show the following.

- Replacing the Lenstra–Pomerance subroutine by a recursive call of the algorithm itself gives $L_N[1/2, \sqrt{(1 + \sqrt{2})\theta} + o(1)]$. The infimum is 0.851 at $d = 9$, and it tends to 0 as d grows.
- A restart schedule over such algorithms yields a single Las Vegas algorithm with expected time $L_N[1/2, o(1)]$; no effective rate is supplied. This is not $L_N[\alpha]$ for any $\alpha < 1/2$.
- The same relation supply computes discrete logarithms in $\mathbb{F}_p^*$ in expected time $L_p[1/2, (1+\theta)/\sqrt{2} + o(1)]$, below Pomerance's $L_p[1/2, \sqrt{2}]$.

PRIVATE DRAFT for expert review.[1] Do not distribute. The author field is intentionally blank. All inputs are published, numbered statements. One of them, Drappeau's Lemme 5 (used in Appendix A), has a published proof that refers to computations in an unpublished preprint of Harper; see the discussion at the start of Appendix A.

Contents

1. Introduction 3
1.1. Background 3
1.2. The result 4
1.3. Further consequences: the recursive shape test and discrete logarithms 5
1.4. Comparison with Lenstra–Pomerance 5

Date: PRIVATE DRAFT for expert review — not for distribution. Version of 20 October 2026.

1.5. The idea in one page 6
1.6. Prior art and attributions 7
1.7. Organisation 8
2. Notation and published inputs 8
2.1. Notation 8
2.2. Smooth numbers 8
2.3. Other published inputs 9
3. Parameters, field law and the algorithm 9
3.1. Parameters 9
3.2. The field law 10
3.3. Candidates and the counted relations 11
3.4. The algorithm 12
4. Overview of the proof 14
5. Split counting 15
6. Fruitfulness: Dixon descent 16
7. Bounded-k splits and field selection 17
8. The first moment 19
8.1. One-point smoothing 19
8.2. Good moduli 19
8.3. Bad moduli are rare 20
8.4. Rational fibres 21
8.5. The first moment 21
9. The second moment and relation supply 22
9.1. Pairs in one lattice 22
9.2. The multiplier 23
9.3. Divisor switching 23
9.4. An all-residue second moment for the family $\ell^2 nq$ 24
9.5. The second moment 25
9.6. Relation supply 27
9.7. A rational pair correlation (not used in the proof) 27
10. Post-processing at fixed degree 27
10.1. Field-level events 27
10.2. Rows and the row identity 28
10.3. Columns 29
10.4. The residual square-class group 29
10.5. Characters 30
10.6. Square roots 31
10.7. Linear algebra and the congruence 32
11. Batch smoothness 33
12. Cost assembly and the constant 34
13. Uniformity, small factors and prime powers; proof of the main theorem 34
14. The recursive shape test 35
14.1. The algorithm and what is used from the main proof 35
14.2. Solving the recursion 36
14.3. A portfolio: $L_N[1/2, o(1)]$ and a genuine $+o(1)$ 38
15. Application: discrete logarithms in $\mathbb{F}_p^*$ 38
15.1. Statement and literature 39
15.2. The algorithm and the transfer of the relation supply 39
15.3. Exact dependencies 40

16. A toy implementation 43
17. Limitations and outlook 44
17.1. Limitations 44
17.2. Outlook: growing degree 45
Acknowledgements 45
Appendix A. Proof of Lemma 8.4 46
A.1. Setting and statement 46
A.2. Cited results 47
A.3. Elementary lemmas 48
A.4. Large conductors 49
A.5. Smooth moduli with a relative saving 50
A.6. Proof of Lemma A.1 50
References 51

1. Introduction

1.1. **Background.** Write $L_N[\alpha, c] = \exp\bigl(c(\log N)^{\alpha}(\log\log N)^{1-\alpha}\bigr)$ and $L = L_N[1/2, 1]$. The fastest factoring algorithms in practice, the quadratic and number field sieves, run in heuristic time $L_N[1/2, 1 + o(1)]$ and $L_N[1/3, (64/9)^{1/3} + o(1)]$ [53, 41, 10]. Their running times are not proved: they rest on assumptions that certain polynomial values behave like random integers with respect to smoothness, and that the resulting congruences of squares are not all trivial.

Rigorous analysis has a separate history.

- Dixon's random-squares method [19] was the first factoring algorithm with a proved subexponential running time, $L_N[1/2, O(1)]$. Its correctness argument is simple and is the starting point of this paper. Draw r uniformly, look for relations $r^2 \equiv \prod p^{e_p} \pmod N$ over a factor base, and combine them into $X^2 \equiv Y^2$. Since Y is computed without looking at the roots r, the class of X/Y in the group of square roots of unity is uniformly distributed, and $\gcd(X - Y, N)$ is a proper factor with probability at least $1/2$. This is the same fibre symmetry that underlies Rabin's reduction from computing square roots to factoring [56].
- Pomerance analysed and refined these methods [53], and proved $L_N[1/2, \sqrt{2} + o(1)]$ for random squares with a rigorous elliptic curve smoothness test [54]. Vallée's sampler of small squares gave $L_N[1/2, \sqrt{4/3} + o(1)]$ [65].
- Class-group methods reached exponent 1. Schnorr and Lenstra [59] did so under heuristic hypotheses. Seysen [60] gave a class-group algorithm rigorous under the generalised Riemann hypothesis, and A. K. Lenstra [40] brought the GRH-conditional constant to 1. Lenstra and Pomerance [44] removed every hypothesis. They proved that the complete prime factorisation of every N can be found in expected time $L_N[1/2, 1 + o(1)]$. That bound has stood since 1992. It is still described as the best rigorous probabilistic bound, for instance by Pomerance [55] and Harvey [29].
- Lee and Venkatesan [39] analysed rigorously an explicitly randomised number field sieve. They proved that it finds congruences of squares in expected time $L_N[1/3, (64/9)^{1/3} + o(1)]$. Whether these congruences yield a factor is left to a conjecture (their Conjecture 7.1), and their factoring conclusion is stated for $N = pq$ with $p \equiv q \equiv 3 \pmod 4$. Their techniques for the first moment of smooth values (a randomisation that makes algebraic values nearly uniform on long progressions, and a restricted Bombieri–Vinogradov estimate for smooth moduli built on Harper's work [27]) and for the post-processing

(random characters, Stark–Landau) are the models for ours. We use their randomisation idea and the Stark and Landau statements as quoted there, but none of their lemmas on smooth numbers.

There is also a folklore heuristic that $L[1/2, 1]$ is a natural barrier for methods that combine random-like auxiliary integers of size at least $N^{1/2}$ into a square. Croot, Granville, Pemantle and Tetali [15] and Balister, Bollobás and Morris [2] made it sharp for random integers. We found no theorem showing that $c = 1$ is a barrier for rigorous factoring itself.

1.2. **The result.** All our parameters are fixed: an odd degree $d \geq 3$, an integer $k > (d+1)/4$, a rational base exponent μ and a rational slack $\eta_0 > 0$. Put $\nu = 1/(4k)$ and

$$h_*(\mu) = \max\bigl(2\nu,\ \mu-(d-2)\nu,\ 1-d\mu+\nu\bigr), \qquad \theta = \theta(d,k,\mu,\eta_0) = 2\mu+h_*(\mu)+\eta_0+(d-1)\nu. \tag{1}$$

The parameters are *admissible* if $\nu + \eta_0 \leq \mu < 1$ and $\theta < 1$.

Theorem 1.1 (main theorem, fixed parameters)**.** *For every admissible (d, k, μ, η_0) there is a Las Vegas algorithm which, given any integer $N \geq 2$, outputs the complete prime factorisation of N. Its expected running time is at most*

$$L_N\bigl[1/2,\ (1+\theta)/\sqrt{2}+o(1)\bigr] \qquad (N \to \infty),$$

where the $o(1)$ depends only on (d, k, μ, η_0).

Minimising θ over μ (Lemma 3.2) gives, at $\eta_0 \downarrow 0$,

$$\theta_{\min}(d,k) = \begin{cases} \dfrac{2}{d} + \dfrac{d^2+d-2}{d}\,\nu, & k < (d^2+1)/4, \\[2ex] \dfrac{3}{d+1} + \dfrac{4d-2}{d+1}\,\nu, & k \geq (d^2+1)/4, \end{cases} \qquad \inf_{k,\eta_0} \frac{1+\theta}{\sqrt{2}} = \frac{1+3/(d+1)}{\sqrt{2}}.$$

Letting d grow gives the uniform statement.

Corollary 1.2. *For every $c > 1/\sqrt{2}$ there is a Las Vegas algorithm which, given any integer $N \geq 2$, outputs its complete prime factorisation in expected time $O(L_N[1/2, c])$.*

Proof. Choose d with $(1+3/(d+1))/\sqrt{2} < c$, then k large and η_0 small with $(1+\theta)/\sqrt{2} < c$, and apply Theorem 1.1. Finitely many N cost $O(1)$. □

Table 1 lists exact values. The constant is below 1 if and only if $\theta < \sqrt{2} - 1 = 0.41421\ldots$.

Remark 1.3 (what is and is not claimed)**.**

(a) *Infimum, not $o(1)$.* At fixed d the value $(1+3/(d+1))/\sqrt{2}$ is an infimum over fixed (k, η_0). It is not attained, and we do not claim $L_N[1/2, (1+3/(d+1))/\sqrt{2}+o(1)]$ for any fixed d. That would require letting $k = k(N) \to \infty$. Our constants depend on k through factors $(\log N)^{O(k)}$, $\binom{2k}{k}$ and thresholds $N_0(k)$, and we have not carried out such a diagonalisation. At $d = 9$ the first parameter choice below 1 ($k = 13$) has $c \approx 0.9972 + \eta_0/\sqrt{2}$; larger fixed k give smaller constants (Table 1), but none attains 0.9192. (Corollary 14.8 obtains a genuine $+o(1)$ by scheduling instead.)

(b) *Complete factorisation.* The algorithm outputs the full factorisation, with every prime certified (AKS [1]). Prime powers, perfect powers and small prime factors are handled before the main splitting routine. The splitting routine is applied at most $\log_2 N$ times (Section 13).

(c) *No hypotheses, no exceptional set.* Nothing is assumed beyond published theorems. In particular no Riemann hypothesis, no hyperelliptic curve method (whose published proof trail is incomplete), and no exceptional set of N. Every input is a published, numbered statement. One of them, Drappeau's Lemme 5, has a published proof that points to computations in Harper's unpublished preprint (Appendix A). The threshold beyond

d	least k	θ there ($\eta_0 \downarrow 0$)	c there	need $\eta_0 \lesssim$	c at $k = 50$	$k = 1000$	$\inf_{k,\eta_0} c$
7	21	139/336	0.99963	0.00052	0.98376	0.97285	0.97227
9	13	16/39	0.99720	0.00396	0.93126	0.91984	0.91924
11	13	9/22	0.99638	0.00512	0.89626	0.88450	0.88388
13	14	73/182	0.99073	0.01311	0.87126	0.85926	0.85863
19	17	257/646	0.98842	0.01638	0.85188	0.81383	0.81317
$d \to \infty$							$1/\sqrt{2} = 0.70711$

TABLE 1. The constant $c = (1+\theta)/\sqrt{2}$ of Theorem 1.1 at the optimal μ. "Least k" is the least k with $c < 1$ for all sufficiently small $\eta_0 > 0$; the exact condition is $\eta_0 < \sqrt{2} - 1 - \theta$ (rounded in the table, e.g. 0.0039571... at $d = 9$). At $d = 9$, $k = 13$: $\mu = 17/156$, $h_* = 1/26$, and $(1+16/39)^2 = 3025/1521 < 2 = 3042/1521$; at $k = 12$ the minimum is $\theta = 23/54 > \sqrt{2} - 1$. For $d \le 5$ the constant is never below 1 ($d = 5$: infimum 1.0607). Computed in exact rational arithmetic by `src/thmA/params.py`; independently re-derived in two referee rounds.

which the time bound holds involves Stark's, Landau's and Lagarias–Odlyzko's effective constants; it is enormous, and effective provided the constants of Drappeau's Lemme 5 are (which we believe but no published source states). The algorithm itself uses no analytic constant (Remark A.17), so Theorem 1.1 is unaffected either way. The bound holds for every N, not on average.

(d) *Status.* The proof is new and has not been peer reviewed. It is long and rests on several delicate interfaces, so it should be treated with caution until experts have checked it.

1.3. **Further consequences: the recursive shape test and discrete logarithms.** The subroutine that costs $L_N[1/2, 1/\sqrt{2}]$ can be replaced by the algorithm itself, called recursively on the two halves. Section 14 derives the following from the proof of Theorem 1.1.[2]

- *Theorem 14.6.* At fixed parameters the expected time is $L_N[1/2, \sqrt{(1+\sqrt{2})\theta} + o(1)]$. The infimum over (k, η_0) is $\sqrt{3(1+\sqrt{2})/(d+1)}$, which is 0.851 at $d = 9$ and 0.777 at $d = 11$, and tends to 0 as $d \to \infty$. These are infima; for example $d = 9$, $k = 13$ gives 0.9952.
- *Corollary 14.8.* A restart schedule over such algorithms gives a single algorithm with expected time $L_N[1/2, o(1)]$; no effective rate is supplied. The same device turns the infimum 0.9192 at $d = 9$ into a genuine $+o(1)$.

Neither gives a rigorous $L_N[\alpha]$ with $\alpha < 1/2$ (Section 17). Section 15 transfers the relation supply to discrete logarithms in $\mathbb{F}_p^*$ (Theorem 15.1): expected time $L_p[1/2, (1+\theta)/\sqrt{2} + o(1)]$ whenever $\theta \le 1/2$. This beats Pomerance's rigorous $L_p[1/2, \sqrt{2}]$ [54], which is the best bound we found for all primes.

1.4. **Comparison with Lenstra–Pomerance.**

[2]Theorem 14.6 and Corollary 14.8 are conditional on Theorem 1.1's proof. They were audited internally by three blind referees (one GPT-6.1-sol, two Claude), all of whom returned "accept with fixes" with no fatal or major finding. They are not peer reviewed.

	Lenstra–Pomerance 1992 [44]	this paper
output	complete factorisation, Las Vegas	complete factorisation, Las Vegas
hypotheses	none	none
expected time	$L_N[1/2, 1 + o(1)]$	$L_N[1/2, (1+\theta)/\sqrt{2} + o(1)]$, $\theta < 1$ fixed
best constant	1	any $c > 1/\sqrt{2}$ (Corollary 1.2)
relations	class-group relations with multipliers	Dixon descent in a randomised NFS
smoothness tests	ECM on "recognisable" smooth numbers	batch smooth parts (no ECM/HECM)
use of the other	—	LP factors the two $\sqrt{N}$-size halves

We use the Lenstra–Pomerance algorithm as a subroutine. It factors integers of size $\sqrt{2N}$ in time $L_{\sqrt{N}}[1/2, 1 + o(1)] = L_N[1/2, 1/\sqrt{2} + o(1)]$, and that is exactly where the constant $1/\sqrt{2}$ comes from. Any improvement of the rigorous constant for factoring $\sqrt{N}$-size integers would improve ours (Section 17).

1.5. **The idea in one page.**

Fruitfulness without characters. Take r uniform in $(\mathbb{Z}/N)^\times$ and $R = r^2$. Suppose we can express R multiplicatively over some factor base, in a way that depends on R and on independent coins but never on r. Then the exponent vectors of many such R can be combined by linear algebra, and the resulting congruence $\mathsf{X}^2 \equiv \mathsf{Y}^2$ has X/Y uniform in $\boldsymbol{\mu}_2 = \{\epsilon : \epsilon^2 = 1\} \cong \{\pm 1\}^{\omega(N)}$. So it is fruitful with probability exactly $1 - 2^{1-\omega(N)}$, whatever the factor base is (Theorem 6.1). Dixon's method expresses R itself as smooth, which is expensive. We *descend* it instead, as the individual-logarithm phase of number field sieve discrete-log algorithms does [25, 34, 14]. Write $R \equiv u/v \pmod N$ with $|u|, |v| \le \sqrt{2N}$, by Gauss reduction of a lattice of determinant N. Then express each prime $\ell \mid uv$ through a special-q relation of a number field sieve.

The shape test. We accept R only if $|u|$ and $|v|$ are products of at most k distinct primes, all in one dyadic interval J around $N^{1/(2k)}$. Deciding this requires complete factorisation of $\sqrt{N}$-size integers whose prime factors are $N^{\Theta(1)}$. We pay for it with the Lenstra–Pomerance algorithm, at $L_N[1/2, 1/\sqrt{2}]$ per trial. With k fixed, a "good" field (one in which a fraction $\delta \ge L^{-o(1)}$ of the primes of J can be descended) accepts a fraction $\ge (\delta/\log N)^{2k} 2^{-\omega(N)} = L^{-o(1)}$ of all R. This follows from a Dixon–Pomerance-type counting lemma in ratio form (Theorems 5.1 and 5.2), which holds for an arbitrary set of admissible primes.

One averaged estimate suffices. Theorem 7.2 shows that complete descents in one field follow from a single averaged statement: $\Pr_{A,\ell}[\ell \text{ descends in field } A] \ge L^{-o(1)}$, where A is a random field and ℓ a random prime of J. No joint event across the k targets of a split is needed, because each field only uses splits inside its own set of descendable primes.

The field and the moments. The field is $F_A = \hat{F} + G \cdot R_c$, a binary form of degree d, where:

- $G(x, y) = g_1 x - g_0 y$ with (g_0, g_1) random of size $M = N^\mu$;
- $\hat{F}$ is the base-(g_0, g_1) expansion of Nw for a random prime w slightly larger than MN^ν;
- R_c has random coefficients of size $H = N^h$.

A prime $\ell \in J$ descends if some primitive $P = (a, b)$ in the special-q lattice $\{\ell \mid G(P)\}$, in a box of side $V \approx N^\nu$, has $G(P)/\ell$ and $F_A(P)$ both B-smooth with $B = L^{\theta/\sqrt{2}}$. The two smooth values have sizes $N^{\mu-\nu}$ and $N^{\mu+h+d\nu}$, with total exponent θ. The Paley–Zygmund inequality reduces the averaged statement to the following two moment bounds.

- A *first moment* (Theorem 8.7). The algebraic values lie in a progression to the modulus $|G(P)|$, which is a multiple of ℓ. The randomisation of R_c makes them nearly uniform on long progressions, as in Lee–Venkatesan. A restricted Bombieri–Vinogradov estimate for the moduli $\ell \cdot (\text{smooth})$ (Lemma 8.4) shows that bad moduli are rare.

- A *second moment* (Theorem 9.7). Pairs of candidates in one lattice are coupled through a kernel lattice of index $|\det(P, P')|$ (Lemma 9.2). The forced residue $F_A(P) \equiv Nw\, b^d g_1^{-d} \pmod{G(P)}$ would couple them through the rational data. The random prime w spreads it (Lemma 9.3, Brun–Titchmarsh). What is left is an all-residue second moment of smooth numbers to the moduli $\ell^2 nq$. We prove it at the precision we need by divisor switching and the large sieve, averaged over ℓ (Theorem 9.5).

Cost. About B rows are needed. Each costs $L^{1/\sqrt{2}}$ for the shape test and $S^2 \approx L^{\theta/(2a)}$ for the special-q searches, where $B = L^a$. Linear algebra costs B^2. The total is $L^{\max(2a,\, a+1/\sqrt{2},\, a+\theta/(2a))}$, minimised at $a = \theta/\sqrt{2}$, which gives $(1+\theta)/\sqrt{2}$. The folklore $L[1/2, 1]$ model does not bind. The two $\sqrt{N}$-size halves are factored, not combined. The relations that enter the linear algebra are descent norms of total size N^θ with $\theta < \sqrt{2} - 1 < 1/2$.

1.6. **Prior art and attributions.** Most components are known. We list them with the precise debt; the combination is what we did not find.

- *Fibre symmetry and fruitfulness.* Rabin [56] and Dixon [19], with Pomerance's analyses [53, 54]. Theorem 6.1 is this argument applied to a descent; we regard it as folklore-level.
- *Descent.* Rational reconstruction of a target as u/v, followed by special-q descent of its large primes, is the initial step of number field sieve individual-logarithm algorithms [25, 34, 14]. These analyses are heuristic. Stange [61] studies index calculus modulo a composite n with random exponents and multiplicative relations, with an NFS variant, under her Hypothesis 3.1. She is the closest conceptual precedent for using hidden random elements as targets.
- *Split counting.* Theorems 5.1 and 5.2 adapt the Dixon–Pomerance random-squares counting [19, 53] to the ratio form $u \equiv Rv$.
- *Non-monic linear side and post-processing.* The linear form $g_1 x - g_0 y$ with a base-(g_0, g_1) expansion is the base-(m_1, m_2) variant of Buhler, Lenstra and Pomerance [10, §§3,12]; see also Murphy [49, §2.3.2] and Kleinjung [36]. Our post-processing follows BLP's obstruction-group analysis, their proposal of random quadratic characters, their p-adic square root and their parity coordinate for the leading coefficient [10, §§6,8,9,12], the last after Weinberger–Rothschild [69]. It also follows Lee–Venkatesan's rigorous treatment of characters through Stark and Landau [39, §6]. What we add is the fixed-degree specialisation with explicit bounds (Section 10).
- *Randomised NFS and the first moment.* The randomisation $f = \hat{f} + (x - m)R$ and the use of a restricted Bombieri–Vinogradov estimate over smooth moduli are due to Lee–Venkatesan [39, §§5,8], building on Harper's Bombieri–Vinogradov and Barban–Davenport–Halberstam theorems for smooth numbers [27]. Our Lemma 8.4 is the analogue for the moduli $\ell \cdot$(smooth). Its proof (Appendix A) is self-contained apart from one published lemma of Drappeau [21, Lemme 5] (also Pascadi [51, Lemma 3.2]) for characters of small conductor, and the classical large sieve [33, Thm. 7.13]. Our one-point smoothing Lemma 8.1 replaces Lee–Venkatesan's Lemmas 5.13–5.20.
- *Second moment.* Lemma 9.4 is a progression form of the standard flexible factorisation of smooth numbers [23, 30] (cf. [28, Smooth Numbers Result 3]). Theorem 9.5 follows the classical large-sieve route to Barban–Davenport–Halberstam upper bounds [24, 45], with a Hooley-type rearrangement [32]. We found no statement of it.
- *The multiplier.* Multipliers are an old device. Knuth–Schroeppel multipliers kN appear in the continued fraction and quadratic sieve methods [48]; Schnorr–Lenstra [59] and Lenstra–Pomerance [44] use multipliers of the discriminant (LP: "the main device ... is the use of multipliers"); and kN is used in NFS polynomial selection [49]. Our use

differs. The multiplier w is a large random prime of the resultant, of the size of the rational modulus, chosen to spread a forced residue in a rigorous second moment. We found no earlier instance of this.

- *Batch smoothness.* Bernstein's smooth-parts algorithm [5, 6]. Bernstein proposed it as a replacement for ECM in rigorous bounds [4]. We apply it to complete special-q boxes (Section 11).
- *Shape test.* Lenstra–Pomerance [44], used as a black box.

We make no novelty claim beyond saying that we did not find this combination, or a rigorous $L_N[1/2, c]$ with $c < 1$ for factoring every integer (unconditional, or conditional on the generalised Riemann hypothesis), in the literature.

1.7. **Organisation.** Section 2 fixes notation and lists the published results we use, with their exact ranges. Section 3 gives the parameters, the field law and the algorithm. Section 4 outlines the proof and gives a dependency diagram. The proof proper is in Sections 5–13, and Appendix A proves Lemma 8.4. Section 14 treats the recursive shape test, and Section 15 discrete logarithms. Section 16 describes a toy implementation, and Section 17 discusses limitations and the growing-degree outlook.

2. Notation and published inputs

2.1. **Notation.** Throughout, log is the natural logarithm and

$$\mathcal{L} = \mathcal{L}_N := \sqrt{\log N \log\log N}, \qquad L = L_N[1/2, 1] = e^{\mathcal{L}}, \qquad L_N[1/2, c] = L^c.$$

We write $L^{o(1)}$ for any quantity of the form $\exp(o(\mathcal{L}))$, with the $o(1)$ uniform in N. Fixed powers L^{ε} and quantities $N^{o(1)}$ are *not* $L^{o(1)}$; when they occur, they are always absorbed by a fixed power of N. Constants $C, c, C_d, \dots$ may depend on the fixed parameters (d, k, μ, η_0) and change from line to line. $(\log N)^{O(1)}$ denotes a fixed power of $\log N$.

$\omega(n)$ is the number of distinct prime factors of n, $\tau(n)$ the number of divisors, φ Euler's function and $P^+(n)$ the largest prime factor of n ($P^+(1) = 1$). $\pi(x)$ counts primes. For the smoothness bound B fixed in Section 3.1 we write

$$\Psi(x) = \#\{1 \le n \le x : P^+(n) \le B\}, \qquad \varrho(x) = \Psi(x)/x,$$
$$\Psi(x; t, r) = \#\{1 \le n \le x : P^+(n) \le B,\ n \equiv r\ (t)\},$$
$$\Psi_t(x) = \#\{1 \le n \le x : P^+(n) \le B,\ (n, t) = 1\},$$

and $\Psi(x; \chi) = \sum_{n \le x,\, P^+(n) \le B} \chi(n)$ for a Dirichlet character χ. When we need another smoothness bound y we write $\Psi(x, y)$, and so on. The set of nonzero B-smooth integers is $\mathcal{S}_B = \{n \in \mathbb{Z} \setminus \{0\} : P^+(|n|) \le B\}$. "Smooth" always means B-smooth.

$\boldsymbol{\mu}_2 = \boldsymbol{\mu}_2(N) = \{\epsilon \in (\mathbb{Z}/N)^\times : \epsilon^2 = 1\}$ and $Q_N = ((\mathbb{Z}/N)^\times)^2$. For odd N, each $(\mathbb{Z}/p^e)^\times$ is cyclic of even order, so $\boldsymbol{\mu}_2 \cong \{\pm 1\}^{\omega(N)}$ and $[(\mathbb{Z}/N)^\times : Q_N] = 2^{\omega(N)}$.

A *gcd check* of an integer x computes $g = \gcd(x, N)$. If $1 < g < N$, the current procedure stops and returns g as a proper factor (a *factor exit*). If $g = N$, the current object is rejected.

2.2. **Smooth numbers.**

Fact 2.1 (Canfield–Erdős–Pomerance [11])**.** *For each $\varepsilon > 0$, uniformly for $u \to \infty$ with $u \le (1 - \varepsilon) \log x / \log\log x$, $\Psi(x, x^{1/u}) = x\, u^{-u(1+o(1))}$.*

Fact 2.2 (Hildebrand–Tenenbaum [31, Thms. 2(i), 3]; [39, Fact 3.15])**.** *Put $u = \log x / \log y$. There is $\alpha = \alpha(x, y)$ such that, uniformly for $x \ge y \ge 2$ and $1 \le c \le y$,*

$$\Psi(cx, y) = \Psi(x, y)\, c^{\alpha} \Big(1 + O\Big(\frac{1}{u} + \frac{\log y}{y}\Big)\Big), \qquad \alpha = \frac{\log(1 + y/\log x)}{\log y} \Big(1 + O\Big(\frac{\log\log(y+1)}{\log y}\Big)\Big).$$

Fact 2.3 (Tenenbaum [64]; [39, Fact 3.16, Cor. 3.17])**.** *Fix $c>0$. Uniformly for $2\le y\le x$ and integers q with $\omega(q)\le y^{c/\log(1+u)}$,*

$$\Psi_q(x,y)=\frac{\varphi(q)}{q}\Psi(x,y)\Big(1+O_c\Big(\frac{\log(1+u)\log(1+\omega(q))}{\log y}\Big)\Big)\Big(1+O\Big(\frac{\omega(q)}{y}\Big)\Big).$$

We use these only in the following range. Fix $0<\xi_1<\xi_2$. We say x is *in range* if $N^{\xi_1}\le x\le N^{\xi_2}$, and we take $y=B=\lfloor L^a\rfloor$ with $a>0$ fixed. Then $u=\log x/\log B\asymp(\log N/\log\log N)^{1/2}$, so $(\log x)^K<B<x^{1/\log\log x}$ for every fixed K and large N.

Lemma 2.4 (local behaviour of ϱ)**.** *Let $C\ge1$ be fixed. For $N\ge N_0(C,\xi_1,\xi_2,a)$ and $z_1\le z_2\le z_1B^C$ with z_1 in range:*

(i) $\varrho(z_2)\le2\varrho(z_1)$*;*
(ii) $\varrho(z_1)\le(\log N)^{O(C)}\varrho(z_2)$*;*
(iii) for fixed $Z\ge1$, $\Psi(z_1(1+1/Z))-\Psi(z_1)\ge\Psi(z_1)/(3Z)$, and $\Psi(z_1/2)\ge\Psi(z_1)/3$;
(iv) if $z=N^{\xi}L^{O(1)}$ with $\xi_1\le\xi\le\xi_2$, then $\varrho(z)=L^{-\xi/(2a)+o(1)}$.

Proof. By Fact 2.2, $\log(1+B/\log x)=\log B-\log\log x+O(\log x/B)$, so $(1-\alpha)\log B=\log\log x+O(\log\log B)=O(\log\log N)$; in particular $\alpha\to1$. Moreover $\alpha<1$: α solves $\sum_{p\le B}\log p/(p^\alpha-1)=\log x$, the left side is decreasing in α, and at $\alpha=1$ it is $O(\log B)<\log x$. For $1\le c\le B$, Fact 2.2 gives $\varrho(cx)/\varrho(x)=c^{\alpha-1}(1+O(1/u))$, which lies between $B^{\alpha-1}(1-O(1/u))\ge(\log N)^{-O(1)}$ and $1+O(1/u)$. Writing z_2/z_1 as a product of at most $C+1$ factors in $[1,B]$ gives (i) and (ii), since $u\to\infty$ uniformly in range. For (iii), $\Psi(z_1(1+1/Z))-\Psi(z_1)=\Psi(z_1)\big((1+1/Z)^\alpha(1+O(1/u))-1\big)\ge\Psi(z_1)(\alpha/(Z+1)-O(1/u))$. Also $\Psi(z_1/2)=2^{-\alpha}\Psi(z_1)(1+O(1/u))$. For (iv), $u=(\xi\log N+O(\mathcal{L}))/(a\mathcal{L}+O(1))=(\xi/a)(\log N/\log\log N)^{1/2}+O(1)$ and $\log u=\frac12\log\log N(1+o(1))$, so Fact 2.1 gives $\log(1/\varrho(z))=u\log u(1+o(1))=(\xi/(2a))\mathcal{L}(1+o(1))$. □

2.3. **Other published inputs.**

Fact 2.5 (Lenstra–Pomerance [44])**.** *There is a Las Vegas algorithm which, for every integer $n>1$, finds the complete prime factorisation of n in expected time at most $L_n[1/2,1+o(1)]$.*

Fact 2.6 (AKS [1])**.** *Primality of n can be decided deterministically in time $(\log n)^{O(1)}$.*

Fact 2.7 (large sieve; Montgomery–Vaughan [47, Thm. 1], see [46])**.** *For complex numbers a_F supported on K consecutive integers and $Q\ge1$, $\sum_{m\le Q}\sum_{(b,m)=1,\,b \bmod m}\big|\sum_F a_Fe(bF/m)\big|^2\le(K+Q^2)\sum_F|a_F|^2$.*

Fact 2.8 (Brun–Titchmarsh; Montgomery–Vaughan [47, Thm. 2])**.** *For $1\le q<y$ and $(a,q)=1$, $\pi(x+y;q,a)-\pi(x;q,a)<2y/(\varphi(q)\log(y/q))$.*

The inputs for the restricted Bombieri–Vinogradov estimate (Drappeau's Lemme 5, de la Bretèche–Tenenbaum's local bound, and the multiplicative large sieve) are stated in Appendix A. The number-field inputs (Lagarias–Odlyzko, Stark, Landau) are stated in Section 10.

3. Parameters, field law and the algorithm

3.1. **Parameters.** *Fixed parameters.* An odd integer $d\ge3$; an integer $k>(d+1)/4$; $\nu:=1/(4k)$; a rational μ with $\nu+\eta_0\le\mu<1$; a rational $\eta_0>0$. Put

$$h:=h_*(\mu)+\eta_0,\qquad h_*(\mu)=\max\big(\underbrace{2\nu}_{\text{kernel}},\ \underbrace{\mu-(d-2)\nu}_{\text{level }1/2},\ \underbrace{1-d\mu+\nu}_{\text{window}}\big),\qquad \theta:=2\mu+h+(d-1)\nu,$$

and assume $\theta<1$. Put $a:=\theta/\sqrt{2}$.

Quantities depending on N (the integer handed to the splitting routine, see Section 13):

- $B := \lfloor L^a \rfloor$, the smoothness bound;
- $X := \lfloor \sqrt{N/2} \rfloor$; j maximal with $2^{kj} \le X$; $Y := 2^{j-1}$; $J := \{\ell \text{ prime} : Y < \ell \le 2Y\}$; $\pi_J := \#J$;
- $S := \lceil L^{1/(2\sqrt{2})} \rceil$, so that $S^2 = L^{\theta/(2a)+o(1)}$;
- for $\ell \in J$: $V_\ell := \lfloor \sqrt{\ell}\, S \rfloor$, with envelopes $V_- := \lfloor \sqrt{Y}\, S \rfloor \le V_\ell \le V_{\max} := \lfloor \sqrt{2Y}\, S \rfloor$; and $T_\ell := dHV_\ell^{d-1}$;
- $M := \lfloor N^\mu \rfloor$, $H := \lfloor N^h \rfloor$, $W := \lceil N^{\mu+\nu+\eta_0/2} \rceil$;
- $X^* := C'_d M H V_{\max}^d$ and $X_r := 2MV_{\max}/Y$, where $C'_d := 2^{2d+5} + 2d$; and $\rho_r := \varrho(X_r)$, $\rho_F := \varrho(X^*)$.

Since $2^{k(j+1)} > X$, we have $X^{1/k}/4 < Y \le X^{1/k}/2$, so $Y = N^{2\nu+o(1)}$ and every product of k primes of J is at most X. $V_\ell = N^\nu L^{1/(2\sqrt{2})+o(1)}$, $X_r = N^{\mu-\nu} L^{O(1)}$ and $X^* = N^{\mu+h+d\nu} L^{O(1)}$. By Lemma 2.4(iv),

(2)
$$\rho_r \rho_F = L^{-(\mu-\nu+\mu+h+d\nu)/(2a)+o(1)} = L^{-\theta/(2a)+o(1)} = L^{-1/\sqrt{2}+o(1)}, \qquad \text{so} \qquad S^2 \rho_r \rho_F = L^{o(1)}.$$

Lemma 3.1 (size relations)**.** *For $N \ge N_0(d, k, \mu, \eta_0)$:*

(a) $B < 2S^2 < Y < \ell$ for $\ell \in J$; $S < \sqrt{Y}$; $V_\ell < \ell$; and $B < V_\ell/2$.
(b) $2MV_{\max} < W$, $W/(2MV_{\max}) \ge N^{\eta_0/3}$, and $2W < N$.
(c) $M/V_{\max} \ge N^{\eta_0/2}$ and $H/V_{\max}^2 \ge N^{\eta_0/2}$.
(d) $M(H+1) < N$ and $2H + 1 < N$.
(e) $2X^2 < N$.

Proof. (a) $B = L^{\theta/\sqrt{2}} \le L^{1/\sqrt{2}} < 2S^2$ since $\theta < 1$; $S^2 = L^{O(1)}$ while $Y = N^{2\nu+o(1)}$; $V_\ell \le \sqrt{\ell} S < \ell$; $V_\ell/2 \ge N^{\nu+o(1)} > B$. (b) $2MV_{\max} = N^{\mu+\nu} L^{O(1)}$. Also $\mu + \nu + \eta_0/2 \le \theta < 1$, because $\mu \ge \nu$, $h \ge 2\nu + \eta_0$ and $d \ge 3$. (c) $\mu - \nu \ge \eta_0$ and $h - 2\nu \ge \eta_0$. (d) $\mu + h = \theta - \mu - (d-1)\nu < \theta < 1$. (e) N is odd in our application, so $2X^2 \le N - 1$. □

At $d = 9$, $k = 13$ the exponents are $\nu = 1/52$, $\mu = 17/156$, $h = 1/26 + \eta_0$; the rational cofactor has exponent $\mu - \nu = 14/156$ and the algebraic value $\mu + h + d\nu = 50/156 + \eta_0$, which add up to $\theta = 16/39 + \eta_0$.

Lemma 3.2 (optimal μ)**.** *At $\eta_0 \downarrow 0$, the minimum of θ over $\mu \in [\nu, 1)$ is attained at $\mu = (1-\nu)/d$ with $\theta = 2/d + (d^2 + d - 2)\nu/d$ if $k < (d^2+1)/4$, and at $\mu = (1 + (d-1)\nu)/(d+1)$ with $\theta = 3/(d+1) + (4d-2)\nu/(d+1)$ if $k \ge (d^2+1)/4$.*

Proof. At zero slack $\theta(\mu) = \max\big(2\mu + (d+1)\nu,\ 3\mu + \nu,\ 1 + (2-d)\mu + d\nu\big)$. The first two are increasing and the third decreasing, so the minimum is at the intersection of the third with the larger of the first two. The third meets the first at $\mu = (1-\nu)/d$ and the second at $\mu = (1+(d-1)\nu)/(d+1)$. Comparing the first two there, the first is the larger iff $\mu \le d\nu$, which at $\mu = (1-\nu)/d$ means $k < (d^2+1)/4$. □

3.2. **The field law.**

Lemma 3.3 (generalised base expansion; cf. [10, §3])**.** *Let $g_1 \ge 2$, $\gcd(g_0, g_1) = 1$, $\gcd(g_1, N) = 1$, and let w be a prime with $w > g_1$. Put $T_d := Nw$ and, for $i = d, d-1, \dots, 1$: let a_i be the integer with $-g_1/2 < a_i \le g_1/2$ and $a_i \equiv T_i g_0^{-i} \pmod{g_1}$, and $T_{i-1} := (T_i - a_i g_0^i)/g_1$. Put $a_0 := T_0$. Then the form $\hat{F}(x, y) := \sum_{i=0}^d a_i x^i y^{d-i}$ satisfies $\hat{F}(g_0, g_1) = Nw$. If $g_0, g_1 \in [M/2, M]$, then $|a_0| \le 2^d Nw/M^d + d2^{d-1}M$, and $a_d \not\equiv 0 \pmod{g_1}$.*

Proof. $T_i \equiv a_i g_0^i \pmod{g_1}$, so $T_{i-1} \in \mathbb{Z}$, and by induction $T_i = \sum_{j \le i} a_j g_0^j g_1^{i-j}$. At $i = d$ this is $\hat{F}(g_0, g_1) = Nw$. Next, $a_0 = (Nw - \sum_{i \ge 1} a_i g_0^i g_1^{d-i})/g_1^d$, with $Nw/g_1^d \le 2^d Nw/M^d$ and

$|a_i|(g_0/g_1)^i \le (M/2)2^i$. Finally $a_d g_0^d \equiv Nw \pmod{g_1}$, and Nw is a unit modulo g_1 because $\gcd(g_1, N) = 1$ and w is a prime larger than g_1. □

Definition 3.4 (field law $\mathcal{A}$)**.** Let $\mathcal{G} := \{(g_0, g_1) \in \mathbb{Z}^2 \cap [M/2, M]^2 : \gcd(g_0, g_1) = 1\}$ and $C_H := [-H, H]^d \cap \mathbb{Z}^d$. A *field* is $A = (g, w, c)$ with three independent components: $g = (g_0, g_1)$ uniform on $\mathcal{G}$; w uniform among the primes in $(W, 2W]$; c uniform on C_H. Put

$$G(x, y) = G_g(x, y) := g_1 x - g_0 y, \quad R_c(x, y) := \sum_{i=0}^{d-1} c_i x^{d-1-i} y^i, \quad F_A := \hat{F}_{g,w} + G \cdot R_c,$$

where $\hat{F}_{g,w}$ is Lemma 3.3 applied to Nw, and $f = f_A(x) := F_A(x, 1) = \sum_i f_i x^i$. For $P = (a, b)$ put $v(P) := (a^{d-1-i} b^i)_{0 \le i < d}$, so that $F_A(P) = \hat{F}(P) + G(P)\langle c, v(P)\rangle$.

So $F_A(g_0, g_1) = Nw$, $f_d = a_d + g_1 c_0 \equiv a_d \not\equiv 0 \pmod{g_1}$, and if $\gcd(g_1, N) = 1$ the map $x \mapsto m := g_0 g_1^{-1} \bmod N$ kills f, because $f(m) \equiv g_1^{-d} F_A(g_0, g_1) = g_1^{-d} Nw \equiv 0$. Thus $(f, g_1 x - g_0)$ is a number field sieve pair for N, with resultant $\pm Nw$.

Lemma 3.5 (sizes)**.** *Let $N \ge N_0$, A a field, $\ell \in J$, $P = (a, b)$ with $|a|, |b| \le V_\ell$, and $|z| \le T_\ell$. Then $|G(P)| \le 2MV_\ell$, $|\hat{F}(P)| \le C_d M H V_\ell^d N^{-\eta_0/2}$ with $C_d = 2^{2d+5}$, and $|\hat{F}(P) + G(P)z| \le X^*$. In particular $|F_A(P)| \le X^*$.*

Proof. $|G(P)| \le |g_1 a| + |g_0 b| \le 2MV_\ell$. By Lemma 3.3, $|\hat{F}(P)| \le (dM/2 + |a_0|)V_\ell^d$ and $|a_0| \le 2^d Nw/M^d + d2^{d-1}M$. Since $w \le 2W \le 3N^{\mu+\nu+\eta_0/2}$ and $M \ge N^\mu/2$, $Nw/M^d \le 3 \cdot 2^d N^{1-(d-1)\mu+\nu+\eta_0/2}$, which is at most $12 \cdot 2^d MHN^{-\eta_0/2}$ because $h \ge 1 - d\mu + \nu + \eta_0$. Also $d2^d M \le MHN^{-\eta_0/2}$ for large N. This gives the bound on $\hat{F}$. Finally $|G(P)z| \le 2MV_\ell \cdot dHV_\ell^{d-1}$, and $|\langle c, v(P)\rangle| \le dHV_\ell^{d-1} = T_\ell$. □

Lemma 3.6 (homogeneity)**.** *For every field and every P, $g_1^d F_A(P) \equiv Nw\, b^d \pmod{G(P)}$, and $F_A(P) \equiv \hat{F}(P) \pmod{G(P)}$.*

Proof. $g_1 P - b(g_0, g_1) = (G(P), 0)$, and $F_A(x_1 + G, x_2) \equiv F_A(x_1, x_2) \pmod{G}$ because F_A has integer coefficients. Hence $g_1^d F_A(P) = F_A(g_1 P) \equiv F_A(bg_0, bg_1) = b^d F_A(g_0, g_1) = b^d Nw$. The second claim holds because $F_A - \hat{F} = G \cdot R_c$. □

3.3. **Candidates and the counted relations.** For $\ell \in J$ and a field A:

- $\mathcal{P}_\ell := \{P = (a, b) \in \mathbb{Z}^2 :\ b \text{ prime},\ V_\ell/2 < b \le V_\ell,\ |a| \le V_\ell,\ \gcd(a, b) = 1\}$;
- $\Lambda_\ell(g) := \{P \in \mathbb{Z}^2 : \ell \mid G_g(P)\}$, a lattice of index ℓ (since $\gcd(g_0, g_1) = 1$, $g \not\equiv 0 \bmod \ell$);
- the *filtered count*

  $$N_\ell(A) := \#\Big\{P \in \mathcal{P}_\ell \cap \Lambda_\ell(g) :\ \tfrac{1}{32} MV_\ell \le |G(P)|,\ \ G(P)/\ell \in \mathcal{S}_B,\ \ F_A(P) \in \mathcal{S}_B\Big\},$$

  and $N_\ell(A) := 0$ if $\ell \mid N$;
- the set of *descendable primes* $\mathcal{D}_A := \{\ell \in J : N_\ell(A) \ge 1\}$, $\delta_A := |\mathcal{D}_A|/\pi_J$, and $\gamma := \mathbb{E}_{A \sim \mathcal{A}} \delta_A = \Pr_{A,\ell}[N_\ell(A) \ge 1]$, with ℓ uniform on J.

The upper filter $|G(P)| \le 2MV_\ell$ holds automatically (Lemma 3.5). For a counted P we call (P, ℓ) an *atom*. Writing $G(P) = \sigma_P \ell n_P$ with $\sigma_P = \pm 1$ and $n_P > 0$ smooth, an atom certifies $\ell \equiv \sigma_P G(P)/n_P \pmod{N}$ together with the smoothness of $F_A(P)$.

Lemma 3.7 (regularity; cf. [52])**.** *Let λ_1 be the length of a shortest nonzero vector of $\Lambda_\ell(g)$. If $\lambda_1 < \sqrt{\ell}/(2S)$, then $\Lambda_\ell(g) \cap \mathcal{P}_\ell = \emptyset$. Otherwise $\#(\Lambda_\ell(g) \cap [-V_\ell, V_\ell]^2) \le 200S^2$.*

Proof. Write $V = V_\ell$. Let w_0 be a shortest vector, with content $e \le |w_0| < \ell$. Then e is invertible modulo ℓ, so $w_0/e \in \Lambda_\ell(g)$, and minimality gives $e = 1$. Thus w_0 is primitive in $\Lambda_\ell(g)$, and the lattice is the union of the lines $jw' + \mathbb{Z}w_0$ ($j \in \mathbb{Z}$), parallel to w_0 at spacing $\ell/|w_0| > 2S\sqrt{\ell} \ge 2V$. Every point of $[-V, V]^2$ has norm at most $\sqrt{2}V < 2V$, so it lies on $\mathbb{Z}w_0$. The only primitive points there are $\pm w_0$, and $|w_0| < \sqrt{\ell}/(2S) < V/2 < b$ for every $(a, b) \in \mathcal{P}_\ell$. In the regular case take a reduced basis $\lambda_1 \le \lambda_2$ with $\lambda_1\lambda_2 \le (2/\sqrt{3})\ell$. Translates of its fundamental parallelogram at the lattice points of $[-V, V]^2$ are disjoint and lie in a square of side $2V + 4\lambda_2$. Since $V^2 \le \ell S^2$ and $\lambda_2^2 \le (4/3)\ell^2/\lambda_1^2 \le (16/3)S^2\ell$, the count is at most $(2V + 4\lambda_2)^2/\ell \le (8V^2 + 32\lambda_2^2)/\ell \le 200S^2$. □

Lemma 3.8 (coprimality and the forced residue)**.** *Let $N \ge N_0$ have no prime factor $\le B$, let $\ell \in J$ with $\ell \nmid N$, and let $P = (a, b) \in \mathcal{P}_\ell \cap \Lambda_\ell(g)$ with $n := G(P)/\ell \in \mathcal{S}_B$. Put $t := |G(P)|$. Then* $\gcd(t,\, bg_1Nw) = 1$ *and*

$$\hat{F}_{g,w}(P) \equiv c_P\, w \pmod{t}, \qquad c_P := Nb^d g_1^{-d} \in (\mathbb{Z}/t)^\times, \tag{3}$$

where c_P depends on (ℓ, P, g) but not on w.

Proof. b is a prime with $B < V_\ell/2 < b \le V_\ell < \ell$ (Lemma 3.1(a)), so $b \nmid \ell|n| = t$. A prime dividing g_1 and $t = |g_1a - g_0b|$ divides g_0b, hence b (as $\gcd(g_0, g_1) = 1$), which is impossible. N has no prime factor $\le B$, n is smooth and $\ell \nmid N$, so $\gcd(t, N) = 1$. Finally $w > W > 2MV_\ell \ge t$ is prime. Now apply Lemma 3.6. □

Remark 3.9. Without the hypothesis that $G(P)/\ell$ is smooth, (3) can fail (for example $\gcd(g_1, t) = b$ is possible). Such candidates never enter the moment sums.

3.4. **The algorithm.**

Algorithm 3.10 (FACTOR(N))**.**

1. If $N = 1$ return the empty list. If N is prime (AKS), return $[N]$. If $N = m^e$ with $e \ge 2$ (integer roots, $e \le \log_2 N$), return FACTOR(m) with all exponents multiplied by e.
2. Trial-divide N by all primes $\le B(N)$. Record the primes found, and let N' be the cofactor. If $N' = 1$, stop. If N' is prime or a perfect power, handle it as in step 1.
3. Run SPLIT(N') and trial division of N' in parallel, alternating steps, until one returns a proper factor g of N'. (This alternation removes the need for any branch "$N' < N_0$": the algorithm then uses no analytic constant at all, and pays at most a factor 2 in expected time; see Remark A.17.)
4. Return FACTOR(g) merged with FACTOR(N'/g), together with the primes from step 2.

Algorithm 3.11 (SPLIT(N); N odd, all prime factors $> B$, not a prime power)**.**

0. *Guards.* If any structural prerequisite fails ($M < 4$, $J = \emptyset$, $(W, 2W]$ contains no prime, $H < 1$, or $2X^2 \ge N$), return FAIL at once. This happens only for N below an explicit bound, where the trial-division branch of FACTOR finishes the job.
1. *Field selection.* Repeat: $A \leftarrow$ SAMPLE_FIELD; if A is rejected, continue; row $\leftarrow$ DESCEND_ONCE(A); if a row is returned, keep A and go to step 2. (Any factor exit returns the factor.)
2. *Rows.* Call DESCEND_ONCE(A) until M_{rows} rows are held, where M_{rows} is defined in (4).
3. *Post-processing.* Run POSTPROCESS(A, rows) (Algorithm 10.14). It returns a nonempty set S of rows and $\mathsf{Y_S} \in (\mathbb{Z}/N)^\times$. Compute $\mathsf{X_S} := \prod_{i \in \mathsf{S}} r_i \bmod N$ and $g := \gcd(\mathsf{X_S} - \mathsf{Y_S}, N)$. If $1 < g < N$ return g.
4. Otherwise discard all rows and go to step 2 (same field, fresh hidden roots).

Algorithm 3.12 (SAMPLE_FIELD)**.** 1. Draw (g_0, g_1) uniformly from $[M/2, M]^2 \cap \mathbb{Z}^2$ until $\gcd(g_0, g_1) = 1$. Draw a uniform prime $w \in (W, 2W]$ (uniform integers and AKS). Draw c uniform in C_H. Gcd-check g_1, w and then $f_d = a_d + g_1 c_0$. (All three are nonzero and smaller than N, by Lemma 3.1, so these checks either exit with a factor or pass.)
2. Compute f. Let e be its content. Then $e \mid F_A(g_0, g_1) = Nw$ and $e \le |f_d| < N$. Gcd-check e. If it passes, $e \in \{1, w\}$. If $e = w$, *reject*.
3. *Irreducibility certificate* (Lemma 10.1): if no prime $r \in (r_0(d), \frac{1}{2} \log H]$ has $f \bmod r$ of degree d and irreducible over $\mathbb{F}_r$, *reject*.
4. Gcd-check $f_d^{d-2} f'(m)$, where $m = g_0 g_1^{-1} \bmod N$; if $N \mid f'(m)$, *reject*.
5. Return A.

Algorithm 3.13 (DESCEND_ONCE(A))**.** 1. Draw r uniform in $[1, N-1]$; gcd-check r (a non-unit r gives a factor exit). Put $R := r^2 \bmod N$.
2. Let (u_0, v_0) be a shortest nonzero vector of $\Lambda_R := \{(u, v) \in \mathbb{Z}^2 : u \equiv Rv \pmod N\}$ (Gauss reduction, ties broken by a fixed rule). Gcd-check u_0 and v_0.
3. *Shape test.* Factor $|u_0|$ and $|v_0|$ completely with the algorithm of Fact 2.5, and certify each output prime with AKS. Require that $|u_0|$ and $|v_0|$ are each a product of at most k *distinct* primes, all in J; otherwise return FAIL.
4. For each prime $\ell \mid u_0 v_0$, gcd-check ℓ and run SQ_SEARCH(A, ℓ) (cached per (A, ℓ)). If any search fails, return FAIL.
5. Return the row $(r, R, \operatorname{sgn}(u_0 v_0), (e_\ell)_\ell, (P_\ell)_\ell)$, where $e_\ell = +1$ if $\ell \mid u_0$, $e_\ell = -1$ if $\ell \mid v_0$, and P_ℓ is the atom found for ℓ, with the factorisations of n_{P_ℓ} and $F_A(P_\ell)$.

Algorithm 3.14 (SQ_SEARCH(A, ℓ); a deterministic function of (A, ℓ))**.** 1. Gauss-reduce a basis of $\Lambda_\ell(g)$. If $\lambda_1 < \sqrt{\ell}/(2S)$, return FAIL (Lemma 3.7).
2. Enumerate $\Lambda_\ell(g) \cap [-V_\ell, V_\ell]^2$ (at most $200S^2$ points). Keep the points $P \in \mathcal{P}_\ell$ (b prime by AKS) with $G(P) \neq 0$, $F_A(P) \neq 0$ and $|G(P)| \ge MV_\ell/32$. Gcd-check $G(P)$ and $F_A(P)$ for each.
3. Compute the B-smooth parts of all the numbers $|G(P)/\ell|$ and $|F_A(P)|$ with one batch (Proposition 11.2).
4. Return the first kept P, in a fixed order, for which both numbers are B-smooth, together with the factorisations of its two values (trial division by the primes $\le B$). If there is none, return FAIL.

POSTPROCESS is Algorithm 10.14 in Section 10. The number of rows is

$$M_{\text{rows}} := C_{\text{col}} + 12, \qquad C_{\text{col}} := 2 + \pi(B) + \#\{(r, s) : r \le B \text{ prime},\ r \nmid D,\ f(s) \equiv 0\ (r)\} + t_\chi, \tag{4}$$

where $D = f_d \operatorname{disc}(f)$ and $t_\chi = O_d(\log N)$ is defined in Lemma 10.11. So $M_{\text{rows}} \le (d+1)\pi(B) + O_d(\log N)$.

Remark 3.15 (on the design)**.** Each design choice serves one step of the proof.

- *b prime in $(V/2, V]$.* This makes every moving residue reduced (Lemma 3.8), removes small b, and gives the effective wedge bound in Lemma 9.2.
- *The filter $|G(P)| \ge MV/32$.* Candidates with small $|G|$ have algebraic values smaller by a power of N. Their smoothness probability could then exceed ρ_F by $L^{\Theta(1)}$, which would spoil the second moment. The first moment only uses $|G| \asymp MV$, so the filter costs nothing.
- *Two-dimensional rational side.* With (g_0, g_1) random, the rational values of a fixed candidate fill whole fibres of a linear map. The rational part of the first moment is then an interval count, and the partner's rational weight in the second moment can be bounded by divisor switching along a fibre (Step 5 of Theorem 9.7).

- *The multiplier w.* By Lemma 3.6 every algebraic value of a candidate lies in one class modulo $G(P)$, and that class is determined by the rational data. With $F(g_0, g_1) = N$ this class moves together with the partner's rational indicator, and we know no way to decouple them. Proposition 3.16 shows that no randomisation of the coefficients at a fixed resultant can move it. Multiplying the resultant by an independent random prime $w > t$ makes the class $c_P w$ almost uniform (Lemma 9.3).
- *The windows.* The window floor $h \geq 1 - d\mu + \nu + \eta_0$ keeps the large constant coefficient $a_0 \approx Nw/M^d$ invisible at the scale MHV^d (Lemma 3.5). The level floor keeps the moduli $|G(P)| \approx N^{\mu+\nu}$ below the square root of the algebraic values, as Lemma 8.4 requires. The kernel floor $H \geq V^2$ is needed by Lemma 9.2.

Proposition 3.16 (no coefficient-only repair; cf. [39])**.** *If two integral binary forms of degree d take the same value at the primitive point (g_0, g_1), their difference is divisible by $G = g_1 x - g_0 y$ in $\mathbb{Z}[x, y]$. Consequently every form F of degree d with $F(g_0, g_1) = N$ satisfies $g_1^d F(P) \equiv Nb^d \pmod{G(P)}$ for all P.*

Proof. The difference D vanishes at (g_0, g_1), so G divides it over $\mathbb{Q}$; G is primitive, so by Gauss's lemma it divides it over $\mathbb{Z}$. The congruence is Lemma 3.6 with $w = 1$. □

4. Overview of the proof

The proof of Theorem 1.1 has three parts. Figure 1 shows how they depend on each other.

Part I: from relation supply to a factor. Suppose the following averaged statement holds.

$$\gamma = \Pr_{A \sim \mathcal{A},\ \ell \in J}\big[N_\ell(A) \geq 1\big] \ \geq \ L^{-o(1)}. \tag{D_k}$$

- Reverse Markov gives many fields with $\delta_A \geq \gamma/4$.
- In such a field the counting lemma (Theorem 5.2) shows that a uniformly random square R splits as u/v with u, v products of k descendable primes with probability $L^{-o(1)}$ (Theorem 7.2).
- The field-selection rule keeps a field with law proportional to its success probability. A harmonic-mean identity, combined with the factor-exit bookkeeping, bounds the expected cost of all the rows (Proposition 7.3).
- The rows are combined by the post-processing of Section 10. Its output depends only on the squares R_i and on independent coins, so Theorem 6.1 makes each round fruitful with probability $\geq 1/2$.

Part II: relation supply. We prove (D_k) by the Paley–Zygmund inequality (Theorem 9.9). Two moment bounds are needed, and averaging over ℓ is essential in both.

- The first moment (Theorem 8.7): $\mathbb{E}_{\ell,A} N_\ell(A) \geq S^2 \rho_r \rho_F / (\log N)^{O(1)}$. It uses the one-point smoothing Lemma 8.1 and the restricted Bombieri–Vinogradov estimate Lemma 8.4.
- The second moment (Theorem 9.7): $\mathbb{E}_{\ell,A}[N_\ell(N_\ell - 1)] \leq (\log N)^{O(1)} S^4 \rho_r^2 \rho_F^2$. It uses Lemma 9.2 (kernel), Lemma 9.3 (multiplier), Lemma 9.4 (divisor switching) and Theorem 9.5 (large sieve).

Part III: the algorithmic side. This covers exact recognition of $N_\ell(A)$ by batch smoothness (Section 11), rejection of bad fields (Lemma 10.1), rigorous post-processing at fixed degree (Section 10), the cost assembly (Section 12), and uniformity in N (Section 13).

Remark 4.1 (two interfaces that need care)**.** (a) *Order of operations.* The moment bounds are proved under the raw product law $\mathcal{A}$. Rejection of bad fields is then accounted for at the level of the bounded quantity $\delta_A \in [0, 1]$, as $\mathbb{E}[\delta_A 1_{\mathcal{R}^c}] \geq \gamma - \Pr(\mathcal{R})$, and only after that is the field-selection rule applied. Conditioning on acceptance before

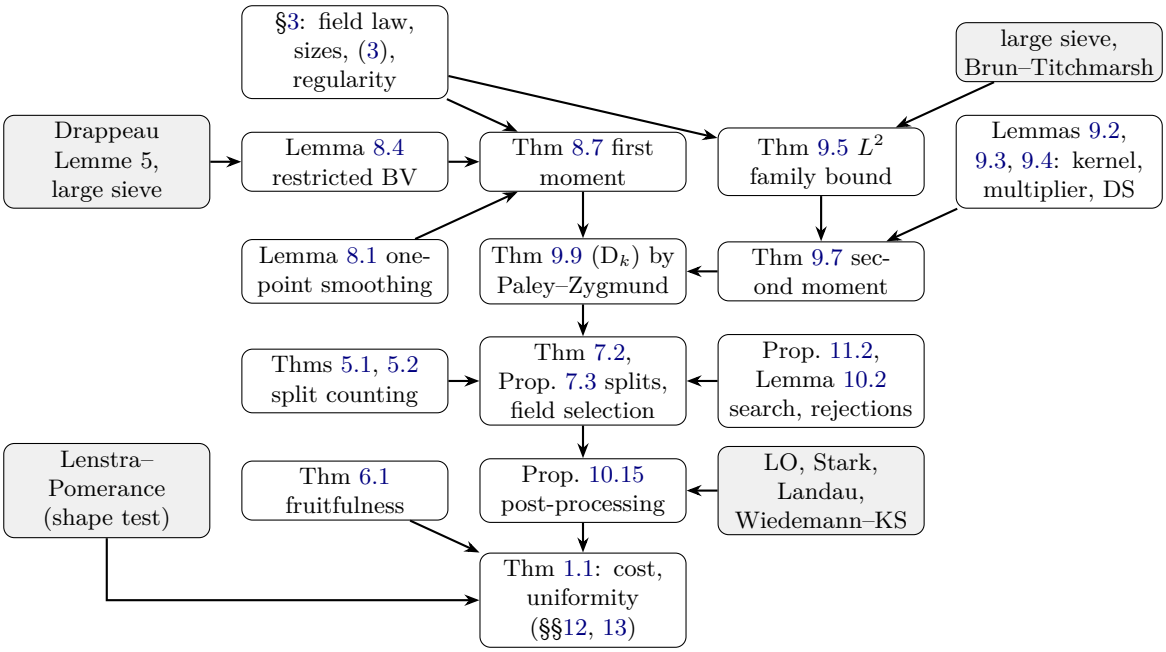


FIGURE 1. Dependencies. Shaded boxes are published results, used only through their statements. Drappeau's Lemme 5 is published, but its proof refers to computations in Harper's unpublished preprint (Appendix A).

the moment estimates would destroy the product structure that the second moment uses.

(b) *Computable S.* The algorithm cannot compute Ψ, so S is an explicit function of N and $\mathbb{E}N_\ell$ need not be ≥ 1. The Paley–Zygmund form $\gamma \geq (\mathbb{E}N)^2/(\mathbb{E}N(N-1)+\mathbb{E}N)$ needs neither $\mathbb{E}N \geq 1$ nor a polylogarithmic lower bound on γ. Indeed $S^2\rho_r\rho_F$ is $L^{o(1)}$ but, by de Bruijn's second-order term, smaller than every negative power of $\log N$. So only $\gamma \geq L^{-o(1)}$ holds, and that suffices because k is fixed.

5. Split counting

In this section N is odd and $h_2 := 2^{\omega(N)} = [(\mathbb{Z}/N)^\times : Q_N]$. For a set $\mathcal{S}$ of integers coprime to N, let $E_{\mathcal{S}}$ be the event that $u \equiv Rv \pmod N$ for some $u, v \in \mathcal{S}$, where R is uniform on Q_N.

Theorem 5.1 (split counting; ratio form of the Dixon–Pomerance lemma)**.** *Let $Y' \geq 2$ with $2Y'^2 < N$, and let $\mathcal{S} \subset [Y'/2, Y'] \cap \mathbb{Z}$ consist of integers coprime to N. Then*

$$\Pr[E_{\mathcal{S}}] \;\geq\; \frac{|\mathcal{S}|^4}{2h_2\varphi(N)Y'^2(1+\log Y')}.$$

Moreover, if $(u,v) \in \mathcal{S}^2$ satisfies $u \equiv Rv$, then (u,v) is an integer multiple of a shortest nonzero vector (u_0, v_0) of $\Lambda_R = \{(x,y) \in \mathbb{Z}^2 : x \equiv Ry \pmod N\}$.

Proof. Put $n(R) := \#\{(u,v) \in \mathcal{S}^2 : u \equiv Rv\}$ for $R \in Q_N$.

First moment. Let $\widehat{G}$ be the group of characters of $(\mathbb{Z}/N)^\times$ trivial on Q_N, so $|\widehat{G}| = h_2$, every $\chi \in \widehat{G}$ is real, and $1[x \in Q_N] = h_2^{-1}\sum_{\chi\in\widehat{G}}\chi(x)$. Then

$$\sum_{R\in Q_N} n(R) = \#\{(u,v) \in \mathcal{S}^2 : uv^{-1} \in Q_N\} = h_2^{-1}\sum_{\chi\in\widehat{G}}\Big(\sum_{u\in\mathcal{S}}\chi(u)\Big)^2 \;\geq\; |\mathcal{S}|^2/h_2,$$

by keeping only the trivial character.

Second moment. If $u \equiv Rv$ and $u' \equiv Rv'$ then $uv' \equiv u'v \pmod N$ and $|uv' - u'v| < Y'^2 < N$, so $uv' = u'v$: the two vectors are proportional. Write $g = \gcd(u, v)$ and $(u, v) = g(u_1, v_1)$. The vectors of $\mathcal{S}^2$ proportional to (u, v) are of the form $j(u_1, v_1)$ with $ju_1 \in [Y'/2, Y']$, and there are at most $Y'/(2u_1) + 1 \le 2g$ of them. Hence

$$\sum_R n(R)^2 \le \sum_{(u,v)\in\mathcal{S}^2} 2\gcd(u, v) \le 2\sum_{g\ge1} g\lfloor Y'/g\rfloor^2 \le 2Y'^2(1 + \log Y').$$

Cauchy–Schwarz. $\#\{R : n(R) > 0\} \ge (\sum n)^2/\sum n^2$. Dividing by $|Q_N| = \varphi(N)/h_2$ gives the bound.

Shortest vector. Λ_R has determinant N. Let $(u, v) \in \mathcal{S}^2 \cap \Lambda_R$ and let (u_0, v_0) be a shortest vector. Both have Euclidean norm at most $\sqrt{2}Y'$, so $|uv_0 - u_0v| \le 2Y'^2 < N$ by Cauchy–Schwarz. The determinant is a multiple of N, hence 0, and the vectors are proportional. A shortest vector is primitive in Λ_R, so $(u, v) \in \mathbb{Z}(u_0, v_0)$. □

Theorem 5.2 (k-products)**.** *Let $\mathcal{D}$ be a set of primes in a dyadic interval, all coprime to N, with $n := |\mathcal{D}| \ge 4k^2$. Let $\mathcal{S}$ be a set of products of exactly k distinct primes of $\mathcal{D}$, contained in a window $[Y'/2, Y']$ with $2Y'^2 < N$, and with $|\mathcal{S}| \ge \binom{n}{k}/(2k)$. Then*

$$\Pr[E_{\mathcal{S}}] \;\ge\; \frac{|\mathcal{S}|^2}{c_k h_2 \varphi(N)}, \qquad c_k := 8k^2\binom{2k}{k}.$$

Proof. As before $\sum_R n(R) \ge |\mathcal{S}|^2/h_2$ and $\sum_R n(R)^2 \le \sum_{(u,v)\in\mathcal{S}^2} \#\{\text{partners of } (u, v)\}$, where the partners are the elements of $\mathcal{S}^2$ proportional to (u, v). Classify (u, v) by the number i of primes that u and v share. Then $(u, v) = (ju_1, jv_1)$ with j the product of the i shared primes. There are at most $\binom{n}{i}\binom{n}{k-i}^2$ such pairs, and each has at most $\binom{n}{i}$ partners $j'(u_1, v_1)$. Now

$$\binom{n}{i}\binom{n}{k-i} = \binom{k}{i}\binom{n}{k}\prod_{t<i}\frac{n-t}{n-k+i-t} \le \binom{k}{i}\binom{n}{k}\Big(1 + \frac{k}{n-k}\Big)^k \le \sqrt{2}\binom{k}{i}\binom{n}{k},$$

since $k^2/(n-k) \le 1/3$ for $n \ge 4k^2$ and $e^{1/3} < \sqrt{2}$. So the i-th class contributes at most $2\binom{k}{i}^2\binom{n}{k}^2 \le 8k^2\binom{k}{i}^2|\mathcal{S}|^2$. Summing, with $\sum_i \binom{k}{i}^2 = \binom{2k}{k}$, gives $\sum n(R)^2 \le c_k|\mathcal{S}|^2$, and Cauchy–Schwarz finishes the proof as before. □

Remark 5.3. No independence, no distribution of the primes and no character sum is used. The square-class condition is absorbed by positivity, which is Dixon's coset pigeonhole written with characters. The factor $2^{\omega(N)}$ is harmless after trial division (Lemma 5.4).

Lemma 5.4 (small factors)**.** *If N has no prime factor $\le B$, then $\omega(N) \le \log N/\log B$ and $2^{\omega(N)} = L^{o(1)}$.*

Proof. $\omega(N)\log B \le \log N$, and $\log N/\log B \asymp (\log N/\log\log N)^{1/2} = o(\mathcal{L})$. □

6. Fruitfulness: Dixon descent

Theorem 6.1 (Dixon descent; Rabin/Dixon fibre symmetry)**.** *Let N be odd with $\omega(N) \ge 2$. Let $r_1, r_2, \ldots$ be independent and uniform on $[1, N-1]$, let $R_i := r_i^2 \bmod N$, and let ξ be further randomness independent of (r_i). Let $\mathcal{F} := \sigma((R_i)_i, \xi)$. Let $\mathcal{E} \in \mathcal{F}$ be an event on which an $\mathcal{F}$-measurable nonempty finite set S of indices with every R_i ($i \in \mathsf{S}$) a unit, and an $\mathcal{F}$-measurable $\mathsf{Y} \in (\mathbb{Z}/N)^\times$, are given, with $\mathsf{X}^2 \equiv \mathsf{Y}^2 \pmod N$ where $\mathsf{X} := \prod_{i\in\mathsf{S}} r_i$. Then, conditional on $\mathcal{F}$, on $\mathcal{E}$ the ratio X/Y is uniform on $\boldsymbol{\mu}_2$, and*

$$\Pr\big[1 < \gcd(\mathsf{X} - \mathsf{Y}, N) < N \,\big|\, \mathcal{F}\big] = 1 - 2^{1-\omega(N)} \ge \tfrac{1}{2} \qquad \text{on } \mathcal{E}.$$

Proof. For every unit square R fix a root $\rho(R)$. Conditional on R_i (a unit), $r_i = \epsilon_i \rho(R_i)$ with ϵ_i uniform on $\boldsymbol{\mu}_2$, because each unit square has exactly $|\boldsymbol{\mu}_2|$ roots and r_i is uniform. These ϵ_i are independent of each other and of ξ. So, conditional on $\mathcal{F}$, they are i.i.d. uniform on $\boldsymbol{\mu}_2$. On $\mathcal{E}$, $\mathsf{X}/\mathsf{Y} = \prod_{i\in\mathsf{S}} \epsilon_i \cdot \prod_{i\in\mathsf{S}} \rho(R_i)\mathsf{Y}^{-1}$, where the second factor is $\mathcal{F}$-measurable and lies in $\boldsymbol{\mu}_2$ (both X/Y and the first factor do). A product of independent uniform elements of a finite group over a nonempty index set is uniform, so X/Y is uniform on $\boldsymbol{\mu}_2$. Finally, for $\zeta \in \boldsymbol{\mu}_2$ and the factorisation $N = \prod p_j^{e_j}$ we have $\zeta \equiv \pm 1 \pmod{p_j^{e_j}}$ for each j, and $\zeta - 1$ is a unit modulo $p_j^{e_j}$ when $\zeta \equiv -1$ there (N is odd). Hence $\gcd(\mathsf{X} - \mathsf{Y}, N) = \gcd(\zeta - 1, N) = \prod_{\zeta\equiv 1\,(p_j^{e_j})} p_j^{e_j}$ with $\zeta = \mathsf{X}/\mathsf{Y}$, which is proper unless $\zeta = \pm 1$. These are 2 of the $2^{\omega(N)}$ elements. □

Remark 6.2 (transcript invariance in Algorithm 3.11)**.** We check that SPLIT satisfies the hypothesis of Theorem 6.1 in each round. Every quantity computed before X_S is a function of the squares R_i and of coins independent of the roots. This covers:

- the field (g, w, c) and its checks;
- the Gauss reduction of Λ_{R_i} (with a fixed tie rule) and the Lenstra–Pomerance factorisation of $|u_0|, |v_0|$, whose coins are fresh;
- the shape decision, and the special-q searches with their cache, which depend on (A, ℓ) only;
- the gcd checks (the unit status of r_i is that of R_i);
- all columns, characters, kernel vectors, the Hensel prime, square-root branches, the choice among three dependencies, and Y_S (Section 10).

The roots enter only through X_S. After a trivial gcd, all rows are discarded and the next round uses fresh roots. For that round we apply Theorem 6.1 with ξ containing everything from earlier rounds, including the old roots, which are independent of the new ones. Field retention biases the law of the field and of the accepted R_i, but not of the fibre signs ϵ_i.

Two cautions. (i) Exactness needs r uniform on all of $[1, N-1]$. Drawing from $[2, N-2]$, for instance, breaks it: at $N = 15$, $R = 1$, the remaining roots $4, 11$ both give proper factors. (ii) Accepted roots are *not* uniform on $(\mathbb{Z}/N)^\times$, because acceptance filters R. What holds, and what is used, is uniformity on the fibre $r_i\boldsymbol{\mu}_2$ given $\mathcal{F}$.

7. Bounded-k splits and field selection

From now on N is the input of SPLIT: odd, with no prime factor $\le B$ and not a prime power, and $N \ge N_0$.

Lemma 7.1 (reverse Markov)**.** *If $Z \in [0, 1]$ and $\mathbb{E}Z \ge \beta$, then $\Pr[Z \ge \beta/2] \ge \beta/2$.*

Proof. $\beta \le \mathbb{E}Z \le \beta/2 + \Pr[Z \ge \beta/2]$. □

Theorem 7.2 (complete descents in one field)**.** *Let A be a field for which* SAMPLE_FIELD *returns A (no rejection and no factor exit), and suppose $|\mathcal{D}_A| \ge 4k^2$. Then one call of* DESCEND_ONCE(A) *returns a row or makes a factor exit with probability at least*

$$c'_k \Big(\frac{\delta_A}{\log N}\Big)^{2k} 2^{-\omega(N)},$$

with $c'_k > 0$ depending only on k.

Proof. The products of k distinct primes of $\mathcal{D}_A \subset J$ lie in $(Y^k, (2Y)^k] \subset (X2^{-2k}, X]$. By pigeonhole one of the $2k$ dyadic windows $[Y'/2, Y']$ with $Y' = X2^{-i}$, $0 \le i < 2k$, contains a set $\mathcal{S}$ of at least $\binom{|\mathcal{D}_A|}{k}/(2k)$ of them, and $2Y'^2 \le 2X^2 < N$ (Lemma 3.1(e)). Every $\ell \in \mathcal{D}_A$ satisfies

$\ell \nmid N$, by definition of N_ℓ. Let E be the event that r is a unit and $R \equiv u/v$ with $u, v \in \mathcal{S}$. Given that r is a unit, R is uniform on Q_N, so Theorem 5.2 gives

$$\Pr_r[E] \geq \frac{\varphi(N)}{N} \cdot \frac{|\mathcal{S}|^2}{c_k h_2 \varphi(N)} = \frac{|\mathcal{S}|^2}{c_k h_2 N}.$$

Now $|\mathcal{S}| \geq (|\mathcal{D}_A|/k)^k/(2k)$ and $|\mathcal{D}_A| = \delta_A \pi_J \geq \delta_A Y/(2 \log Y)$, while $Y^{2k} \geq X^2/16^k \geq N/(4 \cdot 16^k)$ and $2k \log Y \leq \log N$. This gives the displayed lower bound for $\Pr_r[E]$.

On E, Theorem 5.1 shows $(u, v) = j(u_0, v_0)$. Unless a gcd check exits, $|u_0|$ and $|v_0|$ divide u and v, so each is a product of at most k distinct primes of $\mathcal{D}_A$. (This is why the rule says *at most* k: the condition is closed under taking divisors, whereas "exactly k" is not.) So the shape test passes. For each such ℓ we have $N_\ell(A) \geq 1$, and by Proposition 11.2 SQ_SEARCH(A, ℓ) returns an atom unless a gcd check exits first. Hence the call returns a row or exits with a factor. □

Trials and costs. A *trial* is one call of SAMPLE_FIELD followed, if a field is returned, by one call of DESCEND_ONCE. Its outcome is a factor exit, a row, or neither. Let p_F and b be the probabilities of the first two outcomes and $p := p_F + b$. For a field A let ς_A be the probability that DESCEND_ONCE(A) returns a row ($\varsigma_A := 0$ for rejected fields and for fields on which SAMPLE_FIELD exits). Let C_{trial} be a uniform upper bound for the expected cost of one trial and of one call of DESCEND_ONCE, conditional on everything before it. Let T_{post} be a uniform bound for the expected cost of POSTPROCESS. Let $\mathcal{R}$ be the set of fields A with $\gcd(g_1, N) = 1$ that satisfy one of the rejection conditions (i)–(iii) of Lemma 10.2, regardless of other exits. (A field with $\gcd(g_1, N) > 1$ always produces a factor exit.)

Proposition 7.3 (field selection with factor exits)**.** *(i) If $\gamma \geq L^{-o(1)}$ and $\Pr_{\mathcal{A}}(\mathcal{R}) \leq \gamma/2$, then $p \geq L^{-o(1)}$.*

(ii) The expected running time of SPLIT(N) *is at most $3(M_{\text{rows}} C_{\text{trial}} + \mathrm{T}_{\text{post}})/p$.*

Proof. (i) Since $0 \leq \delta_A \leq 1$, $\mathbb{E}[\delta_A 1_{\mathcal{R}^c}] \geq \gamma - \Pr(\mathcal{R}) \geq \gamma/2$. By Lemma 7.1, the set $\mathcal{G}^* := \{A \notin \mathcal{R} : \delta_A \geq \gamma/4\}$ has probability at least $\gamma/4$. For $A \in \mathcal{G}^*$, either SAMPLE_FIELD exits with a factor, or it returns A. In the latter case $|\mathcal{D}_A| \geq \gamma \pi_J/4 \geq 4k^2$ for large N (as $\pi_J = N^{2\nu + o(1)}$), and Theorem 7.2 applies. Hence $p \geq (\gamma/4) c_k'(\gamma/(4 \log N))^{2k} 2^{-\omega(N)} = L^{-o(1)}$, by Lemma 5.4 and because k is fixed.

(ii) Trials are i.i.d., so the number of trials until the outcome is a factor or a row is geometric with mean $1/p$. By Wald's identity the expected cost of selection is at most C_{trial}/p. On a factor exit we are done; in particular, if $b = 0$ the bound is C_{trial}/p. Assume $b > 0$, and let M_{rows} denote a deterministic upper bound for the row target (as in (4)). Otherwise the outcome is a row, and the kept field A^* has law $\Pr[A^* \in dA] = \varsigma_A \Pr_{\mathcal{A}}(dA)/b$. Given $A^* = A$, a round collects M_{rows} rows with an expected $M_{\text{rows}}/\varsigma_A$ calls (or ends earlier with a factor), at expected cost at most $M_{\text{rows}} C_{\text{trial}}/\varsigma_A + \mathrm{T}_{\text{post}}$. By Theorem 6.1 and Remark 6.2, each round that reaches its final gcd succeeds with conditional probability at least $1/2$, whatever happened before. So the expected cost given $A^* = A$ is at most $2(M_{\text{rows}} C_{\text{trial}}/\varsigma_A + \mathrm{T}_{\text{post}})$. Averaging over the law of A^*,

$$\mathbb{E}\big[1/\varsigma_{A^*}\big] = \int \frac{1}{\varsigma_A} \frac{\varsigma_A}{b} \Pr_{\mathcal{A}}(dA) = \frac{\Pr(\varsigma_A > 0)}{b} \leq \frac{1}{b}.$$

The kept-field branch is reached with probability b/p, so its contribution is at most $(b/p) \cdot 2(M_{\text{rows}} C_{\text{trial}}/b + \mathrm{T}_{\text{post}}) \leq 2(M_{\text{rows}} C_{\text{trial}} + \mathrm{T}_{\text{post}})/p$. Factor exits during collection only shorten the run. □

Remark 7.4. No lower bound on ς_A for the kept field is used: the harmonic-mean identity makes the cost depend only on b. Nor do we claim that the law of the non-terminating fields

retains $\gamma/2$. Factor exits are absorbed as successful outcomes, which repairs a gap in an earlier draft of the argument.

8. The first moment

8.1. **One-point smoothing.** The coefficient vector c enters $F_A(P)$ only through $\langle c, v(P)\rangle$. The next lemma replaces Lee and Venkatesan's Lemmas 5.13, 5.18 and 5.20 [39] by a short self-contained argument.

Lemma 8.1 (one-point smoothing)**.** *Let $P = (a, b)$ with $\gcd(a, b) = 1$ and $ab \neq 0$, put $V' = \max(|a|, |b|)$ and $l_i = |a|^{d-1-i}|b|^i$, and let $H \geq 3dV'$. For c uniform on C_H there is a probability measure ϑ on $\mathbb{Z}$ such that:*

(a) $\|\mathrm{Law}(\langle c, v(P)\rangle) - \vartheta\|_{\mathrm{TV}} \leq 3d^2V'/(2H+1)$;
(b) ϑ is the law of $\sum_i Y_i$ with Y_i independent and uniform on integer intervals of lengths $(2H+1)l_i$;
(c) $\vartheta(z)$ is a non-increasing function of $|z - M_0|$ for some $M_0 \in \frac{1}{2}\mathbb{Z}$ with $|M_0| \leq dV'^{d-1}$;
(d) $\max_z \vartheta(z) \leq 1/((2H+1)V'^{d-1})$.

Proof. Write $v = v(P)$. Its entries have $\gcd(a^{d-1}, b^{d-1}) = 1$. The vectors $u_j := b\,e_j - a\,e_{j+1}$ $(0 \leq j \leq d-2)$ lie in $\ker\langle\cdot, v\rangle$. The maximal minor of the matrix with rows u_j that omits column i equals $\pm a^{d-1-i}b^i$, so the gcd of the maximal minors is 1, and the u_j are a basis of the lattice $K := \{x \in \mathbb{Z}^d : \langle x, v\rangle = 0\}$. Each coordinate of a point $\sum s_j u_j$ with $s \in [0, 1)^{d-1}$ has absolute value at most $|a| + |b| \leq 2V'$.

Representations. Let $m \in \mathbb{Z}$ with $|m| \leq \|v\|_2$. Since gcd of the entries is 1, the integer solutions of $\langle x, v\rangle = m$ form a coset $x_0 + K$. The real point $p_m := mv/\|v\|_2^2$ lies on the same affine hyperplane and has $\|p_m\|_\infty \leq |m|/\|v\|_2 \leq 1$. Writing $p_m - x_0 = \sum s_j u_j$ and $x := x_0 + \sum \lfloor s_j \rfloor u_j$, we get $\langle x, v\rangle = m$ and $\|x\|_\infty \leq 1 + 2V'$.

Smoothing. Since $l_i \leq V'^{d-1} \leq \|v\|_2$, for each i we may fix a set S_i of l_i vectors with $\|x\|_\infty \leq 3V'$ on which $\langle\cdot, v\rangle$ is a bijection onto $\{0, \ldots, l_i - 1\}$. Let $x^{(i)}$ be uniform on S_i, independent of each other and of c, and put $c' := c + \sum_i x^{(i)}$. Then $\langle c', v\rangle = \sum_i Y_i$ with $Y_i := c_i v_i + \langle x^{(i)}, v\rangle$. Here $c_i v_i = \pm c_i l_i$ with $\pm c_i$ uniform on $[-H, H]$, so Y_i is uniform on an integer interval of length $(2H+1)l_i$. The Y_i are independent. Let ϑ be the law of $\sum Y_i$. For a fixed shift y, the uniform measures on C_H and $C_H + y$ differ in total variation by at most $\|y\|_1/(2H+1)$. Hence $\|\mathrm{Law}(c) - \mathrm{Law}(c')\|_{\mathrm{TV}} \leq d \cdot 3dV'/(2H+1)$, and (a) follows by applying the map $\langle\cdot, v\rangle$.

(c) A uniform distribution on an integer interval is a log-concave sequence, symmetric about the midpoint of the interval. Convolutions of log-concave sequences without internal zeros are log-concave (see e.g. [62]), and a log-concave sequence is unimodal. A symmetric unimodal sequence is non-increasing in the distance to its centre. The centre of Y_i is $(l_i - 1)/2$ for either sign of v_i, so $|M_0| \leq \sum l_i/2 \leq dV'^{d-1}/2$. (d) A convolution of probability measures has maximum at most the maximum of any factor. Take the factor with $l_i = V'^{d-1}$. □

By (c), ϑ has a layer-cake decomposition

$$\vartheta = \sum_m W_m \,\mathrm{Unif}\big([M_0 - y_m, M_0 + y_m] \cap \mathbb{Z}\big), \qquad W_m \geq 0, \quad \sum_m W_m = 1, \tag{5}$$

because every level set $\{z : \vartheta(z) \geq s\}$ is an integer interval symmetric about M_0.

8.2. **Good moduli.** Put $\varpi := L$ and

$$\mathcal{X} := \big\{X^*\varpi^{-2}(3/2)^i : i \geq 0\big\} \cap \big[X^*\varpi^{-2},\, 16dX^*\big], \qquad |\mathcal{X}| = O(\mathcal{L}).$$

The grid $\mathcal{X}$ does not depend on ℓ, P, g, w, or any residue.

Definition 8.2. A positive integer t is *good* if, for every $x \in \mathcal{X}$ and every r with $(r,t) = 1$, $|\Psi(x;t,r) - \Psi_t(x)/\varphi(t)| \le \Psi_t(x)/(2\varphi(t))$. Otherwise t is *bad*.

Lemma 8.3 (algebraic smoothness at a good modulus)**.** *Let $N \ge N_0$. Let $\ell \in J$, $P = (a,b) \in \mathcal{P}_\ell$, $g \in \mathcal{G}$ and w be such that $t := |G(P)| \in [MV_\ell/32, 2MV_\ell]$, t is good, and $r_0 := \hat{F}_{g,w}(P) \bmod t$ satisfies $(r_0, t) = 1$. Then*

$$\Pr_c\big[F_A(P) \in \mathcal{S}_B\big] \;\ge\; \rho_F/(\log N)^C.$$

Proof. Write $V = V_\ell$, $V' = \max(|a|, b) \in (V/2, V]$, $\hat{F} = \hat{F}_{g,w}(P)$ and $G = G(P)$. Since $F_A(P) = \hat{F} + G\langle c, v(P)\rangle$ and $H \ge 3dV'$, Lemma 8.1 gives $\Pr_c[F_A(P) \in \mathcal{S}_B] \ge \Pr_{Z\sim\vartheta}[\hat{F} + GZ \in \mathcal{S}_B] - 3d^2V/(2H+1)$. Use (5) and put $Y_\varpi := (2H+1)V'^{d-1}/\varpi$.

Small layers. The layers with $y_m < Y_\varpi$ are supported in $|z - M_0| < Y_\varpi$, a set of ϑ-mass at most $(2Y_\varpi + 1)/((2H+1)V'^{d-1}) \le 3/\varpi$ by Lemma 8.1(d).

One large layer. Fix m with $y := y_m \ge Y_\varpi$ and put $y' := \lfloor y/2 \rfloor$. Since $|M_0| \le dV'^{d-1} \le Y_\varpi/2$, the layer's interval contains $[0, y']$ and $[-y', 0]$. Consider $z \in [0, y']$. The values $\hat{F} + Gz$ are all integers $\equiv r_0 \pmod t$ between $\hat{F}$ and $\hat{F} + Gy'$. If $G > 0$ they include every $n \equiv r_0$ with $|\hat{F}| < n \le ty' - |\hat{F}|$. If $G < 0$ they include every $n \equiv r_0$ with $-(ty' - |\hat{F}|) \le n < -|\hat{F}|$, that is, every $-n \equiv -r_0$ in the same range of $|n|$. So, with $r' = \pm r_0$ (a reduced residue),

$$\#\{0 \le z \le y' : \hat{F} + Gz \in \mathcal{S}_B\} \;\ge\; \Psi(x_1; t, r') - \Psi(|\hat{F}|; t, r')$$

for any $x_1 \le ty' - |\hat{F}|$. By Lemma 3.5, $|\hat{F}| \le C_d MHV^d N^{-\eta_0/2}$, while $ty' \ge (MV/32)(Y_\varpi/2 - 1) \ge c_d MHV^d/\varpi$. So $|\hat{F}| \le ty'/4$, and also $ty'/4 \ge X^*\varpi^{-2}$ and $ty'/2 \le 2MV \cdot d(H+1)V^{d-1} \le 16dX^*$. Consecutive points of $\mathcal{X}$ differ by the factor $3/2 < 2$, so we may take $x_1 \in \mathcal{X} \cap (ty'/4, ty'/2]$. Since t is good, and by Fact 2.3 (here $\omega(t) \le \log N$, so its error term is $o(1)$),

$$\Psi(x_1; t, r') \ge \frac{\Psi_t(x_1)}{2\varphi(t)} \ge \frac{\Psi(x_1)}{3t} \ge \frac{y'}{12}\varrho(ty'/4) \ge \frac{y'\rho_F}{(\log N)^C},$$

using Lemma 2.4(i),(ii): $ty'/4$ and X^* differ by a factor $\le \varpi^2 \le B^C$. Trivially $\Psi(|\hat{F}|; t, r') \le |\hat{F}|/t + 1$. The same bound holds for $z \in [-y', 0]$. Dividing by the $2y+1$ points of the layer, the proportion of smooth values in the layer is at least

$$\frac{\rho_F}{3(\log N)^C} - \frac{|\hat{F}|/t + 1}{y} \;\ge\; \frac{\rho_F}{3(\log N)^C} - C\varpi N^{-\eta_0/2}.$$

Conclusion. $\Pr_c \ge (1 - 3/\varpi)\big(\rho_F/(3(\log N)^C) - C\varpi N^{-\eta_0/2}\big) - 3d^2V/(2H+1)$. Here $\rho_F = L^{-O(1)}$, $\varpi = L$ and $V/H \le N^{-\nu-\eta_0/2}$, so the error terms are negligible. □

8.3. **Bad moduli are rare.** The next lemma is a restricted Bombieri–Vinogradov estimate for moduli of the form $\ell \cdot$(smooth). It is the analogue of Lee–Venkatesan's Lemma 8.5 [39] for these moduli. Its proof, in Appendix A, does not use that lemma; it rests on the exact identity of Lemma A.8, the large sieve, and Drappeau's Lemme 5.

Lemma 8.4 (restricted BV for the moduli ℓs; "Lemma 8.5′")**.** *Fix $\xi > 0$ and q_0 with $2\nu < q_0 < \xi/2$. Let $x = N^{\xi+o(1)}$ and $Q = N^{q_0+o(1)}$, and let*

$$\mathcal{T}_Q := \{t = \ell s \in (Q/2, Q] :\; \ell \in J,\; s \ge 1,\; P^+(s) \le B\}.$$

Then, uniformly in such x, Q,

$$\sum_{t\in\mathcal{T}_Q} \max_{(r,t)=1} \Big|\Psi(x;t,r) - \frac{\Psi_t(x)}{\varphi(t)}\Big| \;\le\; \varepsilon_N\, \Psi(x)\, \frac{\#\mathcal{T}_Q}{Q},$$

where $\varepsilon_N = C\exp\bigl(-c_0(\log N)^{1/2}(\log\log N)^{-5/2}\bigr)$ *for some fixed* $C,c_0>0$. *In particular* $\varepsilon_N(\log N)^A\to 0$ *for every fixed* A.

Corollary 8.5. *In the setting of Lemma 8.4 with x ranging over $\mathcal{X}$, the number of bad $t\in\mathcal{T}_Q$ is at most $3|\mathcal{X}|\varepsilon_N\#\mathcal{T}_Q$.*

Proof. If $t\in\mathcal{T}_Q$ is bad, there are $x\in\mathcal{X}$ and $(r,t)=1$ with $|\Psi(x;t,r)-\Psi_t(x)/\varphi(t)|>\Psi_t(x)/(2\varphi(t))\ge\Psi(x)/(3t)\ge\Psi(x)/(3Q)$, by Fact 2.3. Sum Lemma 8.4 over $x\in\mathcal{X}$ and apply Markov's inequality. □

8.4. **Rational fibres.**

Lemma 8.6 (fibres)**.** *Fix $\ell\in J$ and $P=(a,b)\in\mathcal{P}_\ell$, and put $\mathcal{G}_P(n):=\{g\in\mathcal{G}:G_g(P)=\ell n\}$.*

(i) The $g\in\mathbb{Z}^2$ with $G_g(P)=\ell n$ form a line $g^*+\mathbb{Z}(a,b)$, and $|\mathcal{G}_P(n)|\le\sqrt{2}M/V_\ell+1\le 2M/V_\ell$ for every n.

(ii) There is a set I_P of consecutive integers n, of length at least $MV_\ell/(16\ell)-1$, on which ℓn has constant sign, $MV_\ell/32\le\ell|n|\le 2MV_\ell$, and $|\mathcal{G}_P(n)|\ge M/(CV_\ell\log\log N)$.

Proof. Write $V=V_\ell$. (i) $G_g(P)=g_1a-g_0b$ is constant along the direction (a,b), which is primitive with $|(a,b)|\ge b>V/2$, and the square $[M/2,M]^2$ has diameter $M/\sqrt{2}$.

(ii) Let (g_0,g_1) be uniform on the real square $[M/2,M]^2$. Then $G=ag_1-bg_0$ is a sum of independent uniform variables on intervals of lengths $\alpha\le\beta$ with $\beta=V'M/2$, $V'=\max(|a|,b)$. Its density is at least $1/(2\beta)$ on an interval of length β. The density at s equals (length of the chord $\{G=s\}$)$/(|(a,b)|M^2/4)$, so on that interval the chord has length at least $(1/(V'M))\cdot V'\cdot M^2/4=M/4$. Removing $(-MV/32,MV/32)$ from an interval of length $\ge MV/4$ leaves a subinterval of length $\ge MV/16$ of constant sign. Let I_P be the integers n with ℓn in it. For $n\in I_P$ the fibre line meets the closed square in at least $\lfloor(M/4)/|(a,b)|\rfloor\ge M/(8V)$ consecutive lattice points $g^*+j(a,b)$.

Coprimality. Fix $e\ge 1$. Choose x,y with $xa+yb=1$. If $e\mid\gcd(g_0,g_1)$ then $e\mid G=\ell n$. Conversely, if $e\mid\ell n$, the j with $e\mid g^*+j(a,b)$ are exactly those with $j\equiv-(xg_0^*+yg_1^*)\pmod e$: the matrix $\left(\begin{smallmatrix}x&y\\-b&a\end{smallmatrix}\right)$ has determinant 1 and maps g to (xg_0+yg_1,G). By Möbius inversion over $e\mid\ell|n|$, the number of coprime points among $J_0\ge M/(8V)$ consecutive ones is at least $J_0\varphi(\ell|n|)/(\ell|n|)-2^{\omega(\ell|n|)}$. Finally $\varphi(m)/m\gg 1/\log\log m$, and $2^{\omega(m)}=N^{o(1)}$ is small compared with $M/V\ge N^{\eta_0/2}$ (Lemma 3.1(c)). □

8.5. **The first moment.**

Theorem 8.7 (first moment)**.** *For $N\ge N_0$,*

$$\mathbb{E}_{\ell\in J}\mathbb{E}_{A\sim\mathcal{A}}N_\ell(A)\ \ge\ \frac{S^2\rho_r\rho_F}{(\log N)^C}.$$

Proof. At most $O(1)$ primes of J divide N, and we drop them. For $\ell\nmid N$,

$$\mathbb{E}_AN_\ell(A)\ \ge\ \frac{1}{|\mathcal{G}|}\sum_{P\in\mathcal{P}_\ell}\sum_{\substack{n\in I_P,\ n\in\mathcal{S}_B\\ \ell|n|\text{ good}}}\sum_{g\in\mathcal{G}_P(n)}\mathbb{E}_w\Pr_c\bigl[F_A(P)\in\mathcal{S}_B\bigr].$$

This is an inequality, because the filter $|G(P)|\ge MV_\ell/32$ holds on I_P and the other n are dropped. For these (ℓ,P,n,g) and every w, Lemma 3.8 shows that r_0 is reduced, so Lemma 8.3 bounds the inner probability below by $\rho_F/(\log N)^C$. Hence

$$\mathbb{E}_AN_\ell(A)\ge\frac{\rho_F}{(\log N)^C|\mathcal{G}|}\sum_{P\in\mathcal{P}_\ell}\Bigl(\sum_{n\in I_P\cap\mathcal{S}_B}|\mathcal{G}_P(n)|-\frac{2M}{V_\ell}\mathfrak{b}_\ell\Bigr),$$

where $\mathfrak{b}_\ell := \#\{n :\ MV_\ell/32 \le \ell|n| \le 2MV_\ell,\ n \in \mathcal{S}_B,\ \ell|n| \text{ bad}\}$ depends on ℓ only (only $n \in I_P$ enter, and on I_P the filter holds).

Good mass. After reflection, $I_P = (z_1, z_2]$ with $z_2 - z_1 = \lambda \ge MV_\ell/(16\ell) - 1$ and $z_2 \le 2MV_\ell/\ell \le 33\lambda$. If $z_1 \le z_2/2$, Lemma 2.4(iii) with $Z = 1$ gives at least $\Psi(z_2/2)/3$ smooth integers in I_P. Otherwise $z_2 \ge z_1(1 + 1/33)$, and (iii) with $Z = 33$ gives at least $\Psi(z_1)/99$. In both cases, by Lemma 2.4(i),(ii), I_P contains at least $c\lambda\rho_r$ smooth integers. By Lemma 8.6(ii), $\sum_{n \in I_P \cap \mathcal{S}_B} |\mathcal{G}_P(n)| \ge cM^2\rho_r/(\ell \log\log N)$.

Bad mass. The moduli $\ell|n|$ counted by the $\mathfrak{b}_\ell$ lie in $\bigcup_Q \mathcal{T}_Q$ for $O(1)$ dyadic values Q covering $[MV_-/32, 2MV_{\max}]$, all with $Q = N^{\mu+\nu+o(1)}$. Lemma 8.4 applies with $\xi = \mu + h + d\nu$ and $q_0 = \mu + \nu$, because $q_0 - 2\nu = \mu - \nu \ge \eta_0$ and $\xi/2 - q_0 = (h - \mu + (d-2)\nu)/2 \ge \eta_0/2$ by the level floor. By Corollary 8.5,

$$\frac{1}{\pi_J}\sum_{\ell \in J} \mathfrak{b}_\ell \le \frac{2}{\pi_J}\sum_Q 3|\mathcal{X}|\varepsilon_N \#\mathcal{T}_Q \le C|\mathcal{X}|\varepsilon_N \frac{MV_{\max}}{Y}\rho_r(\log N)^C,$$

since $\#\mathcal{T}_Q \le \sum_{\ell \in J} \Psi(Q/\ell) \le C\pi_J(MV_{\max}/Y)\rho_r(\log N)^C$ by Lemma 2.4(i),(ii).

Assembly. $\#\mathcal{P}_\ell \asymp V_\ell^2/\log V_\ell$ (there are $\asymp V/\log V$ primes $b \in (V/2, V]$, and for each at least V values of a coprime to b), $|\mathcal{G}| \asymp M^2$, and $V_\ell^2/\ell \ge S^2/2$. Averaging over ℓ, the good mass contributes at least $cS^2\rho_r\rho_F/(\log N)^{C'}$. The bad mass contributes at most $CS^2\rho_r\rho_F|\mathcal{X}|\varepsilon_N(\log N)^{C''}$, which is smaller by the factor $|\mathcal{X}|\varepsilon_N(\log N)^{O(1)} \to 0$. □

Remark 8.8. The bound is averaged over $\ell \in J$ as well as over the field. The bad-modulus step does not give it for each individual ℓ, and no such statement is needed. The loss is polylogarithmic: $1/\log V$ from prime b, $\log\log N$ from $\varphi(m)/m$, and the $(\log N)^{O(1)}$ of Lemma 2.4(ii).

9. The second moment and relation supply

Throughout this section $\ell \in J$ with $\ell \nmid N$, $V = V_\ell$, and $T = T_\ell = dHV^{d-1}$. For $P \in \mathcal{P}_\ell \cap \Lambda_\ell(g)$ write $Y_P = Y_P(g) := 1$ if $|G_g(P)| \ge MV/32$ and $G_g(P)/\ell \in \mathcal{S}_B$, and $Y_P := 0$ otherwise.

9.1. **Pairs in one lattice.**

Lemma 9.1 (pairs)**.** *Let $P \ne P'$ in $\mathcal{P}_\ell \cap \Lambda_\ell(g)$ and $\Delta := \det(P, P')$. Then $\Delta \ne 0$ and $\ell \mid \Delta$. Put $q := |\Delta|/\ell$; then $q \le 2S^2$. For fixed P and q, at most 12 points $P' \in \mathcal{P}_\ell$ satisfy $|\det(P, P')| = \ell q$.*

Proof. Both points have $b, b' > 0$ and are primitive, so they are not proportional and $\Delta \ne 0$. They are both annihilated modulo ℓ by the nonzero linear form $(x, y) \mapsto g_1 x - g_0 y$, so they are proportional modulo ℓ and $\ell \mid \Delta$. Also $|\Delta| \le 2V^2 \le 2\ell S^2$. For fixed P and D, the solutions of $\det(P, P') = D$ form a line $P'' + \mathbb{Z}P$. It meets $[-V, V]^2$ in a segment of length at most $2\sqrt{2}V$, and consecutive points differ by $|P| \ge V/2$. That gives at most 6 points for each of $D = \pm\ell q$. □

Lemma 9.2 (kernel transfer)**.** *Let $d \ge 3$ and let $P = (a, b) \ne P' = (a', b')$ be primitive, with $\Delta = \det(P, P') \ne 0$. Put $v = v(P)$, $v' = v(P')$ and $\Phi : \mathbb{Z}^d \to \mathbb{Z}^2$, $c \mapsto (\langle c, v\rangle, \langle c, v'\rangle)$. Let κ be an integer with $P' \equiv \kappa P \pmod{\Delta}$.*

- *(i) $\Gamma_\Phi := \Phi(\mathbb{Z}^d) = \{(X, X') \in \mathbb{Z}^2 : X' \equiv \kappa^{d-1}X \pmod{\Delta}\}$, of index $|\Delta|$, and κ is a unit modulo Δ.*
- *(ii) $\ker\Phi = Q_{PP'} \cdot \mathbb{Z}[x, y]_{d-3}$ with $Q_{PP'} = (bx - ay)(b'x - a'y)$, identifying $\mathbb{Z}^d$ with forms of degree $d - 1$.*
- *(iii) If $|a|, |b|, |a'|, |b'| \le V$ and $H \ge 4V^2$, then for every $z \in \mathbb{Z}^2$, $\Pr_{c \in C_H}[\Phi(c) = z] \le C_d|\Delta|/(H^2\|v \wedge v'\|) \le C_d/(H^2(bb')^{d-2})$.*

Proof. For $\gamma \in GL_2(\mathbb{Z})$, $v(\gamma P) = \rho(\gamma)v(P)$ with $\rho(\gamma) \in GL_d(\mathbb{Z})$ (the symmetric power). By Cauchy–Binet, the gcd of the 2×2 minors of the matrix with rows v, v' is unchanged. Choose γ with $\gamma P = (1, 0)$. Then $\gamma P' = (u, \pm\Delta)$ with $\gcd(u, \Delta) = 1$, and the nonzero minors of the matrix with rows $(1, 0, \dots, 0)$ and $v(u, \pm\Delta)$ are $\pm u^{d-1-j}\Delta^j$, $j \geq 1$. Their gcd is $|\Delta|$. Hence $\Phi(\mathbb{Z}^d)$ has index $|\Delta|$. In these coordinates $P' \equiv uP \pmod{\Delta}$, so κ exists, is a unit, and by homogeneity $v' \equiv \kappa^{d-1}v \pmod{\Delta}$. Thus $\Phi(\mathbb{Z}^d)$ lies in the displayed lattice, which also has index $|\Delta|$. This proves (i).

(ii) $\langle c, v(P)\rangle = C(P)$ for the form $C = \sum c_i x^{d-1-i}y^i$. If $C(P) = C(P') = 0$ then $(bx - ay) \mid C$ and $(b'x - a'y) \mid C$ over $\mathbb{Q}$, and the two factors are coprime since $\Delta \neq 0$. By Gauss's lemma the quotient is integral. Conversely every such multiple lies in the kernel.

(iii) The kernel K_2 has rank $d - 2$. It is saturated, and its orthogonal complement in $\mathbb{Z}^d$ is the saturation of $\mathbb{Z}v + \mathbb{Z}v'$. That lattice has determinant $\|v \wedge v'\|$ and index $|\Delta|$ in its saturation (the gcd of the minors, by (i)). A saturated lattice and its orthogonal complement have the same determinant, so $\det K_2 = \|v \wedge v'\|/|\Delta|$. The monomial multiples of $Q_{PP'}$ form a basis of K_2 with entries at most $2V^2$, so a fundamental cell has ℓ^∞-diameter at most $2(d-2)V^2 \leq dH$. A fibre $\Phi^{-1}(z) \cap C_H$ is $(c_0 + K_2) \cap C_H$. The translates of the cell at its points are disjoint and lie in the intersection of the cube $[-(d+1)H, (d+1)H]^d$ with a $(d-2)$-flat, whose $(d-2)$-volume is at most $C_d H^{d-2}$. Hence the fibre has at most $C_d H^{d-2}|\Delta|/\|v \wedge v'\|$ points, and we divide by $(2H+1)^d$. Finally, the minor of columns $d-2, d-1$ equals $ab^{d-2}b'^{d-1} - b^{d-1}a'b'^{d-2} = (bb')^{d-2}\Delta$, so $\|v \wedge v'\| \geq |\Delta|(bb')^{d-2}$. □

For candidates $b, b' > V/2$, so (iii) gives $\Pr_c[\Phi(c) = z] \leq C_d/(H^2V^{2d-4})$. The kernel floor $h \geq 2\nu + \eta_0$ gives $H \geq 4V^2$ for large N.

9.2. **The multiplier.**

Lemma 9.3 (residue spreading)**.** *For $N \geq N_0$, every $1 \leq t \leq 2MV_{\max}$ and every $r \in \mathbb{Z}$,*

$$\Pr_w[w \equiv r \pmod{t}] \leq \frac{C_W}{\varphi(t)}, \qquad C_W := \frac{24}{\eta_0}.$$

Consequently, if (ℓ, P, g) satisfy the hypotheses of Lemma 3.8 and $\Phi_0 : \mathbb{Z}/t \to [0, \infty)$ does not depend on w, then $\mathbb{E}_w \Phi_0(\hat{F}_{g,w}(P) \bmod t) \leq (C_W/\varphi(t)) \sum_{c \in (\mathbb{Z}/t)^\times} \Phi_0(c)$.

Proof. If $(r, t) > 1$ the probability is 0, since every w is a prime larger than t. Otherwise Fact 2.8 with $x = W$, $y = W$ gives $\pi(2W; t, r) - \pi(W; t, r) < 2W/(\varphi(t)\log(W/t))$, while $\pi(2W) - \pi(W) \geq W/(2\log W)$. So the probability is at most $4\log W/(\varphi(t)\log(W/t))$. Here $\log W \leq \log N$, and $\log(W/t) \geq (\eta_0/3)\log N$ by Lemma 3.1(b). For the consequence, $\hat{F}_{g,w}(P) \equiv c_P w$ with c_P a unit independent of w (Lemma 3.8), and $x \mapsto c_P x$ permutes $(\mathbb{Z}/t)^\times$. □

9.3. **Divisor switching.**

Lemma 9.4 (smooth numbers in short progressions)**.** *Fix $C_0 \geq 1$. Let X be in range, $1 \leq q \leq B^{C_0}$, and let $I \subset [XB^{-C_0}, X]$ be an interval of length $A \geq XB^{-C_0}$. Then for every $r \in \mathbb{Z}$,*

$$\#\{n \in I : \ n \equiv r \pmod{q},\ P^+(n) \leq B\} \ \leq \ \frac{A}{q}\,\varrho(X)\,(\log N)^{O(C_0)}.$$

Proof. Let $e = \gcd(r, q)$. If e has a prime factor $> B$ the count is 0. Otherwise write $n = en_1$. Then n_1 lies in I/e, an interval of length A/e, and $n_1 \equiv r/e \pmod{q'}$ with $q' = q/e$ and $\gcd(r/e, q') = 1$, so $(n_1, q') = 1$. Put $D := \min(A/q, X/(eB^{C_0}))$. Then $D \leq \min(I/e)$. Every smooth $n_1 \geq D$ has a divisor $d_1 \in (D/B, D]$: take the largest product of its smallest prime factors not exceeding D. Write $n_1 = d_1 m$. Then $(d_1, q') = 1$, $m \equiv (r/e)d_1^{-1} \pmod{q'}$, and m lies in an interval of length $A/(ed_1)$. So there are at most $A/(d_1 q) + 1 \leq 2A/(d_1 q)$ choices of m

(as $d_1 \le A/q$), and the count is at most $(2A/q)\sum_{D/B<d_1\le D,\, P^+(d_1)\le B} 1/d_1$. Splitting into dyadic blocks and using Lemma 2.4(i), this sum is at most $8\log_2(2B)\varrho(D/B)$. Finally $D \ge XB^{-2C_0}$, so $X/(D/B) \le B^{2C_0+1}$; in particular D/B is in range (with ξ_1 replaced by $\xi_1/2$), and Lemma 2.4(ii) gives $\varrho(D/B) \le \varrho(X)(\log N)^{O(C_0)}$. □

9.4. **An all-residue second moment for the family $\ell^2 nq$.** Put $\mathcal{S}^* := \mathcal{S}_B \cap [-X^*, X^*]$, so $\#\mathcal{S}^* \le 2\Psi(X^*)$, and

$$S(\alpha) := \sum_{F\in\mathcal{S}^*} e(\alpha F), \qquad \mathfrak{Q}(m) := \#\{(F,F') \in \mathcal{S}^{*2} : F \equiv F' \pmod m\},$$

$$f(m') := \sum_{\substack{b \bmod m' \\ (b,m')=1}} |S(b/m')|^2.$$

Orthogonality and the reduction of fractions give the exact identity

$$\mathfrak{Q}(m) = \frac{1}{m}\sum_{h \bmod m} |S(h/m)|^2 = \frac{1}{m}\sum_{m'|m} f(m'), \qquad f(1) = (\#\mathcal{S}^*)^2. \tag{6}$$

Let $X'_+ := 2MV_{\max}/Y = X_r$, and $\mathcal{N}_\ell := \{n \in \mathbb{Z} :\ n \in \mathcal{S}_B,\ MV_\ell/32 \le \ell|n| \le 2MV_\ell\}$, so $\#\mathcal{N}_\ell \le 2\Psi(X'_+)$. Fix C_0 so that

$$R_0 := B^{C_0} \ \ge\ 64S^2L/(\rho_r\rho_F)$$

for large N. This is possible because the right side is $L^{O(1)}$. Note that $R_0 < Y$.

Theorem 9.5 (L^2 bound for the family; "$L2_{\text{fam}}$"). *Uniformly for $1 \le q \le 2S^2$,*

$$\mathbb{E}_{\ell\in J}\sum_{n\in\mathcal{N}_\ell}\ \sum_{m'|\ell^2|n|q,\ m'>1} f(m') \ \le\ \tau(q)(\log N)^C\,\Psi(X'_+)\Psi(X^*)^2, \tag{7}$$

and consequently

$$\mathbb{E}_{\ell\in J}\sum_{n\in\mathcal{N}_\ell}\frac{\mathfrak{Q}(\ell^2|n|q)}{\varphi(\ell|n|)} \ \le\ \tau(q)(\log N)^C\,\frac{\Psi(X'_+)\Psi(X^*)^2}{Yq(MV_-)^2}. \tag{8}$$

Proof. Every prime factor of $|n|q$ is at most $\max(B, 2S^2) < Y$, so $|n|q$ has no prime factor in J.

Rearrangement. Since X^* does not depend on ℓ, neither does f. So the left side of (7) is $\sum_{m'>1} f(m')\mathrm{mult}(m')$ with $\mathrm{mult}(m') := \mathbb{E}_\ell\#\{n \in \mathcal{N}_\ell : m' \mid \ell^2|n|q\}$. Partition the $m' > 1$ into three classes: (A) no prime factor in J and $m' \le R_0$; (B) no prime factor in J and $m' > R_0$; (C) some prime factor in J. The partition does not depend on ℓ. In classes (A) and (B), $m' \mid \ell^2|n|q$ iff $m' \mid |n|q$. In class (C), $m' \mid \ell^2|n|q$ forces the J-primes of m' to be exactly $\{\ell\}$, with exponent $i \in \{1,2\}$.

For any m' write $m_2 := \gcd(m', q)$ and $m_1 := m'/m_2$. Then $m' \mid |n|q$ iff $m_1 \mid n$, because $\gcd(m_1, q/m_2) = 1$. The map $m' \mapsto (m_1, m_2)$ is injective and $m_2 \mid q$.

Class (A). By Parseval over all $b \bmod m'$, $f(m') \le m'\sum_{c \bmod m'} N_c^2$ with $N_c := \#\{F \in \mathcal{S}^* : F \equiv c\ (m')\}$. Lemma 9.4 on $\pm[X^*R_0^{-1}, X^*]$, plus the trivial bound $X^*/(R_0m') + 1$ on $|F| < X^*/R_0$, gives $N_c \le (X^*/m')\rho_F(\log N)^C$ (using $R_0^{-1} \le \rho_F$ and $m' \le R_0 \le X^*\rho_F$). Hence $f(m') \le m'\max_c N_c \cdot \#\mathcal{S}^* \le \Psi(X^*)^2(\log N)^C$. For the multiplicity, $\mathrm{mult}(m') = 0$ unless m_1 is smooth, and then $\mathrm{mult}(m') \le 2\Psi(X'_+/m_1) \le (\Psi(X'_+)/m_1)(\log N)^{O(C_0)}$ by Lemma 2.4(ii). Since $\sum_{m'\le R_0} 1/m_1 \le \tau(q)(1+\log R_0)$, class (A) contributes at most $\tau(q)(\log N)^C\Psi(X'_+)\Psi(X^*)^2$.

Class (B). Here $\mathrm{mult}(m') \le \#\{1 \le |n| \le X'_+ : m_1 \mid n\} \le 2X'_+q/m'$, and $\mathrm{mult}(m') = 0$ unless $m' \le X'_+q$. On dyadic blocks $(R, 2R]$ with $R = R_02^i \le X'_+q$, Fact 2.7 with $K = 2\lfloor X^*\rfloor + 1$ gives $\sum_{m'\in(R,2R]} f(m') \le (3X^* + 4R^2)\cdot 2\Psi(X^*)$. Summing over blocks, class (B) contributes

at most $CX'_+qX^*\Psi(X^*)/R_0 + C(X'_+q)^2\Psi(X^*)$. Relative to $\Psi(X'_+)\Psi(X^*)^2 = X'_+X^*\rho_r\rho_F\Psi(X^*)$ these are $Cq/(R_0\rho_r\rho_F) \le 1$ and $CX'_+q^2/(X^*\rho_r\rho_F) = N^{-h-(d+1)\nu}L^{O(1)} = o(1)$.

Class (C). Let $\ell = \ell(m')$ be the unique J-prime of m' and $m'' := m'/\ell^i$. Only this ℓ contributes, so $\mathrm{mult}(m') = \pi_J^{-1}\#\{n \in \mathcal{N}_\ell : m'' \mid |n|q\}$. We have two bounds: $\mathrm{mult}(m') \le 2\Psi(X'_+)/\pi_J$, and, as $m'' \ge m'/(4Y^2)$, $\mathrm{mult}(m') \le 8X'_+qY^2/(\pi_J m')$. The relevant m' lie in $[Y, 4Y^2X'_+q]$, which is $O(\log N)$ dyadic blocks. On each block use the large sieve bound $(3X^* + 4R^2)\cdot 2\Psi(X^*)$, and the first bound for ν against $3X^*$ and the second against $4R^2$ (if $\mathrm{mult} \le A_1$ and $\mathrm{mult} \le A_2$ then $\mathrm{mult}\cdot(K_1+K_2) \le A_1K_1 + A_2K_2$). Class (C) contributes at most

$$\frac{C\log N}{\pi_J}\Psi(X'_+)X^*\Psi(X^*) + \frac{C}{\pi_J}(X'_+)^2q^2Y^4\Psi(X^*).$$

Relative to the target these are $C\log N/(\pi_J\rho_F) = N^{-2\nu}L^{O(1)}$ and $CX'_+q^2Y^4/(\pi_JX^*\rho_r\rho_F) = N^{(5-d)\nu-h}L^{O(1)}$. The latter is $o(1)$ because $h \ge 2\nu + \eta_0 > (5-d)\nu$ for every $d \ge 3$. This proves (7).

(8). For $n \in \mathcal{N}_\ell$ put $m = \ell^2|n|q = \ell qt$ with $t = \ell|n| \ge MV_-/32$, and note $\varphi(t) \ge t/(C\log\log N)$. By (6), $\mathfrak{Q}(m)/\varphi(t) \le C\log\log N\,\big[(\#\mathcal{S}^*)^2 + \sum_{m'\mid m, m'>1} f(m')\big]/(Yq(MV_-)^2)$. Sum over n, average over ℓ, and use $\#\mathcal{N}_\ell \le 2\Psi(X'_+)$ and (7). □

Remark 9.6. The average over ℓ is essential in class (C). For a single ℓ, the frequencies with denominator ℓ^2 contribute $\ell^2\sum_c N_c^2$ over classes modulo ℓ^2, whose trivial bound is $X^*\Psi(X^*) \gg \Psi(X^*)^2$. Improving it would require L^2-equidistribution of smooth numbers modulo $\ell^2 = B^{\Theta(u)}$, which is out of reach unconditionally. Note also that the statement is an upper bound at the main-term scale, with no saving; it does not follow from Harper's Barban–Davenport–Halberstam theorem over all moduli [27, Thm. 2], whose error term $\Psi^2L^{-o(1)}$ exceeds the family's main term.

9.5. **The second moment.**

Theorem 9.7 (second moment; "LD*")**.** *For $N \ge N_0$,*

$$\mathbb{E}_{\ell\in J}\mathbb{E}_{A\sim\mathcal{A}}\big[N_\ell(A)(N_\ell(A)-1)\big] \;\le\; (\log N)^C\, S^4\rho_r^2\rho_F^2.$$

Proof. Step 0 (law). If $\ell \mid N$ then $N_\ell = 0$. For $\ell \nmid N$, $N_\ell(N_\ell-1)$ counts ordered pairs of distinct counted points, so

$$\mathbb{E}_A[N_\ell(N_\ell-1)] \le \mathbb{E}_{g,w}\sum_{P\ne P'} Y_PY_{P'}\Pr_c\big[F_A(P), F_A(P') \in \mathcal{S}_B\big],$$

the sum running over $P \ne P'$ in $\mathcal{P}_\ell \cap \Lambda_\ell(g)$. By Lemma 9.1, $|\Delta| = \ell q$ with $1 \le q \le 2S^2$. Only $\hat F_{g,w}$ depends on w.

Step 1 (kernel). By Lemma 9.2(iii), $\Pr_c[\Phi(c) = z] \le C_d/(H^2V^{2d-4})$. Since $\Phi(C_H) \subset [-T,T]^2$ and $\Gamma_\Phi = \{z_2 \equiv \kappa^{d-1}z_1\ (\Delta)\}$,

$$\Pr_c[\text{both smooth}] \le \frac{C_d}{H^2V^{2d-4}}\,Z, \qquad Z := \sum_{j \bmod |\Delta|} A_j(P)A_{\kappa^{d-1}j}(P'),$$

where $A_j(P) := \#\{|z| \le T : z \equiv j\ (|\Delta|),\ \hat F(P) + G(P)z \in \mathcal{S}_B\}$.

Step 2 (AM–GM). κ is a unit modulo Δ, so $j \mapsto \kappa^{d-1}j$ permutes the classes and $Z \le \frac12(\mathfrak{q}(P) + \mathfrak{q}(P'))$ with $\mathfrak{q}(P) := \sum_j A_j(P)^2$. This inequality is pointwise but one-sided.

Step 3 (enlargement). $z \mapsto F = \hat F(P) + G(P)z$ is injective, $F \equiv r_0 \pmod t$ with $t = |G(P)|$, and $z \equiv z' \pmod{|\Delta|}$ iff $F \equiv F' \pmod{t|\Delta|}$, for either sign of $G(P)$. All such F lie in $[-X^*, X^*]$

(Lemma 3.5). Hence

$$\mathfrak{q}(P) \le \mathfrak{Q}_0(t, |\Delta|; r_0) := \#\{(F, F') \in \mathcal{S}^{*2} : F \equiv F' \equiv r_0 \ (t),\ F \equiv F' \ (t|\Delta|)\},$$

a function of r_0 alone, which does not depend on w otherwise.

Step 4 (average over w). When $Y_P = 1$, Lemma 9.3 applies to $\Phi_0 = \mathfrak{Q}_0(t, |\Delta|; \cdot)$. Summing over the reduced classes counts a subset of the pairs with $F \equiv F' \pmod{t|\Delta|}$, so $\mathbb{E}_w \mathfrak{Q}_0(t, |\Delta|; r_0) \le (C_W/\varphi(t))\mathfrak{Q}(t|\Delta|)$. The same holds for P'. The pair sum is symmetric under $(P, P') \mapsto (P', P)$, so

$$\mathbb{E}_A[N_\ell(N_\ell - 1)] \le \frac{C_d C_W}{H^2 V_-^{2d-4}}\, \mathbb{E}_g \sum_{P \ne P'} Y_P Y_{P'} \frac{\mathfrak{Q}(\ell q\, t)}{\varphi(t)}, \qquad t = |G_g(P)|,\ q = |\det(P, P')|/\ell. \tag{9}$$

No joint law of the two residues under the shared w is needed.

Step 5 (the partner's rational weight). Fix $P = (a, b)$ and P' with $|\Delta| = \ell q$. The algebraic factor in (9) depends on g only through $n := G_g(P)/\ell$, and $Y_P = 1$ iff $n \in \mathcal{N}_\ell$. For fixed n, the g with $G_g(P) = \ell n$ are $g^* + j(a, b)$ (Lemma 8.6), at most $2M/V$ values of j. Along the fibre, $G_g(P') = G_{g^*}(P') - j\Delta$, so $n' = G_g(P')/\ell$ runs through a progression of difference q in an interval of length at most $2qM/V$. The filter requires $|n'| \in [X_r/256, X_r]$. For each sign, enlarge the n'-interval to length $\max(\cdot, X_r B^{-C_0})$ inside $[X_r B^{-C_0}, X_r]$. Lemma 9.4 (modulus $q \le 2S^2 \le B^{C_0}$) gives

$$\#\{g \in \mathcal{G} : G_g(P) = \ell n,\ Y_{P'}(g) = 1\} \le U := \frac{3M}{V_-} \rho_r (\log N)^C,$$

uniformly in n, using $X_r B^{-C_0}/q \le M/V$. Hence

$$\sum_g Y_P Y_{P'} \frac{\mathfrak{Q}(\ell q t)}{\varphi(t)} \le U \sum_{n \in \mathcal{N}_\ell} \frac{\mathfrak{Q}(\ell^2 |n| q)}{\varphi(\ell |n|)}.$$

Step 6 (partners). For fixed P and q there are at most 12 points P' (Lemma 9.1).

Step 7 (assembly). Bound the ℓ-dependent geometric factors by their envelopes: $\#\mathcal{P}_\ell \le C V_{\max}^2$, $|\mathcal{G}| \ge cM^2$, $V^{-(2d-4)} \le V_-^{-(2d-4)}$. Then apply (8):

$$\mathbb{E}_{\ell, A}[N_\ell(N_\ell - 1)] \le \frac{C}{H^2 V_-^{2d-4}} \cdot \frac{V_{\max}^2}{M^2} \cdot \frac{M \rho_r}{V_-} (\log N)^C \sum_{q \le 2S^2} \tau(q) \frac{\Psi(X'_+) \Psi(X^*)^2}{Y q (MV_-)^2}.$$

Insert $\Psi(X'_+) = X'_+ \rho_r = (2MV_{\max}/Y)\rho_r$ and $\Psi(X^*) = X^* \rho_F \le CMHV_{\max}^d \rho_F$. The powers of M and H cancel exactly:

$$\frac{1}{H^2 V^{2d-4}} \cdot \frac{V^2}{M^2} \cdot \frac{M}{V} \cdot \frac{MV}{Y} \cdot \frac{M^2 H^2 V^{2d}}{Y (MV)^2} = \frac{V^4}{Y^2} \asymp S^4,$$

up to constants from $V_- \asymp V_{\max}$. Finally $\sum_{q \le 2S^2} \tau(q)/q \le (1 + \log 2S^2)^2 = (\log N)^{O(1)}$. □

Remark 9.8. The residue is integrated out over w (Step 4) *before* any arithmetic estimate is applied. After that, the algebraic factor depends on g only through n, and this is what lets Step 5 sum the partner's rational indicator along a fibre by divisor switching. In the design without w (i.e. $F(g_0, g_1) = N$), the forced residue moves with the fibre index together with the partner indicator (Proposition 3.16), and we do not know how to prove the corresponding bound.

9.6. **Relation supply.**

Theorem 9.9 (relation supply)**.** *For $N \geq N_0$, $\gamma = \Pr_{A,\ell}[N_\ell(A) \geq 1] \geq L^{-o(1)}$. More precisely, with $m := S^2\rho_r\rho_F = L^{o(1)}$ and a fixed polylogarithm $\mathcal{P} = (\log N)^C$, $\gamma \geq 1/(\mathcal{P}^3 + \mathcal{P}/m)$.*

Proof. By Cauchy–Schwarz (Paley–Zygmund) for the nonnegative integer variable $N_\ell(A)$ under the product law of (ℓ, A), $\gamma \geq (\mathbb{E}N)^2/\mathbb{E}N^2 = (\mathbb{E}N)^2/(\mathbb{E}N(N-1) + \mathbb{E}N)$. The function $z \mapsto z^2/(K+z)$ is increasing on $z > 0$. By Theorems 8.7 and 9.7, $\mathbb{E}N \geq m/\mathcal{P}$ and $\mathbb{E}N(N-1) \leq \mathcal{P}m^2$. So $\gamma \geq (m/\mathcal{P})^2/(\mathcal{P}m^2 + m/\mathcal{P})$. Finally $m = L^{o(1)}$ by (2). □

Remark 9.10. m is in general smaller than every negative power of $\log N$: the leading terms of $\log(1/\rho_r) + \log(1/\rho_F)$ and $\log S^2$ cancel, but the second-order terms of the de Bruijn expansion do not. So $\gamma \geq 1/(\log N)^{O(1)}$ need not hold for our explicit S. Only $\gamma \geq L^{-o(1)}$ is used.

9.7. **A rational pair correlation (not used in the proof).** For comparison we record the purely rational version of Step 5, which explains why the rational side is randomised in two dimensions. Put $R_1 := \mathbb{E}_{\ell,g} \sum_P Y_P$ and $R_2 := \mathbb{E}_{\ell,g} \sum_{P \neq P'} Y_P Y_{P'}$ (sums over $\mathcal{P}_\ell \cap \Lambda_\ell(g)$).

Proposition 9.11 ("Theorem R")**.** *$R_1 \geq S^2\rho_r/(\log N)^C$ and $R_2 \leq (\log N)^C S^4 \rho_r^2$.*

Proof. The lower bound is the rational part of Theorem 8.7. For the upper bound fix ℓ, P and P' with $|\Delta| = \ell q$. The map $g \mapsto (n, n') = (G_g(P)/\ell, G_g(P')/\ell)$ on $\{g : \ell \mid G_g(P)\}$ is injective, because the matrix $\left(\begin{smallmatrix} a & -b \\ a' & -b' \end{smallmatrix}\right)$ has determinant $-\Delta \neq 0$. Its image is the lattice $\Gamma_q = \{(n, n') : n' \equiv \lambda n \ (q)\}$, $\lambda \equiv b'b^{-1}$, of index q ($b > 2S^2 \geq q$ is prime, so b is a unit modulo q). Indeed $bG_g(P') - b'G_g(P) = -g_1\Delta$, so the image lies in Γ_q, and both lattices have index q. The image of $\mathcal{G}$ lies in a parallelogram Π of area $M^2q/(4\ell)$ inside $[-X_r, X_r]^2$, with one side of length $\asymp MV/\ell$ in a direction making an angle in $[\arctan\frac{1}{2}, \arctan 2]$ with the axes. Cover Π by $O(qS^2)$ squares of side $w_0 = M/V$. In a square $I_1 \times I_2$ away from the axes, Lemma 9.4 in each coordinate (modulus q, length $w_0 \geq X_r/(4S^2)$, since $VV_{\max} \leq 2YS^2$) gives at most $\sum_{r \bmod q}\big((w_0/q)\rho_r(\log N)^C\big)^2 = (w_0^2/q)\rho_r^2(\log N)^{2C}$ smooth points of Γ_q. Near the axes the filter $|n|, |n'| \geq X_r/256$ removes everything. So each pair has at most $(M^2/\ell)\rho_r^2(\log N)^C$ admissible g. There are at most $200S^2$ partners P' per P: they lie in the lattice $\{P' : \ell \mid \det(P, P')\}$, which is $\Lambda_\ell(g)$ for $g = (a, b)$ and contains $P \in \mathcal{P}_\ell$, so it is regular and Lemma 3.7 applies, and $\#\mathcal{P}_\ell \leq CV^2$. Dividing by $|\mathcal{G}| \asymp M^2$ gives $R_2 \leq CV^2S^2\ell^{-1}\rho_r^2(\log N)^C \leq (\log N)^C S^4\rho_r^2$. □

In the one-dimensional design $G = x - my$ the rational values (s, s') of a pair move on a *line* $b's - bs' = \Delta/\ell$ as m varies. Then divisor switching reduces to counting solutions of one linear equation in a box, which has no area main term, and we know no bound of the needed strength. In two dimensions the pair fills a lattice region of area much larger than its perimeter times the lattice scale, and that is all the argument needs.

10. Post-processing at fixed degree

This section adapts the number field sieve post-processing of Buhler–Lenstra–Pomerance [10] and Lee–Venkatesan [39, §6] to our rows. Because d is fixed, every field-dependent quantity is $N^{O(1)}$. This simplifies matters considerably: there are $O(\log N)$ character columns, at most one exceptional square class, and the square root needs no lattice reduction.

10.1. **Field-level events.**

Lemma 10.1 (irreducibility certificate)**.** *Let $g \in \mathcal{G}$ with $\gcd(g_1, N) = 1$, let w be a prime in $(W, 2W]$, and let N have no prime factor $\leq B$. Let $\mathcal{P}_{\mathrm{IR}}$ be the set of primes in $(r_0(d), \frac{1}{2}\log H]$,*

where $r_0(d) \geq 16$ is a constant. For c uniform on C_H, the probability that no $r \in \mathcal{P}_{\mathrm{IR}}$ has $f \bmod r$ of degree d and irreducible over $\mathbb{F}_r$ is at most

$$e\,(1 - 1/(4d))^{|\mathcal{P}_{\mathrm{IR}}|} \;\leq\; \exp\bigl(-\log H/(13d \log\log H)\bigr) \qquad (N \geq N_0).$$

If some r certifies, then f is irreducible over $\mathbb{Q}$.

Proof. Let $Q_{\mathrm{IR}} = \prod_{r \in \mathcal{P}_{\mathrm{IR}}} r = H^{1/2+o(1)} \leq (2H+1)/d$. Every residue class of C_H modulo Q_{IR} has between $\lfloor (2H+1)/Q_{\mathrm{IR}} \rfloor^d$ and $\lceil (2H+1)/Q_{\mathrm{IR}} \rceil^d$ elements, so the law of $c \bmod Q_{\mathrm{IR}}$ is at most $(1 + 1/d)^d \leq e$ times uniform. Under the uniform law the reductions modulo distinct r are independent (CRT). Fix $r \in \mathcal{P}_{\mathrm{IR}}$; note $r < B$, so $r \nmid Nw$. If $r \nmid g_1$, then $c \mapsto f \bmod r$ is an affine bijection from $\mathbb{F}_r^d$ onto the polynomials of degree at most d with $\phi(m_r) = Nwg_1^{-d} \neq 0$, where $m_r = g_0 g_1^{-1} \bmod r$: multiplication by $g_1 x - g_0$ maps polynomials of degree at most $d-1$ bijectively onto those vanishing at m_r. If $r \mid g_1$, then $g_1 x - g_0 \equiv -g_0$ is a unit, and $f \bmod r$ ranges over all polynomials with leading coefficient $a_d \not\equiv 0$ (as $a_d g_0^d \equiv Nw$). Monic irreducible polynomials of degree d over $\mathbb{F}_r$ with prescribed constant term $v \neq 0$ correspond to the elements of degree d of $\mathbb{F}_{r^d}$ with norm $(-1)^d v$. There are at least $((r^d - 1)/(r-1) - 2r^{d/2})/d \geq r^{d-1}/(2d)$ of them. After the translation $x \mapsto x + m_r$ and summing over the $r-1$ leading coefficients, the irreducible fraction is at least $(r-1)/(2dr) \geq 1/(4d)$ in the first case, and at least $1/(2d)$ in the second. Hence the bound, with $|\mathcal{P}_{\mathrm{IR}}| \geq \log H/(3 \log\log H)$. A factorisation of the primitive part of f over $\mathbb{Z}$ would reduce modulo r to a factorisation with positive degrees, since $r \nmid f_d$. □

Lemma 10.2 (rejection events)**.** *Under the raw law $\mathcal{A}$, let $\mathcal{R}$ be the event that $\gcd(g_1, N) = 1$ and one of the following holds: (i) no irreducibility certificate; (ii) $N \mid f'(m)$; (iii) the content of f equals w. Then $\Pr(\mathcal{R}) \leq \exp(-\log H/(13d \log\log H)) + 2/(2H+1) + 1/W = L^{-\omega(1)}$. In particular $\Pr(\mathcal{R}) \leq \gamma/2$ for large N, by Theorem 9.9.*

Proof. (i) is Lemma 10.1. (ii) Modulo N, $f'(m) \equiv \hat{f}'(m) + g_1 R_c(m, 1)$, because the derivative of $(g_1 x - g_0) R_c(x, 1)$ is $g_1 R_c(x,1) + (g_1 x - g_0)\partial_x R_c(x,1)$ and $g_1 m - g_0 \equiv 0$. Here $R_c(m,1) = \sum_i c_i m^{d-1-i}$ has coefficient 1 at c_{d-1} and g_1 is a unit. So for fixed $c_0, \ldots, c_{d-2}$ at most one $c_{d-1} \in [-H, H]$ gives $N \mid f'(m)$, since $2H + 1 < N$. (iii) If the content is w, then $w \mid f_d = a_d + g_1 c_0$, and g_1 is invertible modulo the prime $w > M$. For fixed $(g, w, c_1, \ldots, c_{d-1})$ this fixes c_0 modulo w, so the probability is at most $\lceil (2H+1)/w \rceil/(2H+1)$. Finally $\log H \asymp \log N$, so the bound in (i) is $\exp(-\Omega(\log N/\log\log N)) = L^{-\omega(1)}$. □

Remark 10.3. Rejecting a field with content w loses nothing: then every nonzero value $F_A(P)$ is divisible by the prime $w > B$, so $N_\ell(A) = 0$ for all ℓ. Also, the leading coefficient never causes a degree drop ($f_d \not\equiv 0 \bmod g_1$), and a non-unit leading coefficient is a factor exit because $0 < |f_d| < N$. So no restriction to unit leading coefficients is needed.

10.2. **Rows and the row identity.** Fix an accepted field: $\gcd(g_1 f_d w, N) = 1$, f is irreducible and primitive, and $f_d^{d-2} f'(m)$ is a unit modulo N. Let α be a root of f, $K = \mathbb{Q}(\alpha)$, $\beta := f_d \alpha$, and let $g(y) := f_d^{d-1} f(y/f_d) \in \mathbb{Z}[y]$ be the monic minimal polynomial of β. Put $D := f_d \operatorname{disc}(f)$. Since $\operatorname{disc}(g) = f_d^{(d-1)(d-2)} \operatorname{disc}(f)$, the primes dividing $\operatorname{disc}(g)$ divide D. Define the ring homomorphism

$$\psi : \mathbb{Z}[\beta] \to \mathbb{Z}/N, \qquad \beta \mapsto f_d m, \qquad m = g_0 g_1^{-1} \bmod N.$$

It is well defined because $g(f_d m) = f_d^{d-1} f(m) \equiv f_d^{d-1} g_1^{-d} N w \equiv 0$. For an atom $P = (a, b)$ put $z_P := f_d a - b\beta = f_d(a - b\alpha) \in \mathbb{Z}[\beta]$. Then

$$\mathrm{N}_{K/\mathbb{Q}}(z_P) = f_d^{d-1} F_A(P), \qquad \psi(z_P) = f_d G(P)/g_1.$$

Write $G(P) = \sigma_P \ell n_P$. Then $\ell \equiv \sigma_P g_1 f_d^{-1} \psi(z_P)/n_P \pmod N$. A row i comes from $R_i = r_i^2$ with primitive split $R_i \equiv u_0/v_0$, signs $e_{i\ell} = \pm 1$ and atoms P_ℓ, so

$$r_i^2 \equiv \mathfrak{q}_i \, f_d^{-E_i} \prod_\ell \psi(z_{P_\ell})^{e_{i\ell}} \pmod N,$$
$$\mathfrak{q}_i := \operatorname{sgn}(u_0 v_0) \prod_\ell \big(\sigma_{P_\ell} g_1 / n_{P_\ell}\big)^{e_{i\ell}} \in \mathbb{Q}^\times, \qquad E_i := \sum_\ell e_{i\ell}. \tag{10}$$

The factor $\operatorname{sgn}(u_0 v_0)$ is necessary: for example $R = -1$ can have shortest vector $(1, -1)$, with an empty product of atoms. The same atom can occur in several rows. For a set S of rows put $e_P(\mathsf{S}) := \sum_{i \in \mathsf{S}} e_{iP} \in \mathbb{Z}$ and $E(\mathsf{S}) := \sum_P e_P(\mathsf{S})$.

10.3. **Columns.** The matrix has one row per row of the algorithm and the following columns, all over $\mathbb{F}_2$:

(a) the sign of $\mathfrak{q}_i$;
(b) for each prime $p \le B$: $\sum_\ell e_{i\ell} v_p(n_{P_\ell})$;
(c) one parity column E_i;
(d) for each prime $r \le B$ with $r \nmid D$ and each root s of $f \bmod r$: $\sum_\ell e_{i\ell} \, [a_{P_\ell} \equiv b_{P_\ell} s \ (r)] \, v_r(F_A(P_\ell))$;
(e) t_χ character columns (Section 10.5).

The prime factors of g_1 are never needed: if (a)–(c) vanish on S, then $E(\mathsf{S})$ is even and $\mathfrak{q}_\mathsf{S} := \prod_{i \in \mathsf{S}} \mathfrak{q}_i = g_1^{E(\mathsf{S})} \prod_P n_P^{-e_P(\mathsf{S})}$ is a positive rational square, with the explicit square root $g_1^{E(\mathsf{S})/2} \prod_{p \le B} p^{-\frac{1}{2}\sum_P e_P(\mathsf{S}) v_p(n_P)}$. This parity column also makes $f_d^{-E(\mathsf{S})}$ a square. It is the device of [10, §12.7].

Lemma 10.4 (good primes)**.** *Let r be a prime with $r \nmid D$, and let (a, b) be coprime with $F_A(a, b) \neq 0$.*

(i) If $r \mid b$, then $r \nmid F_A(a, b)$, and $a - b\alpha$ is a unit at every prime of $\mathcal{O}_K$ above r.
(ii) If $r \nmid b$, the only prime $\mathfrak{p} \mid r$ with $\operatorname{ord}_\mathfrak{p}(a - b\alpha) \neq 0$ is the degree-one prime corresponding to the root $s \equiv ab^{-1}$ of $f \bmod r$; it occurs only if $r \mid F_A(a, b)$, and then $\operatorname{ord}_\mathfrak{p}(a - b\alpha) = v_r(F_A(a, b))$.

Consequently, if columns (d) vanish on S, then $\gamma_\mathsf{S} := \prod_P (a_P - b_P \alpha)^{e_P(\mathsf{S})}$ has even valuation at every prime of $\mathcal{O}_K$ not dividing D.

Proof. As $r \nmid f_d$, α is r-integral, and $\mathbb{Z}_{(r)}[\alpha]$ has discriminant $\operatorname{disc}(f) f_d^{-(2d-2)}$, an r-unit. So $\mathbb{Z}_{(r)}[\alpha]$ is the maximal order at r, and by Dedekind–Kummer the primes above r correspond to the irreducible factors ϕ_j of $f \bmod r$ and are unramified. (i) $F_A(a, b) \equiv f_d a^d \not\equiv 0$, and $a - b\alpha \equiv a$ modulo every $\mathfrak{p} \mid r$. (ii) $a - b\alpha \in \mathfrak{p}_j$ iff $\phi_j(ab^{-1}) \equiv 0$, i.e. iff $\phi_j = x - ab^{-1}$. Since $\mathrm{N}(a - b\alpha) = F_A(a, b)/f_d$ and $\mathfrak{p}$ has degree one, $\operatorname{ord}_\mathfrak{p}(a - b\alpha) = v_r(F_A(a, b))$. For the consequence: primes $r > B$ with $r \nmid D$ contribute nothing, since $F_A(P)$ is B-smooth. □

10.4. **The residual square-class group.** Let S_D be the set of primes of $\mathcal{O}_K$ above primes dividing D. Put $V_D := \{\gamma \in K^\times : \operatorname{ord}_\mathfrak{P} \gamma \text{ even for all } \mathfrak{P} \notin S_D\}$ and $\mathsf{H} := V_D/K^{\times 2}$. By Lemma 10.4, every S in the kernel of columns (a)–(d) has $\gamma_\mathsf{S} \in V_D$, and $\phi(\mathsf{S}) := [\gamma_\mathsf{S}] \in \mathsf{H}$ is $\mathbb{F}_2$-linear in S, because $\gamma_{\mathsf{S}+\mathsf{S}'} = \gamma_\mathsf{S} \gamma_{\mathsf{S}'} / \gamma_{\mathsf{S} \cap \mathsf{S}'}^2$.

Lemma 10.5 (size of H; cf. [10, Thm. 6.7])**.** *$|\mathsf{H}| \le 2^{d\,\omega(D) + d} h_K$, and $\log |\mathsf{H}| = O_d(\log N)$.*

Proof. Let $W_0 := \{\gamma : \gamma \mathcal{O}_K = \mathfrak{a}^2\}$ and $Y_0 := \mathcal{O}_K^\times K^{\times 2}$. Then $V_D \supset W_0 \supset Y_0 \supset K^{\times 2}$. The map $V_D \to \bigoplus_{\mathfrak{P} \in S_D} \mathbb{Z}/2$, $\gamma \mapsto (\operatorname{ord}_\mathfrak{P} \gamma \bmod 2)$, has kernel W_0. The map $W_0 \to \mathrm{Cl}(K)[2]$, $\gamma \mapsto [\mathfrak{a}]$, is well defined (the ideal group is torsion-free) with kernel Y_0. And $Y_0/K^{\times 2} \cong \mathcal{O}_K^\times / \mathcal{O}_K^{\times 2}$ has

order $2^{r_1+r_2} \le 2^d$, since the only roots of unity are ± 1 (d odd). Also $|S_D| \le d\,\omega(D)$. For the size: $\|f\|_\infty \le 3M(H+1) < N$ for large N (Lemma 3.3, the window floor and $\mu + h < 1$), and d is fixed, so $\log|D|$ and $\log|\operatorname{disc} g|$ are $O_d(\log N)$. Every ideal class contains an integral ideal of norm at most $\sqrt{|\Delta_K|} \le \sqrt{|\operatorname{disc} g|}$, and there are at most $X(1+\log X)^{d-1}$ integral ideals of norm at most X. So $\log h_K = O_d(\log N)$. □

10.5. **Characters.** For $0 \ne \mathfrak{q} \in \mathsf{H}$ let $L_\mathfrak{q} := K(\sqrt{\mathfrak{q}})$, of degree $2d$. Put $\bar\Delta := |\operatorname{disc} g|^2\, 8^d |D|^d$.

Lemma 10.6. *$|\Delta_{L_\mathfrak{q}}| \le \bar\Delta$, and the quadratic subfield F of $L_\mathfrak{q}$, if any, has $|\Delta_F| \le \bar\Delta^{1/d}$.*

Proof. $\Delta_L = \Delta_K^2\, \mathrm{N}_{K/\mathbb{Q}}(\mathfrak{d}_{L/K})$. A representative $\gamma \in V_D$ has even valuation at every $\mathfrak{P} \nmid 2D$, so L/K is unramified there. At $\mathfrak{P} \mid D$ with $\mathfrak{P} \nmid 2$, $\operatorname{ord}_\mathfrak{P} \mathfrak{d} \le 1$, and at $\mathfrak{P} \mid 2$, $\operatorname{ord}_\mathfrak{P} \mathfrak{d} \le 2\operatorname{ord}_\mathfrak{P}(2) + 1$. So $\mathrm{N}(\mathfrak{d}) \le 8^d |D|^d$, and $|\Delta_K| \le |\operatorname{disc} g|$. For $F \subset L$, $|\Delta_F|^{[L:F]}$ divides Δ_L, and $[L:F] = d$. □

Fact 10.7 (Stark [63, Lemma 8]; [39, Fact 6.22])**.** *Let L be a number field of degree n, and $c(L) = 4$ if $L/\mathbb{Q}$ is normal, $c(L) = 4 \cdot n!$ otherwise. If ζ_L has a real zero β_0 with $1 - (c(L)\log|\Delta_L|)^{-1} \le \beta_0 < 1$, then $\zeta_F(\beta_0) = 0$ for some quadratic field $F \subset L$.*

Fact 10.8 (Landau; see [16, Ch. 14], [39, Fact 6.24])**.** *There is an effective absolute constant $c_L > 0$ such that for distinct real primitive characters χ_1, χ_2 of conductors q_1, q_2, $L(s,\chi_1)L(s,\chi_2)$ has at most one real zero (with multiplicity) in $(1 - c_L/\log(q_1 q_2), 1)$.*

Fact 10.9 (Lagarias–Odlyzko [38, Thm. 1.3])**.** *Let L/K be Galois with group G, $n_L = [L:\mathbb{Q}]$, $d_L = |\Delta_L|$, and $\mathcal{C}$ a conjugacy class. If $x \ge \exp(10 n_L (\log d_L)^2)$, then*

$$\Big|\pi_\mathcal{C}(x) - \frac{|\mathcal{C}|}{|G|}\mathrm{Li}(x)\Big| \le \frac{|\mathcal{C}|}{|G|}\mathrm{Li}(x^{\beta_0}) + c_1 x \exp\big(-c_2 n_L^{-1/2}(\log x)^{1/2}\big),$$

with absolute effective c_1, c_2. The β_0 term is present only if ζ_L has a real zero $\beta_0 \in [1 - (4\log d_L)^{-1}, 1)$. Here $\pi_\mathcal{C}(x)$ counts primes $\mathfrak{p}$ of K unramified in L with $\mathrm{N}\,\mathfrak{p} \le x$ and Frobenius in $\mathcal{C}$.

Put $\nu_1 := \min\big(1/(4(2d)!\log\bar\Delta),\ c_L/(2\log\bar\Delta)\big)$, and let $\mathsf{E} := \{0 \ne \mathfrak{q} \in \mathsf{H} : \zeta_{L_\mathfrak{q}}$ has a real zero in $[1-\nu_1, 1)\}$.

Lemma 10.10 (at most one exceptional class)**.** *$|\mathsf{E}| \le 1$ for every field.*

Proof. Let $\mathfrak{q} \in \mathsf{E}$ with zero β_0. By Fact 10.7 (with $c = 4(2d)!$, valid in all cases) there is a quadratic $F_\mathfrak{q} \subset L_\mathfrak{q}$ with $\zeta_{F_\mathfrak{q}}(\beta_0) = 0$, and as $\zeta(\beta_0) \ne 0$ on $(0,1)$, $L(\beta_0, \chi_{F_\mathfrak{q}}) = 0$. Since $[L_\mathfrak{q}:\mathbb{Q}] = 2d$ with d odd, $L_\mathfrak{q}$ has at most one quadratic subfield (two would generate a biquadratic subfield, and $4 \nmid 2d$), and K has none. Hence $L_\mathfrak{q} = KF_\mathfrak{q}$. The map $\mathfrak{q} \mapsto L_\mathfrak{q}$ is injective by Kummer theory, so $\mathfrak{q} \mapsto F_\mathfrak{q}$ is injective. If $\mathfrak{q} \ne \mathfrak{q}'$ were both in E, then $\chi_{F_\mathfrak{q}}, \chi_{F_{\mathfrak{q}'}}$ would be distinct real primitive characters with conductors at most $\bar\Delta^{1/d}$ (Lemma 10.6), each with a real zero $\ge 1 - \nu_1 > 1 - c_L/\log(\bar\Delta^{2/d})$. That contradicts Fact 10.8. □

Sampler. Let x be the least integer with $\log(x/2) \ge \max(20d(\log\bar\Delta)^2, (\log N)^2)$, so $\log x = O_d((\log N)^2)$. Draw r uniform in $(x/2, x]$. If r is prime (AKS), find the set Z_r of roots of $f \bmod r$ (Las Vegas, e.g. Cantor–Zassenhaus [12]). With probability $|Z_r|/d$ output (r, s) with s uniform in Z_r; otherwise redraw. Each pair (r,s) is output with the same probability, so the output is uniform over the degree-one primes $\mathfrak{p} = (r, \alpha - s)$ with $\mathrm{N}\,\mathfrak{p} \in (x/2, x]$. Since $r > \bar\Delta \ge |D|$, $r > B$, and r exceeds every $|F_A(P)|$, no atom $a - b\alpha$ lies in $\mathfrak{p}$.

The expected number of draws is $O(d \log x)$. To see this, apply Fact 10.9 to K/K (trivial group, $n_L = d$). If ζ_K had a real zero in $[1 - (4 \cdot d!\log|\Delta_K|)^{-1}, 1)$, Fact 10.7 would produce a quadratic subfield of K, which does not exist. A zero in the wider interval of Fact 10.9

contributes at most $x^{1-1/(4d!\log\bar{\Delta})} = o(x/\log x)$ because $\log x \ge (\log N)^2$. So there are at least $x/(3\log x)$ prime ideals of K with norm in $(x/2, x]$; those of degree ≥ 2 number at most $d\sqrt{x}$.

Characters. A degree-one prime $\mathfrak{p} = (r, \alpha - s)$ with $r \nmid D$ defines a quadratic character on H: $\chi_\mathfrak{p}(\mathfrak{q}) :=$ the Legendre symbol modulo $\mathfrak{p}$ of the $\mathfrak{p}$-unit $\gamma\pi^{-\operatorname{ord}_\mathfrak{p}\gamma}$, for a representative $\gamma \in V_D$ and a uniformiser π at $\mathfrak{p}$. Since $\operatorname{ord}_\mathfrak{p}\gamma$ is even, this does not depend on π, and it is well defined on $V_D/K^{\times 2}$. On atoms, $\chi_\mathfrak{p}(a - b\alpha) = \big((a - bs)/r\big)$. The character columns are $\sum_\ell e_{i\ell}(1 - \chi_{\mathfrak{p}_j}(a_{P_\ell} - b_{P_\ell}\alpha))/2 \bmod 2$, for independent samples $\mathfrak{p}_1, \dots, \mathfrak{p}_{t_\chi}$.

Lemma 10.11 (characters separate)**.** *For every field and every $\mathfrak{q} \in \mathsf{H} \setminus (\mathsf{E} \cup \{0\})$, $\Pr_\mathfrak{p}[\chi_\mathfrak{p}(\mathfrak{q}) = -1] \ge 1/(5d)$ for $N \ge N_0(d)$. With $t_\chi := \lceil 5d(\ln|\mathsf{H}|_{\mathrm{bd}} + 2\ln\log N)\rceil = O_d(\log N)$, where $|\mathsf{H}|_{\mathrm{bd}} := 2^d|D|^d X_0(1+\log X_0)^{d-1}$ with $X_0 := \lceil\sqrt{|\operatorname{disc} g|}\,\rceil$ is a computable upper bound for $|\mathsf{H}|$ (proof of Lemma 10.5, using $2^{\omega(D)} \le |D|$), the probability that $\mathsf{H}_0 := \mathsf{H} \cap \bigcap_j \ker\chi_{\mathfrak{p}_j}$ is not contained in $\mathsf{E} \cup \{0\}$ is at most $(\log N)^{-2}$. Otherwise $|\mathsf{H}_0| \le 2$.*

Proof. Apply Fact 10.9 to $L_\mathfrak{q}/K$ with $G = C_2$, $\mathcal{C} = \{\sigma\}$, $n_L = 2d$, $d_L \le \bar{\Delta}$. For $\mathfrak{p} \nmid 2D$, $\chi_\mathfrak{p}(\mathfrak{q}) = -1$ iff the Frobenius of $\mathfrak{p}$ is σ. The main term contributes $\frac{1}{2}(\mathrm{Li}(x) - \mathrm{Li}(x/2)) \ge x/(4\log x)(1-o(1))$ inert primes of norm in $(x/2, x]$. A real zero, if present, is $< 1-\nu_1$ because $\mathfrak{q} \notin \mathsf{E}$, and contributes at most $2x^{1-\nu_1} \le 2x\exp(-\nu_1(\log N)^2) = o(x/\log x)$, since $\nu_1 \gg_d 1/\log N$. The error term is $2c_1 x\exp(-c_2\log N/\sqrt{2d}) = o(x/\log x)$, as $\log x = O((\log N)^2)$. Primes of degree ≥ 2 and ramified primes are negligible. There are at most $d(\pi(x) - \pi(x/2)) \le 0.6dx/\log x$ degree-one primes in all, so the ratio is at least $1/(5d)$. Independence and a union bound over H give the probability. Finally H_0 is a subgroup contained in $\mathsf{E} \cup \{0\}$, so $|\mathsf{H}_0| \le 2$. □

10.6. **Square roots.**

Algorithm 10.12 (SQ(Γ); after [10, §9])**.** Input: $\Gamma = g'(\beta)^2\prod_P z_P^{|e_P|} \in \mathbb{Z}[\beta]$, and the bound A_{sq} of Lemma 10.13.

1. Let p be the least prime in $(B, 2B]$ with $p \nmid D$, found by sieving $(B, 2B]$ once per field. Then $g \bmod p$ is squarefree, and Γ is a unit at p, because $\mathrm{N}(\Gamma) = \pm\operatorname{disc}(g)^2\prod(f_d^{d-1}F_A(P))^{|e_P|}$ has only prime factors $\le B$ or dividing D.
2. Factor $g \bmod p = \phi_1\cdots\phi_\tau$ over $\mathbb{F}_p$. In each $\mathbb{F}_p[y]/(\phi_j)$ test whether Γ is a square and compute a square root ρ_j (Tonelli–Shanks). If some component is a non-square, output NONSQUARE.
3. For each of the $2^{\tau-1}$ sign vectors ϵ with $\epsilon_1 = +1$: combine the $\epsilon_j\rho_j$ by the Chinese remainder theorem to $\rho_\epsilon \in \mathbb{Z}[\beta]/p$; Newton-lift the inverse square root from $z_0 = \rho_\epsilon^{-1}$ by $z \mapsto z + z(1 - \Gamma z^2)/2$ to precision $p^e > 2A_{\mathrm{sq}}$; put $\eta_\epsilon := \Gamma z \bmod p^e$ with centred coefficients. If $\eta_\epsilon^2 = \Gamma$ exactly, output η_ϵ.
4. Output NONSQUARE.

Lemma 10.13 (square roots)**.** *Let* $\mathrm{Mh} := \lfloor\|g\|_2\rfloor + 1$, $g_{\mathrm{d}} := d\|g\|_1\mathrm{Mh}^{d-1}$, $U := g_{\mathrm{d}}^2\prod_P(\|z_P\|_1\mathrm{Mh}^{d-1})^{|e_P|}$ *and* $A_{\mathrm{sq}} := d\lceil U^{1/2}\rceil 2^{d-1}\mathrm{Mh}\, g_{\mathrm{d}}^{d-1}$. *Then:*

(i) Any output η satisfies $\eta^2 = \Gamma$. If $\Gamma \in K^{\times 2}$, then $\Gamma = \eta^2$ with $\eta \in \mathbb{Z}[\beta]$, $\|\eta\|_\infty \le A_{\mathrm{sq}}$, and SQ outputs $\pm\eta$.

(ii) $\log A_{\mathrm{sq}} = O_d((1 + \sum_P|e_P|)\log N)$, and SQ runs in $O_d(\mathsf{M}(\log A_{\mathrm{sq}})\log\log A_{\mathrm{sq}}) + (\log N)^{O(1)}$ bit operations (Las Vegas), where M is the cost of integer multiplication.

Proof. (i) If $\Gamma = \delta^2$ with $\delta \in K$, then $(\delta/g'(\beta))^2 = \prod z_P^{|e_P|} \in \mathcal{O}_K$, so $\delta/g'(\beta) \in \mathcal{O}_K$. Moreover $g'(\beta)\mathcal{O}_K \subset \mathbb{Z}[\beta]$, because $\mathcal{O}_K \subset \mathcal{O}_K^\vee \subset \mathbb{Z}[\beta]^\vee = g'(\beta)^{-1}\mathbb{Z}[\beta]$ (Euler's lemma). So $\eta := \delta \in \mathbb{Z}[\beta]$. Since $p \nmid \operatorname{disc} g$, $\mathbb{Z}_p[\beta] \cong \prod_j \mathbb{Z}_p[y]/(G_j)$ with unramified local factors lifting the ϕ_j. The reduction of η is $\pm\rho_\epsilon$ for exactly one ϵ with $\epsilon_1 = +1$. A unit square root with prescribed

residue lifts uniquely (p odd), so $\eta_\epsilon \equiv \pm\eta \pmod{p^e}$, and centring recovers $\pm\eta$ since $p^e > 2A_{\mathrm{sq}}$. NONSQUARE at step 2 is correct, since η would reduce to a square in each residue field $\mathcal{O}_K/\mathfrak{P} = \mathbb{F}_p[y]/(\phi_j)$ (p is prime to the index). The coefficient bound: the conjugates satisfy $|\sigma_k(z)| \le \|z\|_1 \mathrm{Mh}^{d-1}$ and $|g'(\beta_k)| \le g_{\mathrm{d}}$, so $|\sigma_k(\eta)| \le U^{1/2}$. The coefficients of η are recovered with the inverse Vandermonde matrix, $(V^{-1})_{jk} = [y^j](g(y)/(y-\beta_k))/g'(\beta_k)$. Its numerators are the coefficients of $\prod_{j\neq k}(y-\beta_j)$, which are at most $2^{d-1}\prod_j \max(1,|\beta_j|) = 2^{d-1}\mathrm{M}(g) \le 2^{d-1}\mathrm{Mh}$ in absolute value (Mahler measure, $\mathrm{M}(g) \le \|g\|_2$), and $\prod_k |g'(\beta_k)| = |\operatorname{disc} g| \ge 1$ gives $|g'(\beta_k)| \ge g_{\mathrm{d}}^{-(d-1)}$. (ii) The bound on $\log A_{\mathrm{sq}}$ is immediate. The cost consists of a product tree for Γ, Newton iteration doubling the precision with $O(d^2)$ multiplications per step, and $2^{d-1} = O(1)$ sign vectors. □

10.7. **Linear algebra and the congruence.**

Algorithm 10.14 (POSTPROCESS(A, rows)).
1. Sample $\mathfrak{p}_1, \dots, \mathfrak{p}_{t_\chi}$ and build the $M_{\mathrm{rows}} \times C_{\mathrm{col}}$ matrix over $\mathbb{F}_2$ of columns (a)–(e).
2. Sample two uniform left-kernel vectors s_1, s_2 (Wiedemann's method over $\mathbb{F}_{2^e}$ with the Kaltofen–Saunders preconditioners [70, 35]). Verify them, and retry until they are linearly independent.
3. For $\mathsf{S} = s_1, s_2, s_1 + s_2$ in turn, run SQ on $\Gamma_{\mathsf{S}} := g'(\beta)^2 \prod_P z_P^{|e_P(\mathsf{S})|}$, and stop at the first S for which it returns η. If none does, go back to step 1 with fresh characters.
4. Return S and
$$\mathsf{Y}_{\mathsf{S}} := \sqrt{\mathfrak{q}_{\mathsf{S}}} \cdot f_d^{-E(\mathsf{S})/2} \cdot \psi(\eta) \cdot \big(\psi(\gamma_{\mathsf{S}}^-)\psi(g'(\beta))\big)^{-1} \bmod N, \qquad \gamma_{\mathsf{S}}^- := \prod_{e_P(\mathsf{S})<0} z_P^{|e_P(\mathsf{S})|}.$$

Proposition 10.15 (post-processing). *Let $N \ge N_0$ and fix an accepted field and $M_{\mathrm{rows}} \ge C_{\mathrm{col}} + 12$ rows. Algorithm 10.14 is Las Vegas. Its expected time is $M_{\mathrm{rows}}^2 (\log N)^{O_d(1)} + B^{1+o(1)}$, and it returns a nonempty set S of rows and $\mathsf{Y}_{\mathsf{S}} \in (\mathbb{Z}/N)^\times$ with $\prod_{i\in\mathsf{S}} r_i^2 \equiv \mathsf{Y}_{\mathsf{S}}^2 \pmod N$. Its output and all its intermediate quantities are functions of the field, the certificates $(R_i, u_0, v_0, \text{atoms})$ and its own coins.*

Proof. Kernel sampling. The matrix has at most $M_{\mathrm{rows}}(2k\log_2 N + t_\chi + 2)$ nonzero entries. Wiedemann's method with the Kaltofen–Saunders preconditioners, run over $\mathbb{F}_{2^e}$ with $2^e \ge M_{\mathrm{rows}}^4$, returns a uniformly random kernel element of a black-box singular system in expected $\tilde{O}(M_{\mathrm{rows}} \cdot \#\text{entries} + M_{\mathrm{rows}}^2)$ field operations [35, Thms. 2–4]. Here is how to sample uniformly from the binary kernel. Draw x uniform in $\mathbb{F}_2^{M_{\mathrm{rows}}}$, solve $A^\top y = A^\top x$ (padding $A^\top$ to a square matrix), and return the coordinate of $x - y$ at the basis element 1 of a fixed $\mathbb{F}_2$-basis of $\mathbb{F}_{2^e}$. Given $A^\top x$, the vector x is uniform on a coset of the kernel, and the solver sees only $A^\top x$; so the output is uniform on the binary kernel. Every output is verified. The kernel has dimension $r \ge 12$, so two samples are independent with probability $(1-2^{-r})(1-2^{1-r}) \ge 1 - 3\cdot 2^{-12}$.

One of three is a square. Columns (a)–(d) vanish on s_1, s_2, so $\phi(s_1), \phi(s_2) \in \mathsf{H}$. Columns (e) vanish too, so they lie in H_0. If $|\mathsf{H}_0| \le 2$ (probability $\ge 1 - (\log N)^{-2}$ by Lemma 10.11), either one of $\phi(s_1), \phi(s_2)$ is 0 or both equal the nonzero element, and then $\phi(s_1 + s_2) = 0$. In every case some $\mathsf{S} \in \{s_1, s_2, s_1 + s_2\}$ is nonempty with $\gamma_{\mathsf{S}} \in K^{\times 2}$. Then $\gamma'_{\mathsf{S}} := \prod_P z_P^{e_P(\mathsf{S})} = f_d^{E(\mathsf{S})}\gamma_{\mathsf{S}}$ is a square ($E(\mathsf{S})$ even), hence so is $\Gamma_{\mathsf{S}} = g'(\beta)^2\gamma'_{\mathsf{S}}(\gamma_{\mathsf{S}}^-)^2$, and SQ returns η by Lemma 10.13. A failed round has probability at most $(\log N)^{-2} + 3\cdot 2^{-12}$, so the expected number of rounds is $O(1)$.

The congruence. ψ is a ring homomorphism, so $\psi(\eta)^2 = \psi(g'(\beta))^2 \prod_P \psi(z_P)^{|e_P|}$. Hence
$$\mathsf{Y}_{\mathsf{S}}^2 = \mathfrak{q}_{\mathsf{S}} f_d^{-E(\mathsf{S})} \prod_P \psi(z_P)^{|e_P|} / \psi(\gamma_{\mathsf{S}}^-)^2 = \mathfrak{q}_{\mathsf{S}} f_d^{-E(\mathsf{S})} \prod_P \psi(z_P)^{e_P(\mathsf{S})} \equiv \prod_{i\in\mathsf{S}} r_i^2$$

by (10). Every quantity inverted is a unit modulo N. Indeed $\psi(z_P) = f_d G(P)/g_1$ with $G(P) = \pm \ell n_P$ coprime to N; $\psi(g'(\beta)) = f_d^{d-2} f'(m)$ by the field check; and $\sqrt{\mathsf{qs}}$ involves only g_1 and primes $\le B$.

Cost. Building the matrix and the characters costs $M_{\text{rows}}(\log N)^{O(1)}$, and kernel sampling $M_{\text{rows}}^2(\log N)^{O(1)}$. Each SQ call costs $O(\mathsf{M}(kM_{\text{rows}} \log N) \log\log N)$ by Lemma 10.13. The Hensel prime costs $\tilde{O}(B)$, and $\sqrt{\mathsf{qs}}$ costs $M_{\text{rows}}(\log N)^{O(1)}$. No square root modulo N is ever extracted. The last claim of the proposition holds by construction: no step reads a root r_i. □

Remark 10.16 (relation to the literature)**.** The obstruction-group chain is [10, Thm. 6.7], with the index $[\mathcal{O}_K : \mathbb{Z}[\alpha]]$ replaced by $2^{|S_D|}$. The random characters were suggested in [10, §8] and carried out rigorously in [39, §6]. The p-adic square root with exact verification is [10, §9], after [69]. Lee–Venkatesan's degree is $\asymp (\log N)^{1/3}$, which leads to $O(\log\log N)$ exceptional classes and an exceptional set of fields. At fixed degree, Lemma 10.10 gives at most one exceptional class for every field, and three tests suffice. We know of no explicit statement of this specialisation. Lenstra–Pomerance [44, §§3–4] use a similar explicit Hecke-character estimate with at most one exceptional zero, and random primes up to $\exp(c(\log|\Delta|)^2)$ to avoid GRH.

11. Batch smoothness

Lemma 11.1 (smooth parts; Bernstein [5, 6])**.** *Let $\mathcal{P}_B := \prod_{p \le B} p$, and let $x \ge 2$ and $2^e \ge \log_2 x$. The B-smooth part of x is $\gcd\big(x,\ (\mathcal{P}_B \bmod x)^{2^e} \bmod x\big)$. For given $x_1, \dots, x_n$, all their smooth parts can be computed in $O(\mathsf{M}(b) \log b)$ bit operations, where $b = \log_2 \mathcal{P}_B + \sum_i \log_2 x_i$.*

Proof. Write $x = s\,u$ with s smooth and $\gcd(u, \mathcal{P}_B) = 1$. Every $p^j \parallel s$ has $j \le \log_2 x \le 2^e$, so $s \mid \mathcal{P}_B^{2^e}$, while $\gcd(u, \mathcal{P}_B^{2^e}) = 1$. For the cost, form a product tree of the x_i, reduce $\mathcal{P}_B$ modulo the product, and descend a remainder tree; each level costs $O(\mathsf{M}(b))$ [68, §§10.1–10.3]. Then perform $e = O(\log\log x)$ modular squarings and a gcd for each x_i. □

Proposition 11.2 (the special-q search recognises $N_\ell(A)$ exactly)**.** SQ_SEARCH(A, ℓ) *is a deterministic function of (A, ℓ) and runs in time $\tilde{O}(B + S^2 \log N)$. If $\ell \nmid N$, then either it makes a factor exit, or it returns an atom if and only if $N_\ell(A) \ge 1$, and every atom it returns is a counted point.*

Proof. If $N_\ell(A) \ge 1$, the lattice is regular (Lemma 3.7), so step 1 does not fail. Step 2 enumerates at most $200S^2$ points from the reduced basis. It keeps exactly the points of $\mathcal{P}_\ell \cap \Lambda_\ell(g) \cap [-V, V]^2$ with $G(P) \ne 0$, $F_A(P) \ne 0$ and $|G(P)| \ge MV/32$. Every counted point is kept, and its gcd checks pass (its values are coprime to N). Step 3 computes all smooth parts exactly (Lemma 11.1). So step 4 returns a counted point whenever one exists, and only counted points. Lemma 11.1 applies with $b = O(B + S^2 \log N)$, since every value has $O(\log N)$ bits (Lemma 3.5). Trial division of the two chosen values costs $\tilde{O}(B)$. □

Remark 11.3. Since $B \le S^2$ when $\theta \le 1$, the batch costs $\tilde{O}(S^2)$ per box, the same as the enumeration. No smoothness test is probabilistic, and in particular no hyperelliptic curve method is used. (The published proof of the hyperelliptic curve method's running time [42, 43] defers details to a third part that has not appeared; Lee–Venkatesan's congruence theorem relies on it.) The Lenstra–Pomerance algorithm used in the shape test does not use the hyperelliptic method either; it uses the elliptic curve method on "recognisable" smooth numbers [44, §7].

12. Cost assembly and the constant

Proposition 12.1 (cost of one split)**.** *For $N \geq N_0$ as in Section 7, the expected running time of* SPLIT(N) *is at most*

$$L^{\max(a+1/\sqrt{2},\ a+\theta/(2a),\ 2a)+o(1)} = L^{(1+\theta)/\sqrt{2}+o(1)}.$$

Proof. By Lemma 10.2 and Theorem 9.9, the hypotheses of Proposition 7.3(i) hold, so $p \geq L^{-o(1)}$. We bound the remaining quantities in Proposition 7.3(ii).

- *Rows.* $M_{\text{rows}} \leq (d+1)\pi(B) + O_d(\log N) = L^{a+o(1)}$.
- *One trial.* SAMPLE_FIELD costs $(\log N)^{O(1)}$ in expectation. DESCEND_ONCE costs one Gauss reduction, two runs of the Lenstra–Pomerance algorithm on integers $\leq \sqrt{2N}$, and at most $2k$ special-q searches. The first has expected cost $L_{\sqrt{2N}}[1/2, 1+o(1)] = L^{1/\sqrt{2}+o(1)}$ (Fact 2.5), uniformly over the inputs. Each search costs $\tilde{O}(B + S^2 \log N) = L^{1/\sqrt{2}+o(1)}$ (Proposition 11.2; recall $S^2 = L^{\theta/(2a)+o(1)}$ and $\theta/(2a) = 1/\sqrt{2}$). AKS certification of the output primes is polynomial. So $C_{\text{trial}} = L^{\max(1/\sqrt{2},\theta/(2a))+o(1)}$.
- *Post-processing.* $\mathrm{T}_{\text{post}} = M_{\text{rows}}^2 (\log N)^{O(1)} + B^{1+o(1)} = L^{2a+o(1)}$ (Proposition 10.15).

Hence the expected time is at most $L^{o(1)}(L^{a+o(1)} L^{\max(1/\sqrt{2},\theta/(2a))} + L^{2a})$. With $a = \theta/\sqrt{2}$ the first two exponents both equal $(1+\theta)/\sqrt{2}$, and $2a = \sqrt{2}\theta < (1+\theta)/\sqrt{2}$ because $\theta < 1$. □

Remark 12.2 (the choice of a)**.** For a general smoothness exponent a (with $S^2 = L^{\theta/(2a)}$) the same proof gives $L^{c(a)+o(1)}$ with $c(a) = \max(2a, a + 1/\sqrt{2}, a + \theta/(2a))$. For $a > \theta/\sqrt{2}$ the term $a + 1/\sqrt{2}$ exceeds $(1+\theta)/\sqrt{2}$. On $(0, \theta/\sqrt{2})$ the function $a + \theta/(2a)$ is decreasing, because $\theta/\sqrt{2} < \sqrt{\theta/2}$ for $\theta < 1$. So $a = \theta/\sqrt{2}$ is optimal, and the minimum is $(1+\theta)/\sqrt{2}$.

13. Uniformity, small factors and prime powers; proof of the main theorem

Proof of Theorem 1.1. Correctness. Every prime in the output is certified by AKS. Every proper factor returned by SPLIT, or by a factor exit, divides the current integer. FACTOR recurses on both cofactors and stops at certified primes. So the output is the complete factorisation, and the algorithm is Las Vegas: it never errs, only its running time is random.

Preprocessing. Step 1 of FACTOR (primality, perfect powers) costs $(\log N)^{O(1)}$. Trial division to $B(N)$ costs $L^{a+o(1)}$. The cofactor N' handed to SPLIT is odd, has no prime factor $\leq B(N)$, and is neither a prime nor a perfect power, so $\omega(N') \geq 2$. Since $B(N') \leq B(N)$, N' has no prime factor $\leq B(N')$, and all statements of Sections 5–12 apply to N' with its own parameters. A non-squarefree N' (for example $p^2 q$) needs no special treatment: Theorem 6.1 uses $|\boldsymbol{\mu}_2| = 2^\omega$, and the other sections use only $\gcd(N', t) = 1$ and $N' \mid F_A(g_0, g_1)$.

Small N'. Step 3 of FACTOR runs SPLIT and trial division in parallel, so its cost is at most twice the smaller of their costs. For $N' < N_0$, trial division bounds it by $O(\sqrt{N_0}) = O(1)$. For $N' \geq N_0$, Proposition 12.1 bounds it by $L_{N'}^{(1+\theta)/\sqrt{2}+o(1)} \leq L_N^{(1+\theta)/\sqrt{2}+o(1)}$. Here $N_0 = N_0(d, k, \mu, \eta_0)$ is effective provided the constants η_D, c_D, δ_D of Fact A.2 are (Remark A.17); it involves the effective constants of Facts 10.7–10.9 and of the smooth-number facts, but it need not be known to the algorithm.

Recursion. Every call of SPLIT produces a proper factor, and the factorisation tree has fewer than $\log_2 N$ internal nodes. So there are at most $\log_2 N$ calls of SPLIT and of trial division, each on an integer $\leq N$, plus polynomial bookkeeping. The bound of Proposition 12.1 is uniform in the input: write its bound as $L_n^{c+\epsilon(n)}$ with $c = (1+\theta)/\sqrt{2}$ and $\epsilon(n) \to 0$. For fixed $\epsilon > 0$ there is n_ϵ with $\epsilon(n) \leq \epsilon$ for $n \geq n_\epsilon$. So every call on an $N' \leq N$ costs at most $C_\epsilon + L_N^{c+\epsilon}$, since $L_{N'} \leq L_N$.

Summing over at most $\log_2 N$ calls, the expected total time is at most $\log_2 N\,(C_\epsilon + L_N^{c+\epsilon})$ for every $\epsilon > 0$, that is, $L_N^{c+o(1)}$.

Uniformity. All $L^{o(1)}$ factors are uniform in N: $(\log N)^{O(k)}$, $\binom{2k}{k}$, $2^{O(\omega(N))}$ (Lemma 5.4), $\gamma^{-O(k)}$ (Theorem 9.9), and the Wiedemann and Hensel polylogarithms. No exceptional set of N arises. The analytic inputs need only $\gcd(N, t) = 1$, the absence of prime factors $\le B$, and the coprimality statements of Lemma 3.8. They need no squarefreeness and no condition on the residue classes of the prime factors of N. □

Remark 13.1 (expected time versus probability)**.** Every loop of the algorithm is Las Vegas with verified outputs: primes are certified, kernel vectors checked, $\eta^2 = \Gamma$ tested, and gcds checked. So only expected costs have to be bounded. The Lenstra–Pomerance algorithm is used only through its expected-time bound, so no cap on its running time is needed.

14. The recursive shape test

Status. Theorem 14.6 and Corollary 14.8 are conditional on the proof of Theorem 1.1. They add no new analytic input. They were audited internally by three blind referees, all of whom returned "accept with fixes", with no fatal or major finding. They are not peer reviewed. Note that they raise the stakes on Theorem 1.1: if it is correct, so is a rigorous $L_N[1/2, o(1)]$.

The constant $1/\sqrt{2}$ in Theorem 1.1 enters at exactly one place: the Lenstra–Pomerance factorisation of the two $\sqrt{N}$-size halves u_0, v_0 in the shape test (Proposition 12.1). In this section that subroutine is replaced by a recursive call of the algorithm itself.

14.1. **The algorithm and what is used from the main proof.** Fix admissible (d, k, μ, η_0) with $\theta < 1$, and fix any $a > 0$. Let $\mathrm{FACTOR}_{\mathrm{rec}}[a]$ be Algorithm 3.10 with two changes:

(i) $B := \lfloor L^a \rfloor$ and $S := \lceil L^{\theta/(4a)} \rceil$, so that $S^2 \rho_r \rho_F = L^{o(1)}$ as before;
(ii) in step 3 of DESCEND_ONCE, $|u_0|$ and $|v_0|$ are factored by $\mathrm{FACTOR}_{\mathrm{rec}}[a]$ itself, with fresh coins. Its output is returned in canonical form: the sorted list of primes with multiplicities. The loop over $\ell \mid u_0 v_0$ in step 4 runs in increasing order.

For a procedure O that, given $m \ge 1$, always returns the canonical factorisation of m, terminates almost surely, and uses coins independent of everything else, write SPLIT^O for SPLIT with shape test O. Call such an O *admissible*.

Lemma 14.1 (the main proof at a general a)**.** *For every fixed $a > 0$, Sections 5–11 hold verbatim with $B = \lfloor L^a \rfloor$ and $S = \lceil L^{\theta/(4a)} \rceil$, with N_0 depending also on a. Consequently, for preprocessed $N \ge N_0$ and every admissible O:*

(A1) the expected number τ of calls of DESCEND_ONCE in $\mathrm{SPLIT}^O(N)$ is at most $L_N^{a+o(1)}$, and each call makes at most two calls of O, on integers at most $\sqrt{2N}$;
(A2) the expected work of $\mathrm{SPLIT}^O(N)$ outside the calls of O is at most $L_N^{\max(2a,\, a+\theta/(2a))+o(1)}$;
(A3) $\mathrm{SPLIT}^O(N)$ returns a proper factor almost surely.

Proof. The value $a = \theta/\sqrt{2}$ is used only in the cost optimisation (Proposition 12.1). The one size relation tied to it, $B < 2S^2$ in Lemma 3.1(a), is used only to bound the cost of batch smoothness. That cost is $\tilde{O}(B + S^2)$, which is $L^{\max(a,\theta/(2a))+o(1)}$ for any a. Every other estimate is stated for an arbitrary fixed a. This covers Lemma 2.4, the constants C_0 and R_0 of Section 9, the grid $\mathcal{X}$, Lemma 8.4 (Appendix A), and Lemma 5.4. Then (A1)–(A3) follow from the proof of Proposition 7.3, with $M_{\mathrm{rows}} = L^{a+o(1)}$. In that proof, charge at most $2k$ special-q searches to every call of DESCEND_ONCE. The harmonic-mean identity applies as before, and the post-processing costs $L^{2a+o(1)}$. □

Lemma 14.2 (oracle substitution; "Lemma OS"). *For all admissible O, the joint law of the outer transcript of* $\mathrm{SPLIT}^O(N)$ *is the same. The outer transcript consists of the outer coins and every value computed outside O, but not running times. In particular τ, the work outside O, and the output have the same law for every admissible O.*

Proof. The outer computation is a deterministic function of the outer coins and of the values returned by O. Each call of O on m returns the canonical factorisation of m, which is a function of m alone. No step of the algorithm depends on a running time. Induction over the steps of the outer computation shows that every outer value is a function of the outer coins only, and this function does not depend on O. □

Lemma 14.3 (Wald for adaptive sub-calls). *Suppose a randomised procedure makes sub-calls $i = 1, 2, \dots$. Suppose the decision to make call i, and its input, are measurable with respect to the history $\mathcal{F}_{i-1}$. Suppose call i uses fresh coins and has cost C_i with $\mathbb{E}[C_i \mid \mathcal{F}_{i-1}, \mathit{input}] \le \bar{F}$. Then $\mathbb{E}[\sum_{i \le \tau} C_i] \le \bar{F}\,\mathbb{E}[\tau]$.*

Proof. $\mathbb{E}\sum_i C_i 1_{\tau \ge i} = \sum_i \mathbb{E}\big[1_{\tau \ge i}\mathbb{E}[C_i \mid \mathcal{F}_{i-1}, \mathrm{input}]\big] \le \bar{F}\sum_i \Pr[\tau \ge i]$, by monotone convergence. □

Lemma 14.4 (transcript invariance under recursion; "Lemma $\mathrm{T}_{\mathrm{rec}}$"). *The hypothesis of Theorem 6.1 holds for* SPLIT^O *with* $O = \mathrm{FACTOR}_{\mathrm{rec}}[a]$.

Proof. The inputs $|u_0|, |v_0|$ of the inner calls are functions of R, and their coins are fresh. Each inner call's whole transcript is a function of its input and its coins. This includes the inner call's own hidden roots, fields and dependencies, and its inner gcds, which are taken with $|u_0|$, $|v_0|$ and their divisors, not with N. So the outer transcript together with all inner transcripts is a function of $(R_i)_i$ and of coins independent of the outer roots $(r_i)_i$. The inner transcripts are not independent of the outer squares, but they are invariant under $r_i \mapsto \epsilon r_i$ on each outer fibre, which is all Theorem 6.1 needs. Its proof then applies verbatim. □

14.2. **Solving the recursion.** Write $L_x := \exp\big(\sqrt{\log x \cdot \max(1, \log\log x)}\big)$ for all $x \ge 2$.

Lemma 14.5 ("Lemma F"). *Let $F : [2, \infty) \to [0, \infty)$ be nondecreasing and finite, and suppose that for $x \ge x_0$,*

$$F(x) \le L_x^{p+\epsilon(x)} + L_x^{\delta+\epsilon(x)} F(\sqrt{2x}), \qquad \epsilon(x) \to 0.$$

Then $F(x) \le L_x^{c+o(1)}$ with $c = \max\big(p,\ \delta/(1 - 1/\sqrt{2})\big) = \max(p, (2+\sqrt{2})\delta)$.

Proof. Scaling. For $y = \sqrt{2x}$, $\log y \log\log y \le \frac{1}{2}(\log x + \log 2)\log\log x$, so $L_y^b \le e\, L_x^{b/\sqrt{2}}$ for $x \ge x_1(b)$. *Margin.* Fix $\eta > 0$ and put $\eta' := \eta(1 - 1/\sqrt{2})/4$. Choose $X \ge \max(x_0, x_1(c+\eta), 4)$ such that for $x \ge X$: $\epsilon(x) \le \eta'$, $L_x^{\eta - \eta'} \ge 2$ and $L_x^{3\eta'} \ge 2e$. Put $K := \max(1, F(X))$. Then $F(x) \le K \le K L_x^{c+\eta}$ on $[2, X]$. *Induction.* Suppose $F \le K L_x^{c+\eta}$ on $[2, Y]$ with $Y \ge X$, and let $x \in (Y, Y^2/2]$. Then $\sqrt{2x} \le Y$, and

$$F(x) \le L_x^{p+\eta'} + K\, e\, L_x^{\delta+\eta'+(c+\eta)/\sqrt{2}}.$$

Since $\delta \le c(1 - 1/\sqrt{2})$, the second exponent is at most $c + \eta/\sqrt{2} + \eta' = c + \eta - 3\eta'$, so that term is at most $(K/2)L_x^{c+\eta}$. The first term is at most $L_x^{c+\eta}/2$. As $Y^2/2 > Y$, the bound holds for all x, and η was arbitrary. □

Theorem 14.6 (recursive shape test; "Theorem RA"). *For admissible (d, k, μ, η_0) with $\theta < 1$ and fixed $a > 0$,* $\mathrm{FACTOR}_{\mathrm{rec}}[a]$ *is a Las Vegas algorithm that outputs the complete factorisation of every $N \ge 2$ in expected time*

$$L_N\big[1/2,\ c(a) + o(1)\big], \qquad c(a) := \max\big(a + \theta/(2a),\ (2+\sqrt{2})a\big).$$

The minimum over a is attained at $a^* = \sqrt{(\sqrt{2}-1)\theta/2}$, with value $c^* = \sqrt{(1+\sqrt{2})\theta}$.

Proof. Let $F(x) := \sup_{2 \le n \le x} \mathbb{E}[\text{cost of FACTOR}_{\text{rec}}[a](n)]$.

Finiteness. We argue by strong induction on n. For $n < N_0$ the cost is $O(1)$, because FACTOR alternates SPLIT with trial division (step 3). For $n \ge N_0$ the inequality below calls F only at $\sqrt{2n} < n$.

Recursion. The factorisation tree of N has fewer than $2\log_2 N$ splitting nodes, plus $O(\log N)$ perfect-power steps. Each node costs $L^{a+o(1)}$ for preprocessing, plus one race between the core SPLIT^O, with $O = \text{FACTOR}_{\text{rec}}[a]$, and trial division. The race costs at most twice the core. Inside the strong induction, O is applied only to inputs at most $\sqrt{2n} < n$, on which it is already known to be admissible (exact, canonical output, almost surely terminating; nested calls share no mutable state). So Lemma 14.2 applies to the core, and (A1)–(A3) hold for this O. Every call of O is on an integer at most $\sqrt{2N}$ with fresh coins, so its conditional expected cost is at most $F(\sqrt{2N})$. By Lemma 14.3,

$$F(x) \le 2\log_2 x \Big[L_x^{\max(2a,\, a+\theta/(2a))+o(1)} + 2L_x^{a+o(1)} F(\sqrt{2x}) \Big].$$

Lemma 14.5 with $p = \max(2a, a+\theta/(2a))$ and $\delta = a$ gives $c = \max(2a, a+\theta/(2a), (2+\sqrt{2})a) = c(a)$.

Correctness. Every output prime is certified by AKS. Termination is almost sure because the expected time is finite. Fruitfulness holds by Lemma 14.4.

Optimisation. $a + \theta/(2a)$ is decreasing on $(0, \sqrt{\theta/2}]$ and $(2+\sqrt{2})a$ is increasing. They meet where $(1+\sqrt{2})a = \theta/(2a)$, that is at $a^* < \sqrt{\theta/2}$. Then $c^{*2} = (2+\sqrt{2})^2(\sqrt{2}-1)\theta/2 = (1+\sqrt{2})\theta$. □

Remark 14.7.

(a) $c^* < 1$ iff $(1+\theta)^2 < 2$, the same threshold as Theorem 1.1, and $c^* < (1+\theta)/\sqrt{2}$ iff $\theta < \sqrt{2}-1$. So recursion helps exactly when Theorem 1.1 already beats Lenstra–Pomerance. At $d = 9$, $k = 13$, $c^* = 0.99521$ (against 0.99720).

(b) The infima over (k, η_0) are $c^* = \sqrt{3(1+\sqrt{2})/(d+1)}$: 0.9515 ($d = 7$), 0.8510 ($d = 9$), 0.7769 ($d = 11$), 0.7193 ($d = 13$), 0.6018 ($d = 19$). They tend to 0 as $d \to \infty$.

(c) *What binds.* After recursion the binding pair is the recursive shape test and the special-q search. Linear algebra and batch smoothness are slack by the factor $2a^*/c^* = 2-\sqrt{2}$.

(d) *Fixed-depth recursion.* With Lenstra–Pomerance at the bottom, the constants are $c_{j+1} = (\theta/c_j + c_j)/\sqrt{2}$ with $c_0 = 1$, so $c_1 = (1+\theta)/\sqrt{2}$ is Theorem 1.1. At $d = 9$ (with $\theta \to 3/10$) the sequence is $1, 0.9192, 0.8808, 0.8637, \cdots \to 0.8510$.

(e) *Sanity check against heuristics and folklore.* The heuristic fixed-degree number field sieve (balanced non-monic pair, [10, §12]) has constant $2/\sqrt{d+1}$, which is 0.632 at $d = 9$. The infimum in (b) is exactly $\sqrt{3(1+\sqrt{2})/4} \approx 1.346$ times larger, for every d. So Theorem 14.6 claims less than the heuristic picture, and $L_N[1/2, o(1)]$ is weaker than the heuristic $L_N[1/3]$. The folklore "constant 1" is not a theorem. It is the random-squares ledger $\max(a + \xi/(2a), 2a)$ for relations built from random-like integers of size N^ξ, evaluated at $\xi = 1/2$. That ledger is sharp for *independent uniform* integers [15, 2], but it does not bind here. The $\sqrt{N}$-size halves are never required to be smooth: they must be products of k primes of a fixed size, an event of probability $L^{-o(1)}$ for fixed k, and they are then factored completely. Only the descent values, of total size N^θ, enter the matrix. Applied to those values the ledger gives the heuristic floor $\sqrt{2\theta}$. The constant $\sqrt{(1+\sqrt{2})\theta}$ lies above it, because the recursive shape test costs $L^{a+c/\sqrt{2}}$. This

evasion is legitimate only to the extent that the relation-supply estimates (Theorems 8.7 and 9.7) are correct.

(f) *Prior art.* We found no rigorous analysis that calls a factoring algorithm recursively on auxiliary $\sqrt{N}$-size numbers, and no fixed-point equation for the L-constant. The nearest precedent is Pomerance's use of his own rigorous factoring algorithm on the auxiliary $p-1$ inside a discrete-logarithm algorithm [54, pp. 138–139]. That is a single call, not a recursion. The recursive variant was first computed in this project, and GPT-6.1-sol re-derived it independently.

14.3. **A portfolio: $L_N[1/2, o(1)]$ and a genuine $+o(1)$.**

Corollary 14.8 (portfolio; "Corollary P")**.** *There is a single Las Vegas algorithm that outputs the complete factorisation of every $N \geq 2$ in expected time $L_N[1/2, o(1)]$, that is, $\exp\big(o(\sqrt{\log N \log\log N})\big)$. No effective rate for the $o(1)$ is supplied. Likewise, running the same schedule over k at fixed $d = 9$, without recursion, gives a single algorithm with expected time $L_N[1/2, 0.9192\ldots + o(1)]$.*

Proof. Let $(A_j)_{j\geq 1}$ be the algorithms $\mathrm{FACTOR}_{\mathrm{rec}}[a_j]$ for a computable sequence of admissible parameters with $c(a_j) \to 0$. For example, take $d_j = 2j+7$, $k_j = \lceil d_j^2/4\rceil + j$, μ_j the (rational) optimum of Lemma 3.2, $\eta_{0,j} := 1/d_j^2$, and a_j rational with $|a_j - a_j^*| \to 0$. These tuples are admissible for every j (the inequalities are monotone in j; checked exactly for $j \leq 60$, where $c^* = 0.239$ at $j = 60$). The slack must tend to 0, since $\theta \geq \eta_0$. The map $j \mapsto A_j$ is computable. No threshold $N_0(j)$ enters the algorithms, because each A_j races its core against trial division at every node (step 3 of Algorithm 3.10); so $N_0(j)$ need not be computable. Let $T_j(N)$ be the running time of A_j on N.

Restart schedule. In phase $s = 1, 2, \ldots$, run A_i for each $i \leq s$, from scratch with fresh coins, for at most 2^{s-i} steps. Stop at the first certified output. Phase s costs at most 2^s steps, up to a polylogarithmic bookkeeping factor.

Bound. Fix j and put $s_0 := j + \lceil \log_2(2\mathbb{E}T_j(N))\rceil$. In each phase $s \geq s_0$, the run of A_j fails to finish within $2^{s-j} \geq 2^{s-s_0}\cdot 2\mathbb{E}T_j$ steps with probability at most $2^{-(s-s_0)-1}$, by Markov. Runs in different phases are independent. So the expected time is at most a polylogarithmic factor times

$$2^{s_0+1} + \sum_{m\geq 1} 2^{s_0+m} 2^{-m(m+1)/2} = O\big(2^j\, \mathbb{E}T_j(N)\big).$$

Hence $\limsup_N \log \mathbb{E}T(N)/\mathcal{L}_N \leq c(a_j)$ for every j. The second statement is the same argument over k at $d = 9$, with Theorem 1.1 in place of Theorem 14.6. Since j is fixed in each step, no uniformity in j is needed; for the same reason no effective rate is supplied. □

Corollary 14.8 answers Remark 1.3(a): the infimum 0.9192 at $d = 9$ becomes a genuine $+o(1)$ for one algorithm, through scheduling rather than through uniform estimates in k. It is *not* an $L_N[\alpha]$ bound for any $\alpha < 1/2$.

15. Application: discrete logarithms in $\mathbb{F}_p^*$

Status. Theorem 15.1 is a reduction to the relation supply of Theorem 1.1: Sections 5, 7–9 and 11 with $N := p$. The DL-specific part is proved below and needs no further hypothesis. Its components have prior art (Remark 15.11). It was reviewed by GPT-6.1-sol (one blind and one adversarial round) and by two Claude critics, with all fixes applied. It has not had an independent end-to-end audit, and it is not peer reviewed.

15.1. **Statement and literature.** Let p be prime, $n := p - 1$, g a primitive root and $h \in \mathbb{F}_p^*$. We want $X = \log_g h \bmod n$.

Theorem 15.1 (discrete logarithms; "Theorem DL"). *Fix admissible (d, k, μ, η_0) with $\theta < 1$. There is a Las Vegas algorithm which, given a prime p, a primitive root g and $h \in \mathbb{F}_p^*$, computes $\log_g h \bmod (p-1)$ in expected time $L_p[1/2, c_{\mathrm{DL}}(\theta) + o(1)]$, where*

$$c_{\mathrm{DL}}(\theta) := \min_{a>0} \max\bigl(a + 1/\sqrt{2},\ a + \theta/(2a),\ 3a\bigr) = \begin{cases} (1+\theta)/\sqrt{2}, & \theta \le 1/2, \\ 3\sqrt{\theta}/2, & 1/2 < \theta < 1. \end{cases}$$

No factorisation of $p - 1$, no unit or class group computation, no Schirokauer map and no power-residue character is used.

So c_{DL} equals the factoring constant whenever $\theta \le 1/2$, which covers every tuple with $(1 + \theta)/\sqrt{2} < 1$. It is 0.99720 at $d = 9$, $k = 13$ (exactly $55\sqrt{2}/78$, plus $\eta_0/\sqrt{2}$), and its infimum is $13\sqrt{2}/20 = 0.91924$ at $d = 9$ and $1/\sqrt{2}$ over all d. It is below $\sqrt{2}$ iff $\theta < 8/9$, which includes every optimised tuple; at $d = 3$ the cubic term binds and $c_{\mathrm{DL}} = 3\sqrt{3}/4 \approx 1.299$.

Corollary 15.2 (arbitrary base). *For $g, h \in \mathbb{F}_p^*$ with $h \in \langle g \rangle$, $\log_g h \bmod \mathrm{ord}(g)$ can be computed in the same expected time.*

Proof. Draw random auxiliary bases t. A primitive root occurs with probability $\varphi(n)/n \ge 1/(1 + \log_2 n)$. Run $\mathrm{DLOG}(p, t, g)$ and $\mathrm{DLOG}(p, t, h)$ under a budget that doubles at each stage, with $c(1 + \log_2 n)$ fresh t per stage, where c is a constant fixed so that a stage whose budget exceeds 4 times the expected time for a primitive root fails with probability at most $1/4$ (Markov). Every output is verified ($t^u = g$, $t^v = h$). Then solve $ux \equiv v \pmod n$, discard the result if $\gcd(u, n) \nmid v$, and check $g^x = h$ before output. The doubling series converges. □

Literature. The best rigorous bound for discrete logarithms in $\mathbb{F}_p^*$, for all primes p, that we found is still Pomerance's [54]: $L_p[1/2, \sqrt{2} + o(1)]$ for precomputation and $L_p[1/2, 1/\sqrt{2} + o(1)]$ per logarithm. We found nothing below $\sqrt{2}$ for all primes or for almost all primes, unconditionally or under GRH. For comparison:

- Finite fields $\mathbb{F}_{p^n}$ with small characteristic have stronger rigorous results: Bender–Pomerance [3] and Kleinjung–Wesolowski's quasi-polynomial algorithm for fixed characteristic [37].
- Class groups of imaginary quadratic fields have $L[1/2, \sqrt{2}]$ under ERH (Hafner–McCurley [26]). Under GRH the constant $3/\sqrt{8}$ is proved by Biasse and coauthors [7] for computing the group structure, not for discrete logarithms. Vollmer's earlier claim of that constant (ANTS IV, 2000 [66]) was retracted in a corrigendum in his ANTS V paper [67], as recorded by our literature reviewer (we did not re-check the corrigendum).
- Enge and Gaudry [22] give a general framework with hidden exponents.
- The ERH-rigorous number-field toolkit of de Boer, Pellet-Mary and Wesolowski [17] explicitly leaves finite-field discrete logarithms aside.

This is a search result, not a certification.

15.2. **The algorithm and the transfer of the relation supply.** The parameters are those of Section 3.1 with $N := p$, except that $a > 0$ is a free fixed constant (Lemma 14.1). For a field A write $m := g_0 g_1^{-1} \bmod p$ and $\psi(a - b\alpha) := a - bm$.

Algorithm 15.3 ($\mathrm{DLOG}(p, g, h)$). 1. Select a field A exactly as in step 1 of Algorithm 3.11, with targets drawn as in step 2 below.

2. *Round.* Repeat until $\bar{C} + d$ rows are held, where $\bar{C} := (d+1)\pi(B) + 1$: draw (a_i, b_i) uniform in $(\mathbb{Z}/n)^2$, put $T_i := g^{a_i} h^{b_i}$, and run DESCEND_ONCE(A) with $R := T_i$ (the Gauss split $T_i \equiv u_0/v_0$, the shape test, and a special-q atom for each large prime). Store the row together with its label (a_i, b_i).
3. Run EXACT_DEPENDENCE (Proposition 15.9). It returns a primitive $x^* \in \mathbb{Z}^M$ and $\lambda = \log_g \prod_i T_i^{x_i^*}$.
4. Put $A_* := \sum_i x_i^* a_i$ and $B_* := \sum_i x_i^* b_i \bmod n$. If $\gcd(B_*, n) = 1$, put $X := (\lambda - A_*) B_*^{-1} \bmod n$. If $g^X = h$, return X. Otherwise start a new round in the same field, with fresh rows.

Each row satisfies, with $G(P_\ell) = \sigma_\ell \ell n_\ell$ and $G(P) \equiv g_1 \psi(a - b\alpha) \pmod p$,

$$T \equiv \varepsilon \prod_{q \le B} q^{e_q} \, g_1^E \prod_\ell \psi(a_\ell - b_\ell \alpha)^{e_\ell} \pmod p, \qquad \varepsilon = \pm 1, \quad E = \sum_\ell e_\ell, \tag{11}$$

where $e_\ell = \pm 1$ according as $\ell \mid u_0$ or $\ell \mid v_0$, and $e_q = v_q(u_0) - v_q(v_0) - \sum_\ell e_\ell v_q(n_\ell)$.

Lemma 15.4 (transfer)**.** *With $N := p$, the following hold for every $p \ge p_0(d, k, \mu, \eta_0, a)$:*

(i) Lemmas 3.3–3.8, Proposition 11.2, Theorems 8.7, 9.7 and 9.9, and Lemma 10.2 hold verbatim. No factor exit can occur, since every gcd-checked quantity is nonzero and smaller than p.

(ii) Theorem 5.2 holds with $h_2 = 1$ for T uniform on $\mathbb{F}_p^*$: $\Pr[\exists u, v \in \mathcal{S} : u \equiv Tv] \ge |\mathcal{S}|^2/(c_k(p-1))$. Consequently Theorem 7.2 holds without the factor $2^{-\omega}$, and Proposition 7.3 holds with $p_F = 0$.

(iii) (Hidden labels.) T_i is uniform on $\mathbb{F}_p^*$, and given T_i the label b_i is uniform on $\mathbb{Z}/n$. Within a round, every accepted row and every computation other than the final A_*, B_* is a function of the T's and of coins independent of the labels. So, conditional on the round's transcript $\mathcal{T}_j$ (all T's, all coins, and all labels and outcomes of earlier rounds), the b_i of round j are i.i.d. uniform.

Proof. (i) The prime N enters those sections only through $\gcd(N, t) = 1$ for $0 < t \le 2MV < N$, through $N \mid F(g_0, g_1)$, and through preprocessing, which is not needed ($\omega = 1$). By Lemma 3.5 and Lemma 3.1, $|G(P)|$, $|F_A(P)| \le X^*$ and $0 < |f_d|$ are all less than p. (ii) In the proof of Theorem 5.2, $\sum_T n(T) = |\mathcal{S}|^2$ exactly, because every pair determines $T = uv^{-1}$ and there is no square-class condition. The second moment bound and Cauchy–Schwarz are unchanged. (iii) $(a, b) \mapsto (\log_g T, b) = (a + Xb, b)$ is a bijection of $(\mathbb{Z}/n)^2$. □

15.3. **Exact dependencies.** Fix an accepted field: f irreducible with content 1, $0 < |f_d| < p$, and m a simple root of $f \bmod p$ (the rejection $p \mid f'(m)$). Let $K = \mathbb{Q}(\alpha)$, $t := f_d$, $\mathfrak{a} := \mathcal{O}_K + \alpha \mathcal{O}_K$, and for a primitive atom $P = (a, b)$ let $\mathfrak{J}_P := (a - b\alpha)\mathfrak{a}^{-1}$.

Lemma 15.5 (ideal bookkeeping)**.** *(i) $\mathfrak{J}_P$ is integral, with $\mathrm{N}\,\mathfrak{J}_P = |F(a, b)|$. (ii) If $r \nmid t \operatorname{disc} f$, the primes dividing $\mathfrak{J}_P$ above r are of degree one, namely $(r, \alpha - ab^{-1})$, with valuation $v_r(F(a, b))$. (iii) For each of the $O(d \log p)$ primes $r \le B$ dividing $t \operatorname{disc} f$, an r-maximal order, the primes above r and the valuations of $\mathfrak{J}_P$ can be computed in time $\mathrm{poly}(d, \log r, \log p)$.*

Proof. (i) $a - b\alpha \in \mathfrak{a}$. Since f is primitive, Gauss's content lemma gives $\mathrm{N}\,\mathfrak{a} = |t|^{-1}$, and $|\,\mathrm{N}(a - b\alpha)| = |F(a, b)|/|t|$. (ii) For $r \nmid t$, $\mathfrak{a}$ is trivial at r; for $r \nmid t \operatorname{disc} f$, argue as in Lemma 10.4. (iii) This is the standard local computation (Round 2 at a known prime, decomposition, valuations; [13, §§4.8, 6.1–6.2], [9]). The bad $r \le B$ are found by reducing $t \operatorname{disc} f$ modulo each $r \le B$, in time $\tilde{O}(B)$. □

Columns. There is one column for each prime $q \le B$ (entry e_q), one column for g_1 (entry E), and one column for each prime ideal of $\mathcal{O}_K$ above some $r \le B$ that occurs (entry $\sum_\ell e_\ell v(\mathfrak{J}_{P_\ell})$). There are at most $\bar{C}$ columns. Let $E \in \mathbb{Z}^{M\times C}$ be the exact exponent matrix, $\Lambda_1 := \{x \in \mathbb{Z}^M : x^\top E = 0\}$, and $\gamma_x := \prod_P (a_P - b_P\alpha)^{y_P(x)}$, where $y_P(x) := \sum_i x_i \sum_{\ell: P_\ell = P} e_{i\ell}$.

Lemma 15.6. *(i) For $x \in \Lambda_1$, $\gamma_x \in \mathcal{O}_K^\times$.*

(ii) Let $\mathfrak{p}$ be the degree-one prime above p with $\alpha \equiv m \pmod{\mathfrak{p}}$, and $\mathrm{red}_\mathfrak{p} : \mathcal{O}_{K,\mathfrak{p}}^\times \to \mathbb{F}_p^\times$. For $x \in \Lambda_1$, $W(x) := \prod_i T_i^{x_i} \equiv \pm \mathrm{red}_\mathfrak{p}(\gamma_x) \pmod p$.

(iii) $\Lambda_2 := \{x \in \Lambda_1 : \gamma_x \in \mu_K\}$ is saturated in $\mathbb{Z}^M$, and $\operatorname{rank}\Lambda_2 \ge M - C - (r_1 + r_2 - 1) \ge 1$ when $M \ge C + d$. Every primitive vector of Λ_2 is primitive in $\mathbb{Z}^M$.

(iv) For $x \in \Lambda_2$, $W(x) = \pm 1$ (d odd, so $\mu_K = \{\pm 1\}$), and $\lambda := \log_g W(x) \in \{0, n/2\}$ is read off.

Proof. (i) $(\gamma_x) = \prod_P \mathfrak{J}_P^{y_P} \mathfrak{a}^{\sum_P y_P}$. The g_1-column gives $\sum_P y_P = 0$, and the ideal columns give valuation 0 at every prime above $r \le B$. Every prime dividing some $\mathfrak{J}_P$ lies above some $r \le B$, because $\mathrm{N}\,\mathfrak{J}_P = |F(P)|$ is B-smooth. (ii) Each $a - b\alpha$ is a $\mathfrak{p}$-unit ($|F(P)| < p$ and $p \nmid t$), and $\mathrm{red}_\mathfrak{p}(a - b\alpha) = \psi(a - b\alpha)$. Raise (11) to the power x_i and multiply: the rational and g_1 factors cancel because $x^\top E = 0$, and the signs give ± 1. (iii) $u : x \mapsto [\gamma_x] \in \mathcal{O}_K^\times/\mu_K \cong \mathbb{Z}^{r_1+r_2-1}$ is a homomorphism with kernel Λ_2; if $jx \in \Lambda_2$ then $\gamma_x^j \in \mu_K$, so $\gamma_x \in \mu_K$. (iv) By (ii), $\mathrm{red}_\mathfrak{p}$ maps μ_K into $\{\pm 1\}$. □

Lemma 15.7 (unit elimination)**.** *Given $s := r_1 + r_2 \le d$ independent $x_1, \dots, x_s \in \Lambda_1$ with entries of at most β_x bits, a nonzero $\mathbf{n} \in \mathbb{Z}^s$ with $\sum_j n_j x_j \in \Lambda_2$ can be found in time $\tilde{O}_d\big(M(\beta_x + \log p) + (\beta_x + \log p)^2\big)$.*

Proof. Let $\mathrm{Log} : K^\times \to \mathbb{R}^{r_1+r_2}$ be the usual logarithmic embedding, and $v_j := \mathrm{Log}(\gamma_{x_j}) = \sum_P y_P(x_j)\mathrm{Log}(a_P - b_P\alpha)$. These vectors lie in the unit lattice Λ_U, of rank $r_U := r_1 + r_2 - 1$, and $\mathrm{Log}\, z = 0$ iff $z \in \mu_K$ (Kronecker).

Gap. By Blanksby–Montgomery [8], a non-torsion algebraic integer of degree $d' \le d$ has Mahler measure $> 1 + 1/(52d' \log 6d')$. For a unit z of K, $\log \mathrm{M}(z) \le d\|\mathrm{Log}\, z\|_\infty$, so every nonzero $w \in \Lambda_U$ has $\|w\|_\infty \ge \beta_d := \log(1 + 1/(52d \log 6d))/d$.

Small relations exist. $\mathrm{Rel} := \{\mathbf{n} : \sum n_j v_j = 0\}$ has rank at least 1. In an LLL-reduced basis of Λ_U the Gram–Schmidt norms are at least $2^{-r_U/2}\beta_d$, so the integer coordinates of v_j are at most $2^{O(d^2)}\|v_j\|/\beta_d$. Cramer's rule then gives a nonzero relation of norm at most $R := \max\big(1, d!(2^{O(d^2)} \max_j \|v_j\|/\beta_d)^d\big)$. Each $|\sigma_i(a - b\alpha)|$ lies between $p^{-O(d)}$ and $p^{O(1)}$ (the lower bound from $|\mathrm{N}(a - b\alpha)| \ge 1/|t|$), so $\|\mathrm{Log}(a - b\alpha)\| = O_d(\log p)$, $\|v_j\| \le 2^{\beta_x} M \cdot 2k \cdot O_d(\log p)$, and $\log R = O_d(\beta_x + \log p)$.

LLL. Let $b_j := (e_j, \lfloor 2^\kappa v_j \rceil)$, with each coordinate within 1 of $2^\kappa v_j$. If $\mathbf{n} \in \mathrm{Rel}$, then $\|\sum n_j b_j\| \le (1 + s\sqrt{d})\|\mathbf{n}\|$. If $\mathbf{n} \notin \mathrm{Rel}$, then $\|\sum n_j b_j\| \ge 2^\kappa \beta_d - s\sqrt{d}\|\mathbf{n}\|$. LLL returns a first vector of norm at most $R' := 2^{(s-1)/2}(1 + s\sqrt{d})R$. Choosing $2^\kappa \beta_d > (1 + s\sqrt{d})R'$, its coefficient vector (of norm at most R') is a nonzero relation. So $\kappa = O_d(\beta_x + \log p)$. The roots of f are needed to $\kappa + O(\beta_x + \log p)$ bits, the logarithms to the same precision, and the sums over at most $2kM$ atoms cost $\tilde{O}(M(\kappa + \beta_x))$. LLL in dimension at most $2d$ with $O(\kappa)$-bit entries costs $\tilde{O}(\kappa^2)$ [50]. *Certification.* Accept $\mathbf{n}$ only if $\|\sum n_j v_j\|_\infty < \beta_d/2$ is certified by interval arithmetic. A true relation passes, and by the gap a non-relation cannot. □

Lemma 15.8 (hidden exponents)**.** *Let $x^* \in \mathbb{Z}^M$ be primitive and a function of the round's transcript $\mathcal{T}_j$. Then $B_* = \sum_i x_i^* b_i$ is uniform on $\mathbb{Z}/n$ given $\mathcal{T}_j$, so $\Pr[\gcd(B_*, n) = 1 \mid \mathcal{T}_j] = \varphi(n)/n \ge 1/(1 + \log_2 n)$. If $\gcd(B_*, n) = 1$, then $X = (\lambda - A_*)B_*^{-1} \bmod n$, exactly for every prime power of n at once.*

Proof. By Lemma 15.4(iii) the b_i are i.i.d. uniform given $\mathcal{T}_j$, and $b \mapsto b \cdot x^* \bmod n$ is a surjective homomorphism (a Bezout vector exists), so the image of the uniform law is uniform. Also $\log_g W(x^*) = \sum_i x_i^*(a_i + Xb_i) = A_* + XB_*$. The bound on $\varphi(n)/n$ is elementary. □

Proposition 15.9 ("Proposition XK")**.** *The following procedure EXACT_DEPENDENCE is Las Vegas. On $M \geq \bar{C} + d$ rows with $M = L^{a+o(1)}$ it outputs a primitive $x^* \in \Lambda_2$ and $\lambda = \log_g W(x^*)$ in expected time $L^{3a+o(1)}$.*

1. *Compute E exactly (Lemma 15.5).*
2. *Pick a random prime $q \in [Q, 2Q]$ with $Q := M(\log p)^3$, and compute by Gaussian elimination modulo q a rank profile (I, J) of size ρ.*
3. *For d rows $i \notin I$, solve $y^\top E_{I,J} = -E_{i,J}$ over $\mathbb{Q}$ by p-adic lifting [20], and clear denominators: $x_i := D(y, e_i)$.*
4. *Verify $x_i^\top E = 0$ on all columns. If this fails, return to step 2.*
5. *Apply Lemma 15.7 to s of the x_i. Put $x^* := \sum n_j x_j / \gcd$.*
6. *Compute $W(x^*)$ and λ (Lemma 15.6(iv)).*

Proof. Correctness. This follows from Lemmas 15.6 and 15.7. The verified x_i are independent through their free coordinates e_i, which is all Lemma 15.7 needs; step 4 does not certify that ρ is the rational rank, and nothing requires it. *Step 2.* A fixed maximal nonsingular minor of E has determinant at most $(O(k \log p))^{1.5\rho}$, since every row has at most $2k(1 + \log_2 p)$ nonzero entries of size $O(k \log p)$. So at most $O(\rho \log \log p)$ primes divide it, while $[Q, 2Q]$ contains $\gg \rho \log p$ primes. With probability at least $1/2$ the minor survives, and then verification passes. So the expected number of restarts is at most 2, and each costs $O(MC\rho) = L^{3a+o(1)}$ operations modulo q. *Step 3.* One inverse modulo a prime costs ρ^3, and $O(\rho \log(\rho \|E\|))$ lifting steps cost $O(\rho^2)$ each, for d right-hand sides. In total this is $\tilde{O}(\rho^3) = L^{3a+o(1)}$, and the outputs have $\beta_x = O(\rho \log \log p) = L^{a+o(1)}$ bits (Hadamard). *Steps 4–6.* These cost $L^{2a+o(1)}$ (Lemma 15.7 with this β_x, and M modular exponentiations with $L^{a+o(1)}$-bit exponents). □

Proof of Theorem 15.1. By Lemma 15.4 and Proposition 7.3, with $p_F = 0$ and the per-row cost of Proposition 12.1 at a general a (Lemma 14.1), the rows of a round cost $L^{a+\max(1/\sqrt{2}, \theta/(2a))+o(1)}$. Post-processing costs $L^{3a+o(1)}$ (Proposition 15.9). By Lemma 15.8 each round succeeds with conditional probability at least $1/(1 + \log_2 n)$, independently of earlier rounds, so the expected number of rounds is $L^{o(1)}$. Wald's identity gives the total $L^{\max(a+1/\sqrt{2},\, a+\theta/(2a),\, 3a)+o(1)}$. At $a = \theta/\sqrt{2}$ the first two terms are $(1+\theta)/\sqrt{2}$, and $3a \leq (1+\theta)/\sqrt{2}$ iff $\theta \leq 1/2$; no a does better, by the factoring optimisation. For $\theta > 1/2$ the minimum is at $a = \sqrt{\theta}/2$, where $a + \theta/(2a) = 3a = 3\sqrt{\theta}/2 \geq a + 1/\sqrt{2}$. □

Remark 15.10 (interaction with recursion)**.** The shape test in DLOG is integer factorisation of $|u_0|, |v_0| \leq \sqrt{2p}$. Logarithms cannot be computed recursively, because the logarithms of the halves live in the same group $\mathbb{F}_p^*$, but the factoring subroutine can be replaced. With a factoring constant τ for that subroutine,

$$c_{\mathrm{DL}} = \min_a \max\bigl(a + \tau/\sqrt{2},\ a + \theta/(2a),\ 3a\bigr) = \begin{cases} s + \theta/(2s), & s := \tau/\sqrt{2} \geq \sqrt{\theta}, \\ 3\sqrt{\theta}/2, & s < \sqrt{\theta}. \end{cases}$$

With the recursive constant $\tau = c^* = \sqrt{(1+\sqrt{2})\theta}$ of Theorem 14.6, $s \geq \sqrt{\theta}$ and $c_{\mathrm{DL}} = c^*$. With the portfolio of Corollary 14.8 as the subroutine ($\tau \to 0$), the cubic post-processing term $3a$ takes over and $c_{\mathrm{DL}} = 3\sqrt{\theta}/2$. This is 0.822 as an infimum at $d = 9$, and it tends to 0 as d grows, so a portfolio over d gives $L_p[1/2, o(1)]$ for discrete logarithms as well. We record these as consequences only; they were not separately reviewed.

Remark 15.11 (prior art and what does not transfer)**.** The unit elimination by Log, LLL and a minimum-length bound is the classical route to unit relations in class group and regulator computations (see [13] and, for a modern rigorous treatment, [17, Part 3]). Explicit units in place of Schirokauer maps are known when units are computable [58, Ex. 2.6]. Hidden exponents are classical [54, 22]. What we did not find elsewhere is the combination that computes no unit or class group at all, uses one exact dependency with exact cancellation of the ideals $(a-b\alpha)\mathfrak{a}^{-1}$ and removes units only on a d-dimensional image, has a quantified L^{3a} cost (so the constants of Theorem 1.1 survive), and is applied to descended uniform targets.

The square-class machinery of Section 10 does not extend uniformly to ℓ-th powers for large $\ell \mid p-1$. The Kummer extensions have degree growing with ℓ, each character value would need an order-ℓ discrete logarithm to be linearised, and Schirokauer maps need a Leopoldt/Wieferich-type condition that is not known for a prescribed ℓ. Exact dependencies avoid all of this, because they never form a modular kernel.

Toy check. The DL pipeline ran on 36 pre-registered instances with 30–50-bit primes ($d = 3$). It returned the correct logarithm in 36/36 cases, with no row, kernel or torsion violation. All 3465 certified torsion dependencies had $W = \pm 1$. None of the 707 non-torsion kernel vectors with nontrivial algebraic part passed the auxiliary torsion test, which shows that the unit step is necessary. The hidden-label law held for $\ell = 2, 3, 5, 7$. One pre-registered descriptive prediction (full unit rank in at least 9/12 runs) failed at 5/12, because the toy's pair box was too small. This is a design shortfall, not a violation. The code is in `src/dlog/`.

16. A toy implementation

We implemented the whole pipeline at toy scale as a correctness check. It is not evidence for any asymptotic claim. The code is in `src/thmA/` with tests in `tests/test_thmA.py`; the run was pre-registered (commit `9c67cc6`) before it was made.

What is implemented. The pipeline implements the following, and receives only N, a public configuration and a seed.

- The random field with (g_0, g_1) random and $F(g_0, g_1) = Nw$ for a random prime w.
- Dixon descent of uniform squares: Gauss split, shape test, and a special-q atom for every large prime, found with batch smoothness per box.
- The row identity (10) including the split sign.
- The columns of Section 10 with one parity column, first-degree primes $r \nmid f_d \operatorname{disc} f$, and 40 quadratic characters.
- A binary kernel, the Hensel square root in $\mathbb{Z}[\beta]$, the congruence, and the final gcd.

Deviations from the theorem (all disclosed in the pre-registration).

- $d = 3$, $\mu = 1/3$, $H = 4$, and the floors on h are not imposed.
- w is a random prime in $(2M, 4M]$ rather than in $(W, 2W]$. So the algebra of $F(g_0, g_1) = Nw$ is tested, but the analytic role of Lemma 9.3 is not.
- The split shape is relaxed: each half is a smooth part times at most two primes in $(B, 4B]$, with at least one large prime per row. The theorem's shape (at most k primes near $N^{1/(2k)}$) is far too rare at 30–50 bits.
- The shape test uses Pollard–Brent rather than Lenstra–Pomerance, and primality uses deterministic Miller–Rabin rather than AKS.
- The kernel is computed by Gaussian elimination.

Results. Run E1 used 48 instances: 20 balanced semiprimes, 12 three-prime integers, 4 unbalanced, 4 small prime times semiprime, 4 of the form p^2q, and 4 powers or primes, at 30–50 bits. Run E1b used 12 further instances in a diagnostic mode that continues after accidental factors.

registered prediction	result
correct complete factorisation (E1)	48/48 (E1b: 12/12)
row identity (10)	0 violations
congruence $\mathsf{X}_\mathsf{S}^2 \equiv \mathsf{Y}_\mathsf{S}^2$	0 violations
tested kernel vectors that are squares in K	352/352
Theorem 6.1, $\omega = 2$ (predicted 1/2)	$137/288 = 0.476$ ($z = -0.82$)
Theorem 6.1, $\omega = 3$ (predicted 3/4)	$49/64 = 0.766$ ($z = +0.29$)
accidental exits within the null bound	$32 \le 118$
preprocessing (small factors, powers, primes)	12/12

Descriptively, the special-q success rate was 0.06–0.13, and 35–260 fields were tried per kept field. Each split used 83–743 rows and took 0.15–3.1 s. The toy confirms the algebra and the hidden-root symmetry end to end. It cannot test relation supply (Theorems 8.7 and 9.7), which is where the content of the theorem lies, and its running times say nothing about the constant c.

The recursive variant. A second pre-registered toy run (commit `ea5953b`; code in `src/recurse/`) ran 22 instances twice: once with Pollard–Brent as the shape test, and once with the pipeline itself as the shape test (recursion depth 1, inner inputs of 25–28 bits).

- Both runs produced correct complete factorisations in 22/22 instances.
- The outer transcripts were bit-identical in 22/22 instances, as Lemma 14.2 predicts.
- There were no row-identity violations.
- The fruitful fraction for $\omega = 2$ was $2390/4936 = 0.484$ ($z = -2.22$, inside the pre-registered $|z| \le 3$).

Again, this checks code and the exact Lemma 14.2 only.

17. Limitations and outlook

17.1. **Limitations.**

- *Not refereed.* The proof was produced and checked with AI assistance only (see the Acknowledgements). An earlier version of Lemma 8.4 imported a bound from inside the proof of Lee–Venkatesan's Lemma 8.5; reviewers named it the weakest link, and it has been replaced by the self-contained proof of Appendix A. The one remaining external input there, Drappeau's Lemme 5, is published, but its proof refers to an unpublished preprint of Harper. The second moment (Theorem 9.7) is the most novel analytic component, and the field-selection bookkeeping (Proposition 7.3) is the most delicate probabilistic one. Sections 14 and 15 received internal AI audits only.
- *Infimum versus a single algorithm.* At fixed d, Theorem 1.1 gives $(1+\theta)/\sqrt{2} + o(1)$ for each fixed k; the infimum over k is not attained by any one of these algorithms. Corollary 14.8 obtains it by a restart schedule. A uniform version, with $k = k(N) \to \infty$ and effective rates, would need control of the k-dependence of every constant.
- *Effectivity and practicality.* The constants (Stark, Landau, Lagarias–Odlyzko) are effective but astronomically large. The effectivity of the constants of Drappeau's Lemme 5 is not stated in the literature; this affects only the threshold N_0, not Theorem 1.1. No effective rate is supplied for Corollary 14.8. The algorithm is a vehicle for a rigorous bound, not a practical method; the heuristic number field sieve is far faster.
- *Other problems.* For discrete logarithms in $\mathbb{F}_p^*$ the same devices give Theorem 15.1. Whether they improve rigorous bounds for class groups, or for discrete logarithms in other finite fields, is open.

17.2. **Outlook: growing degree.** Recursion (Section 14) removes every cost barrier that comes from testing smoothness or shape. The argument of Lemma 14.5 works at any level $L[\alpha]$, with scaling factor $2^{-\alpha}$ in place of $1/\sqrt{2}$. A shape test on $\sqrt{N}$-size halves then costs $L_N[\alpha, 2^{-\alpha}c]$, and intermediate smoothness tests on values $\exp((\log N)^\gamma)$, $\gamma < 1$, cost $L_N[\alpha]^{o(1)}$. In particular the barrier that rigorous hyperelliptic smoothness tests impose on one-level descent ($L[4/9]$ at a splitting bound $L[2/3]$) disappears. We nevertheless do *not* obtain a rigorous $L_N[\alpha]$ with $\alpha < 1/2$. What binds is the geometry of the descent and the precision of the relation-supply estimates. The following are heuristic targets in a size model, not theorems.

- *One-level descent,* 2/5. With d and k growing ($d \asymp (\log N/\log\log N)^{1/5}$, $k \asymp (\log N/\log\log N)^{2/5}$), split acceptance requires $k(\log\log N + \log(1/\gamma)) = O((\log N)^\alpha(\log\log N)^{1-\alpha})$, and the design floors force d and k to grow. Together these give $\alpha \geq 2/5$, provided the first and second moments hold with polylogarithmic precision uniform in (d, k).
- 3/7. If the provable moments lose $(\log N)^{O(d)}$, the floor becomes $\alpha \geq (2+\tau)/(5+2\tau)$ with $\tau = 1$, that is 3/7. This is the realistic first target.
- *Multi-level descent,* 1/3. With $O(\log\log N)$ levels of degree-one descent in one field of degree $\asymp (\log N/\log\log N)^{1/3}$, whose sizes contract by constant factors, the size model reaches $L[1/3]$. This is the heuristic picture of number field sieve discrete-logarithm descent. A rigorous version needs a *nested descendability* statement: lower bounds for polynomial values whose prime factors lie in a field-dependent set of relative density $(\log N)^{-O(1)}$. We know of no tool for this. The nearest are Lenstra–Pomerance's lower bound for integers in an interval composed of primes from an arbitrary set [44, Thm. 6.1], and the ERH-conditional ideal sampling of de Boer, Pellet-Mary and Wesolowski [17].

Below $L[1/2]$ trial division no longer makes $2^{\omega(N)}$ negligible. This one obstruction can be removed at once: restricting the split count to coprime pairs makes each square hit at most once, which removes the factor $2^{-\omega}$ of Theorems 5.1–5.2 ("Lemma S″" in the project notes). The decisive missing input is an explicit-error version of the first and second moments, and of Lemma 8.4 at $Y = L[3/5]$, uniform in (d, k). Two further pieces need repair: the post-processing at growing degree, where the factorial in Stark's threshold must be absorbed by the character-prime size, and a slack $\eta_0 = \eta_0(N)$ that shrinks with N.

Acknowledgements

This work was carried out with heavy AI assistance. We describe the process so that readers can weigh it.

- *The human author* set the research question and direction and decided which threads to pursue.
- *An orchestrating Claude session* (Anthropic, Claude Opus 5.5) planned the work and reviewed results.
- *Claude Opus 5.5 subagents* produced the proofs, the code, the toy experiments and the project notes on which this paper is based. They worked under written rules (exact arithmetic, no access to the factors in operational code, status labels, and pre-registration of experiments). This paper was drafted by such a subagent from those notes.
- *GPT-6.1-sol* (OpenAI, through the Codex CLI with web search) was used for blind attempts, literature searches and adversarial review. In two places it found the key device independently: the two-dimensional rational randomisation with divisor switching, and the random multiplier of the resultant.

Each component passed at least one GPT-6.1-sol referee and one independent Claude critic. The assembled theorem was audited by two blind GPT-6.1-sol referees and three blind Claude critics, all of whom returned "accept with fixes" with no fatal or major finding. Reviews found and corrected real errors, among them:

- a claimed shape-test cost of $L^{o(1)}$ (actually $L_N[1/2, 1/\sqrt{2}]$), which changed the constant from $\sqrt{2\theta}$ to $(1+\theta)/\sqrt{2}$;
- a second-moment comparison that failed for candidates with small rational value, repaired by the filter $|G(P)| \geq MV/32$;
- a missing sign in the row identity;
- an inference about field selection that ignored factor exits;
- a polylogarithmic lower bound on γ that is false for a computable S;
- the conflation of an infimum over k with an $o(1)$.

No human has checked the proofs line by line. AI review is not peer review.

Appendix A. Proof of Lemma 8.4

This appendix gives a self-contained proof of Lemma 8.4. It does not use Lee–Venkatesan's Lemma 8.5, neither its statement nor its proof. It also does not use Harper's preprint [27] directly.

- For moduli $t = \ell s$ with ℓ a prime $> B$, the part of the error coming from characters principal at ℓ is *exactly* the error modulo s divided by $\ell - 1$ (Lemma A.8). So we need a restricted Bombieri–Vinogradov estimate over smooth moduli with a relative saving, which we prove as Lemma A.16.
- Characters of large conductor are handled by the classical multiplicative large sieve and a Perron bound (Lemma A.14).
- Characters of small conductor are handled by one published lemma of Drappeau (Fact A.2).

The proof follows the project note `notes/lem85_results.md`, which was checked by one GPT-6.1-sol and two Claude referees.

Honest statement of the one external input. Fact A.2 is a published, numbered statement: Drappeau [21, Lemme 5], restated by Pascadi [51, Lemma 3.2]. In both sources, however, its proof consists of a reference to the computations of §3.3 of Harper's preprint [27], which has not been published. Drappeau writes "il découle des calculs de la section 3.3 de [Har12b]". Our argument never enters the preprint, but the preprint stands behind that published statement. Removing it would mean re-proving Harper's zero-density and zero-free-region analysis of $\Psi(x, y; \chi)$, with Page's theorem, and we have not done that. We apply Fact A.2 with $A = 1/2$ and discard the factor $(\log x)^{-A}$. So Siegel's theorem is never needed, and only the effective part of Harper's computation is involved.

A.1. **Setting and statement.** Throughout, $y := B = \lfloor L^a \rfloor$, and $S(z, y)$ is the set of y-smooth $n \leq z$. We write $\Psi(z, y) = \#S(z, y)$, $\varrho(z, y) = \Psi(z, y)/z$ and $\Psi(z, y; \chi) = \sum_{n \in S(z,y)} \chi(n)$. For $(a, q) = 1$ put $E(z, y; q, a) := \Psi(z, y; q, a) - \Psi_q(z, y)/\varphi(q)$, and let $\Delta(q) := \varphi(q)^{-1} \sum_{\chi \neq \chi_0 \bmod q} |\Psi(z, y; \chi)|$. Let $H(u) := \exp(u/\log^2(u+1))$ and $u = \log x/\log y$. Let $\alpha(z, y)$ be the Hildebrand–Tenenbaum saddle point. $P^+(n)$ and $P^-(n)$ denote the largest and smallest prime factors, with $P^-(1) := +\infty$.

Regime (R). Fix $a > 0$, a compact $\Xi \subset (0, \infty)$ with $\xi_{\max} := \max \Xi$, and $\lambda \in (0, 1/4)$. Let x be real with $\log x/\log N \in \Xi$. Put $K := 6 + \lceil 2\xi_{\max}/a^2 \rceil$, $R_1 := y^K$, and $\varepsilon_N := \exp\bigl(-c_0(\log N)^{1/2}/(\log\log N)^{5/2}\bigr)$, where $c_0 := \delta_D \min \Xi/(2a)$ and δ_D is the constant of Fact A.2.

Lemma A.1 (general form of Lemma 8.4)**.** *In regime (R), let Y, Q satisfy $y^K \le Y$ and $x^\lambda Y \le Q \le x^{1/2-\lambda}$. Put $J_Y := \{\ell \text{ prime} : Y < \ell \le 2Y\}$, $\mathcal{T}_Q := \{t = \ell s \in (Q/2, Q] : \ell \in J_Y,\ P^+(s) \le y\}$ and $\delta_{\mathcal{T}} := \varrho(Q/Y, y)/\log Y$. Then there are N_0, C, depending only on (a, Ξ, λ) and on the constants of the cited results, such that for $N \ge N_0$:*

(i) $Q\delta_{\mathcal{T}}/10 \le \#\mathcal{T}_Q \le 3Q\delta_{\mathcal{T}}$;
(ii) $\sum_{t\in\mathcal{T}_Q} \max_{(r,t)=1} |E(x, y; t, r)| \le C\,\Psi(x, y)\,\delta_{\mathcal{T}}\,\varepsilon_N$.

Deduction of Lemma 8.4. In Lemma 8.4, $Y = N^{2\nu+o(1)} \ge y^K$ and $Q/Y = N^{q_0-2\nu+o(1)}$. Take Ξ a small closed neighbourhood of ξ and $\lambda := \min(q_0-2\nu,\ \xi/2-q_0)/(4\xi)$. Then $x^\lambda Y \le Q \le x^{1/2-\lambda}$. By (i), $\delta_{\mathcal{T}} \le 10\#\mathcal{T}_Q/Q$, so (ii) gives Lemma 8.4 with this ε_N. In the regime of Lemma 8.4, $u \asymp (\log N/\log\log N)^{1/2}$, so ε_N beats every power of $\log N$. The bound is uniform over the grid $\mathcal{X}$ and the $O(1)$ dyadic values of Q used in Theorem 8.7. □

A.2. **Cited results.**

Fact A.2 (Drappeau [21, Lemme 5]; Pascadi [51, Lemma 3.2])**.** *There exist constants $\eta_D, c_D, \delta_D > 0$ such that, whenever $(\log x)^{c_D} \le y \le x$ and $Q \le x$, for every fixed $A > 0$,*

$$\sum_{q\le Q} \frac{1}{\varphi(q)} \sum_{\substack{\chi \bmod q \\ 1<\mathrm{cond}(\chi)\le x^{\eta_D}}} \Big| \sum_{n\in S(x,y)} \chi(n)\Big| \ll_A \Psi(x, y)\big\{H(u)^{-\delta_D}(\log x)^{-A} + y^{-\delta_D}\big\}.$$

The implied constant is effective if $A < 1$.

Drappeau states effectivity for the implied constant when $A < 1$, but not for η_D, c_D, δ_D. Harper fixes these explicitly from effective inputs, so we believe they are effective, but no published source says so. Our use lies within Harper's literal setting: we take $Q = R_1 = y^K \le \sqrt{\Psi}$ and discard $(\log x)^{-A}$. The algorithm never uses these constants (see Remark A.17).

Fact A.3 (Drappeau [21, Lemme 1(i)], from de la Bretèche–Tenenbaum [18, Thm. 2.4(i)])**.** *For $2 \le y \le x$ and $1 \le d \le x$, $\Psi(x/d, y) \le C_{\mathrm{dBT}}\Psi(x, y)d^{-\alpha(x,y)}$, with C_{dBT} absolute.*

Fact A.4 (Iwaniec–Kowalski [33, Thm. 7.13])**.** *For complex a_n supported on $M < n \le M + N'$ and integers $Q \ge 1$, $\sum_{q\le Q} \frac{q}{\varphi(q)} \sum^*_{\chi \bmod q} \big|\sum_n a_n\chi(n)\big|^2 \le (Q^2 + N' - 1)\sum |a_n|^2$, where $\sum^*$ runs over primitive characters.*

Fact A.5 (Rosser–Schoenfeld [57])**.** *$\pi(z) < 1.25506z/\log z$ for $z > 1$; $\pi(2z)-\pi(z) > 3z/(5\log z)$ for $z \ge 20.5$; and $n/\varphi(n) < e^\gamma \log\log n + 2.51/\log\log n$ for $n \ge 3$.*

Hildebrand–Tenenbaum (Fact 2.2) gives, in regime (R) and for $N \ge N_0$:

(α1) $3/4 \le \alpha(z, y) \le 1$ for $2y \le z \le x$;
(α2) $1 - \alpha(z, y) \le (\log(u\log y) + C_\alpha \log\log y)/\log y$ for $2y \le z \le x$.

To see (α2), note that $\log(1 + y/\log z) \ge \log y - \log\log z$ and $\log z \le u\log y$. (α1) follows in regime (R).

Lemma A.6 (truncated Perron kernel)**.** *For $c > 0$, $T > 0$ and $w > 0$ with $w \ne 1$, $\big|\frac{1}{2\pi i}\int_{c-iT}^{c+iT} w^s \frac{ds}{s} - 1_{w>1}\big| \le \frac{w^c}{\pi T|\log w|}$.*

Proof. The integrand w^s/s has one simple pole, at $s = 0$, with residue 1. If $w > 1$, close the contour with a rectangle whose left side is at $\Re s = -\sigma_0$, $\sigma_0 \to \infty$. The left side contributes at most $w^{-\sigma_0}2T/\sigma_0 \to 0$, and each horizontal side at most $(2\pi)^{-1}\int_{-\infty}^{c} w^\sigma d\sigma/T = w^c/(2\pi T\log w)$. If $w < 1$, close to the right; there is no pole. □

A.3. **Elementary lemmas.**

Lemma A.7 (orthogonality)**.** *For* $(a,q) = 1$, $E(z,y;q,a) = \varphi(q)^{-1}\sum_{\chi\neq\chi_0}\bar\chi(a)\Psi(z,y;\chi)$. *Hence* $\max_{(a,q)=1}|E(z,y;q,a)| \le \Delta(q)$.

Proof. Orthogonality of characters, together with $\Psi(z,y;\chi_0) = \Psi_q(z,y)$. □

Lemma A.8 (lift averaging)**.** *Let* $\ell > y$ *be prime, let* $\ell \nmid s$, *and put* $t = \ell s$. *For every* $(r,t) = 1$,

$$E(z,y;t,r) = \frac{E(z,y;s,r)}{\ell-1} + \mathrm{Rem}(t,r), \qquad |\mathrm{Rem}(t,r)| \le \Delta_\ell(t) := \frac{1}{\varphi(t)}\sum_{\chi \bmod t,\ \ell\mid \mathrm{cond}\chi}|\Psi(z,y;\chi)|.$$

Proof. Every character modulo t is $\chi_\ell\chi_s$, and $\ell \mid \mathrm{cond}\chi$ iff χ_ℓ is non-principal. Every y-smooth n is coprime to ℓ, so $\Psi(z,y;\chi_0^{(\ell)}\chi_s) = \Psi(z,y;\chi_s)$. Also $\Psi_t = \Psi_s$ and $\varphi(t) = (\ell-1)\varphi(s)$. In Lemma A.7 for the modulus t, the terms with χ_ℓ principal and χ_s non-principal add up to $E(z,y;s,r)/(\ell-1)$. □

Lemma A.9 (coprimality by Möbius)**.** *Let* $\chi \bmod q$ *be induced by a primitive* $\chi^* \bmod r$, *and put* $q_y := \prod_{p\mid q,\, p\le y} p$. *Then* $\Psi(z,y;\chi) = \sum_{d\mid q_y,\,(d,r)=1}\mu(d)\chi^*(d)\Psi(z/d,y;\chi^*)$.

Proof. For y-smooth n, $(n,q) = 1$ iff $(n,q_y) = 1$. Expand $1_{(n,q_y)=1} = \sum_{d\mid(n,q_y)}\mu(d)$ and write $n = dn'$ with $n' \in S(z/d,y)$. Terms with $(d,r) > 1$ vanish. □

Lemma A.10 (factorisation)**.** *Let* $M_1 \ge 1$. *Every* $n \in S(z,y)$ *with* $n > M_1$ *has exactly one factorisation* $n = mk$ *with* $M_1 < m \le P^+(m)M_1$, $P^+(m) \le P^-(k)$ *and* $k \in S(z,y)$. *Moreover* $k \le z/M_1$.

Proof. List the primes of n in increasing order and let m be the shortest initial product exceeding M_1. Conversely, the two conditions force the primes of m to come first, and $m/P^+(m) \le M_1 < m$ fixes their number. □

Lemma A.11 (constants)**.** *(i)* $\sum_{j\le X}1/\varphi(j) \le 2(1+\log X)$ *for* $X \ge 1$.

(ii) Over squarefree d: $\sum_d 1/(d\varphi(d)) < 1.95$, $\sum_d 1/(\varphi(d)d^{1/2}) < 6$, *and* $\sum_d d^{1-\beta-\beta'}/\varphi(d) < 6$ *if* $\beta,\beta' \ge 3/4$.

(iii) $n/\varphi(n) \le 5\log\log(n+16)$ *for all* $n \ge 1$.

Proof. (i) $j/\varphi(j) = \sum_{e\mid j}\mu^2(e)/\varphi(e)$, so the sum is at most $(1+\log X)\sum_e \mu^2(e)/(e\varphi(e))$, and $\sum_e \mu^2(e)/(e\varphi(e)) = \zeta(2)\zeta(3)/\zeta(6) < 1.95$. (ii) Over squarefree d, $\sum_d 1/(\varphi(d)d^{1/2}) = \prod_p(1 + p^{-1/2}/(p-1)) \le \prod_p(1+2p^{-3/2}) \le \exp(2\sum_p p^{-3/2}) < e^{1.7} < 6$; and $\sum_d d^{1-\beta-\beta'}/\varphi(d) \le \sum_d d^{-1/2}/\varphi(d) < 6$ for $\beta,\beta' \ge 3/4$. (iii) Fact A.5 for $n \ge 16$, and inspection for $n < 16$. □

Lemma A.12 (density lower bound)**.** *For* $2y \le z \le x$ *and* $N \ge N_0$, $\varrho(z,y) \ge y^{-1}\exp(-u\log(4u\log y))$. *In regime (R),* $\varrho(z,y) \ge y^{-1-\xi_{\max}/a^2-o(1)}$.

Proof. Let $k_0 := \lfloor \log z/\log y\rfloor \le u$, and let $\mathcal{P}$ be the set of primes in $(y/2,y]$, so $\#\mathcal{P} \ge 3y/(10\log y) \ge k_0$ by Fact A.5. Products of k_0 distinct primes of $\mathcal{P}$ are distinct, y-smooth and at most $y^{k_0} \le z$. So $\Psi(z,y) \ge \binom{\#\mathcal{P}}{k_0} \ge (z/y)(3/(10k_0\log y))^{k_0}$. In regime (R), $u\log(4u\log y) = (\xi/a^2)(1+o(1))\log y$. □

Lemma A.13 (weights over smooth moduli)**.** *Let* $2y \le Q' \le x$, *let* $r,d \ge 1$ *with* $(r,d) = 1$, *and put* $\beta' := \alpha(Q',y)$. *Then*

$$\sum_{\substack{Q'/2<s\le Q',\ P^+(s)\le y\\ rd\mid s}}\frac{1}{\varphi(s)} \ \le\ 10C_{\mathrm{dBT}}\log\log(Q'+16)\,\varrho(Q',y)\frac{(rd)^{1-\beta'}}{\varphi(r)\varphi(d)}.$$

Proof. Write $s = rdj$, so $1/\varphi(s) \le 1/(\varphi(r)\varphi(d)\varphi(j))$, with j y-smooth in $(J'/2, J']$, $J' := Q'/(rd)$. By Lemma A.11(iii), $\sum_j 1/\varphi(j) \le 5\log\log(Q'+16)(2/J')\Psi(J', y)$. By Fact A.3, $\varrho(J', y) = \Psi(Q'/(rd), y)rd/Q' \le C_{\mathrm{dBT}}\varrho(Q', y)(rd)^{1-\beta'}$. □

A.4. **Large conductors.**

Lemma A.14 (large conductors)**.** *Let $y \ge 2$, $z \ge \max(4, y^{3/2})$ and $R \ge 1$. Then*

$$\sum_{R<r\le 2R} \frac{r}{\varphi(r)} \sideset{}{^*}\sum_{\chi \bmod r} |\Psi(z, y; \chi)| \;\le\; 30\pi(y)\log(2zR)\big(y^{3/2}z + y^{1/2}z^{1/2}R^2\big).$$

Proof. Write $w_\chi := r/\varphi(r)$, so $\sum_\chi w_\chi \le 4R^2$. Put $M_1 := z^{1/2}$. By Lemma A.10,

$$\Psi(z, y; \chi) = \sum_{n\le M_1,\, n\in S(z,y)} \chi(n) + \sum_{p\le y}\sum_{mk\le z} a_m^{(p)} b_k^{(p)} \chi(m)\chi(k),$$

with $a_m^{(p)} := 1$ if $P^+(m) = p$ and $M_1 < m \le pM_1$, and $b_k^{(p)} := 1$ if $P^-(k) \ge p$, $P^+(k) \le y$, $k \le z/M_1$ (and 0 otherwise). Neither sequence depends on χ, and their supports have sizes at most yM_1 and z/M_1.

Short part. Cauchy–Schwarz and Fact A.4 with $Q = \lfloor 2R\rfloor$ give at most $(4R^2)^{1/2}((4R^2 + M_1)M_1)^{1/2} \le 2Rz^{1/2} + 4R^2z^{1/4}$.

Bilinear part. Let $z' := \lfloor z\rfloor + \frac12$, $c := 1/\log(2z)$ and $T := (2zR)^4$. For integers $m, k \ge 1$ we have $1_{mk\le z} = 1_{z'/(mk)>1}$. On the support, $mk \le yz$ and $|z' - mk| \ge 1/2$, so $|\log(z'/mk)| \ge 1/(4yz)$ and $(z'/mk)^c \le e$. By Lemma A.6 each pair contributes an error at most $4eyz/(\pi T)$. Summing over at most yz pairs per p, at most y primes p, and total weight $4R^2$, the error is at most $16ey^3z^2R^2/T \le 3$. The main term is $(2\pi i)^{-1}\int_{c-iT}^{c+iT} A_p(\chi, s)B_p(\chi, s)z'^s\, ds/s$, with $A_p(\chi, s) = \sum_m a_m^{(p)}\chi(m)m^{-s}$ and B_p similarly. For each t, Cauchy–Schwarz and Fact A.4 give

$$\sum_\chi w_\chi |A_pB_p|(c+it) \le (yM_1+4R^2)^{1/2}(z/M_1+4R^2)^{1/2}(yM_1)^{1/2}(z/M_1)^{1/2} \le y^{3/2}z + 4y^{1/2}z^{1/2}R^2,$$

because $M_1 = z^{1/2}$ and $y \ge 1$ give $(z/M_1 + 4R^2) \le (yM_1 + 4R^2)$ and $(yM_1)(z/M_1) = yz$. Also $\int_{-T}^{T} dt/|c + it| \le 2/c + 2\log T \le 10\log(2zR)$ and $z'^c \le e$. Summing over $p \le y$ and adding the short part gives the bound. □

Corollary A.15 (moduli with smooth cofactor)**.** *Let $z_0 \ge \max(4, y^{3/2})$ and $1 \le R_0 \le Q_0$ with $z_0/Q_0 \ge \max(4, y^{3/2})$. For each $q \le Q_0$ let $\mathcal{X}(q)$ be any set of characters modulo q whose conductors lie in $(R_0, Q_0]$. Then*

$$\sum_{q\le Q_0} \frac{1}{\varphi(q)} \sum_{\chi\in\mathcal{X}(q)} |\Psi(z_0, y; \chi)| \;\le\; 1.5\cdot 10^3\pi(y)(\log z_0)^2\big(y^{3/2}z_0/R_0 + y^{1/2}z_0^{1/2}Q_0\big).$$

Proof. Group (q, χ) by the inducing primitive χ^* mod r and apply Lemma A.9: $|\Psi(z_0, y; \chi)| \le \sum_{d|q_y,(d,r)=1} |\Psi(z_0/d, y; \chi^*)|$, with d squarefree and $d \le Q_0$. By Lemma A.11(i) and $\varphi(rdj) \ge \varphi(r)\varphi(d)\varphi(j)$, we have $\sum_{q\le Q_0,\, rd|q} 1/\varphi(q) \le 2(1+\log Q_0)/(\varphi(r)\varphi(d))$. On each dyadic block $r \in (R', 2R']$ with $R_0 \le R' \le Q_0$, use $1/\varphi(r) \le (r/\varphi(r))/R'$ and Lemma A.14 at $z = z_0/d \ge z_0/Q_0$. Here $\log(2zR') \le 2\log z_0$. Sum over the blocks ($\sum 1/R' \le 2/R_0$, $\sum R' \le 2Q_0$) and then over d with Lemma A.11(ii). Finally $1 + \log Q_0 \le \log z_0$. □

A.5. **Smooth moduli with a relative saving.**

Lemma A.16 (smooth moduli)**.** *In regime (R), let $2y \le Q' \le x^{1/2-\lambda}$ and $\mathcal{S}' := \{s \in (Q'/2, Q'] : P^+(s) \le y\}$. Then for $N \ge N_0$,*

$$\sum_{s\in\mathcal{S}'} \max_{(a,s)=1} |E(x,y;s,a)| \le \sum_{s\in\mathcal{S}'} \Delta(s) \le C_{\mathrm{SM}} \Psi(x,y)\varrho(Q',y)\varepsilon_N.$$

Proof. The first inequality is Lemma A.7. Split $\Delta(s)$ by the conductor r of χ: case (a) $1 < r \le R_1$, case (b) $r > R_1$.

(a) Small conductors. Apply Lemma A.9 (s is y-smooth) and exchange the sums. Only $d \le Q'$ occurs. By Lemma A.13 with $\beta' = \alpha(Q', y)$,

$$(a) \le 10 C_{\mathrm{dBT}} \log\log(Q'+16)\varrho(Q',y) R_1^{1-\beta'} \sum_{d\le Q'} \frac{d^{1-\beta'}}{\varphi(d)} \mathsf{D}(x/d),$$

$$\mathsf{D}(z) := \sum_{1<r\le R_1} \frac{1}{\varphi(r)} \sideset{}{^*}\sum_{\chi \bmod r} |\Psi(z,y;\chi)|,$$

with d squarefree, y-smooth and coprime to r. Apply Fact A.2 to the primitive characters only (the terms $q = r$ of its left side, with $Q := R_1$), with $A = 1/2$ and argument x/d. Its hypotheses hold: $x/d \ge x^{1/2}$, so $(\log(x/d))^{c_D} \le y$, and $R_1 = N^{o(1)} \le (x/d)^{\eta_D}$. Hence $\mathsf{D}(x/d) \le C_D \Psi(x/d, y)(H(u/2)^{-\delta_D} + y^{-\delta_D})$, because $u_d := \log(x/d)/\log y \ge u/2$ and H is increasing on $[4,\infty)$. By Fact A.3, $\Psi(x/d,y) \le C_{\mathrm{dBT}}\Psi(x,y)d^{-\beta}$ with $\beta = \alpha(x,y)$. By (α1) and Lemma A.11(ii), $\sum_d d^{1-\beta-\beta'}/\varphi(d) < 6$. By ($\alpha$2), $R_1^{1-\beta'} \le (\log x)^{K(1+C_\alpha)}$. Therefore

$$(a) \le 60 C_{\mathrm{dBT}}^2 C_D \log\log x\, (\log x)^{K(1+C_\alpha)} \big(e^{-\delta_D(u/2)/\log^2(u/2+1)} + y^{-\delta_D}\big)\Psi(x,y)\varrho(Q',y).$$

In regime (R), $\delta_D(u/2)/\log^2(u/2+1) = (2\delta_D\xi/a)(\log N)^{1/2}(\log\log N)^{-5/2}(1+o(1))$. This beats $(\log x)^{O(1)}$ and gives $(a) \le \Psi(x,y)\varrho(Q',y)\varepsilon_N$.

(b) Large conductors. If $Q' \le R_1$, case (b) is empty. Otherwise apply Corollary A.15 with $z_0 = x$, $Q_0 = Q'$, $R_0 = R_1$, keeping only the moduli in $\mathcal{S}'$. This gives $(b) \le 1.5 \cdot 10^3 \pi(y)(\log x)^2 (y^{3/2}x/y^K + y^{1/2}x^{1/2}Q')$. By Lemma A.12, $\Psi(x,y)\varrho(Q',y) \ge xy^{-2-2\xi_{\max}/a^2 - o(1)}$. The first term is then at most $\Psi\varrho(Q')y^{-1}$ by the choice of K, and the second at most $\Psi\varrho(Q')x^{-\lambda}y^{O(1)}$, since $Q' \le x^{1/2-\lambda}$. □

The argument is tailored to regime (R). The large sieve needs $\log R_1 \gtrsim u\log u$, while the small-conductor transfer costs $R_1^{1-\beta'} \le \exp\big(K[\log(u\log y) + C_\alpha \log\log y]\big)$ by (α2). With $\log R_1 \asymp \log y$ both are affordable when $u\log u \ll \log y$, which holds for $y = L^{\Theta(1)}$.

A.6. **Proof of Lemma A.1.** *(i).* Since $\ell > y \ge P^+(s)$, each t determines (ℓ, s), so $\#\mathcal{T}_Q = \sum_{\ell\in J_Y}(\Psi(Q/\ell, y) - \Psi(Q/2\ell, y))$. The upper bound follows from $\#J_Y \le 1.256 \cdot 2Y/\log 2Y$ and $\Psi(Q/\ell, y) \le (Q/Y)\varrho(Q/Y, y)$. For the lower bound put $z := Q/(2\ell) \ge x^\lambda/4$. Then Fact 2.2 with $c = 2$ and $c = 4$ (and $\alpha \to 1$) gives $\Psi(2z,y) - \Psi(z,y) \ge 0.8\Psi(z,y)$ and $\Psi(Q/Y,y) \le 4.1\Psi(z,y)$. So each ℓ contributes at least $0.19(Q/Y)\varrho(Q/Y,y)$, and $\#J_Y \ge 3Y/(5\log Y)$.

(ii). By Lemma A.8, $\max_{(r,t)=1}|E(x,y;t,r)| \le \max_{(a,s)=1}|E(x,y;s,a)|/(\ell-1) + \Delta_\ell(t)$.

Part 1 (characters principal at ℓ). The sum of the first terms is

$$\sum_{\ell\in J_Y} \frac{1}{\ell-1} \sum_{s\in(Q/2\ell, Q/\ell],\, P^+(s)\le y} \max_a |E(x,y;s,a)|.$$

Lemma A.16 applies with $Q' = Q/\ell$, since $Q/\ell \geq x^{\lambda}/2 \geq 2y$ and $Q/\ell \leq x^{1/2-\lambda}$. It gives at most $C_{\mathrm{SM}}\Psi(x,y)\varepsilon_N \sum_{\ell} \varrho(Q/\ell, y)/(\ell-1)$. Now $\varrho(Q/\ell, y) \leq 2\varrho(Q/Y, y)$ and $\sum_{\ell \in J_Y} 1/(\ell-1) \leq 2.6/\log Y$, so Part 1 is at most $5.2 C_{\mathrm{SM}}\Psi(x,y)\delta_{\mathcal{T}}\varepsilon_N$.

Part 2 (characters with $\ell \mid$ conductor). Every such conductor exceeds Y. Apply Corollary A.15 with $z_0 = x$, $Q_0 = Q$, $R_0 = Y$, and $\mathcal{X}(t) = \{\chi \bmod t : \ell_t \mid \mathrm{cond}\chi\}$ for $t \in \mathcal{T}_Q$ (empty otherwise). This gives at most $1.5 \cdot 10^3 \pi(y)(\log x)^2(y^{3/2}x/Y + y^{1/2}x^{1/2}Q)$. By Lemma A.12, $\Psi(x,y)\delta_{\mathcal{T}} \geq xy^{-2-2\xi_{\max}/a^2 - o(1)}/\log Y$. Since $Y \geq y^K$ and $Q \leq x^{1/2-\lambda}$, Part 2 is at most $\Psi(x,y)\delta_{\mathcal{T}} y^{-1/2} \leq \Psi(x,y)\delta_{\mathcal{T}}\varepsilon_N$.

Adding the two parts proves (ii). □

Remark A.17 (exceptional zeros, effectivity, and the algorithm)**.** Only case (a) of Lemma A.16 could see a Siegel or Page zero, and it is handled entirely through Fact A.2 with $A = 1/2$. Every other step uses no information about zeros of L-functions. There is no exceptional set of N and no exceptional set of moduli. The algorithm uses none of the analytic constants: its parameters B, S, M, H, W, J and $\mathcal{X}$ are explicit functions of N and of (d, k, μ, η_0). The only place where a threshold N_0 could enter the algorithm is the small-N branch, and FACTOR removes it by alternating trial division with SPLIT step by step (Algorithm 3.10, step 3). This costs at most a factor 2 in expected time. So even if some constant behind Fact A.2 were ineffective, the algorithm would be explicit and Theorem 1.1 would be unaffected; only the threshold beyond which its time bound applies would become ineffective.

References

[1] Manindra Agrawal, Neeraj Kayal, and Nitin Saxena. PRIMES is in P. *Annals of Mathematics*, 160(2):781–793, 2004.

[2] Paul Balister, Béla Bollobás, and Robert Morris. The sharp threshold for making squares. *Annals of Mathematics*, 188(1):49–143, 2018.

[3] R. Lovorn Bender and Carl Pomerance. Rigorous discrete logarithm computations in finite fields via smooth polynomials. In *Computational Perspectives on Number Theory*, volume 7 of *AMS/IP Studies in Advanced Mathematics*, pages 221–232. 1998.

[4] Daniel J. Bernstein. How to find small factors of integers. Unpublished draft, `https://cr.yp.to/factorization.html`, 2002. Draft dated 2002-09-23.

[5] Daniel J. Bernstein. How to find smooth parts of integers. Unpublished draft, `https://cr.yp.to/factorization.html`, 2004. Draft dated 2004-05-10.

[6] Daniel J. Bernstein. Fast multiplication and its applications. In *Algorithmic Number Theory: Lattices, Number Fields, Curves and Cryptography*, volume 44 of *MSRI Publications*, pages 325–384. Cambridge University Press, 2008.

[7] Jean-François Biasse et al. A proof of the conjectured run time of the Hafner–McCurley class group algorithm. *Advances in Mathematics of Communications*, 17(6):1388–1405, 2023.

[8] P. E. Blanksby and Hugh L. Montgomery. Algebraic integers near the unit circle. *Acta Arithmetica*, 18:355–369, 1971.

[9] Johannes A. Buchmann and Hendrik W. Lenstra, Jr. Approximating rings of integers in number fields. *Journal de Théorie des Nombres de Bordeaux*, 6:221–260, 1994.

[10] Joe P. Buhler, Hendrik W. Lenstra, Jr., and Carl Pomerance. Factoring integers with the number field sieve. In *The Development of the Number Field Sieve*, volume 1554 of *Lecture Notes in Mathematics*, pages 50–94. Springer, 1993.

[11] E. Rodney Canfield, Paul Erdős, and Carl Pomerance. On a problem of Oppenheim concerning "factorisatio numerorum". *Journal of Number Theory*, 17(1):1–28, 1983.

[12] David G. Cantor and Hans Zassenhaus. A new algorithm for factoring polynomials over finite fields. *Mathematics of Computation*, 36(154):587–592, 1981.

[13] Henri Cohen. *A Course in Computational Algebraic Number Theory*, volume 138 of *Graduate Texts in Mathematics*. Springer, 1993.

[14] An Commeine and Igor Semaev. An algorithm to solve the discrete logarithm problem with the number field sieve. In *Public Key Cryptography – PKC 2006*, volume 3958 of *Lecture Notes in Computer Science*, pages 174–190. Springer, 2006.

[15] Ernie Croot, Andrew Granville, Robin Pemantle, and Prasad Tetali. On sharp transitions in making squares. *Annals of Mathematics*, 175(3):1507–1550, 2012.
[16] Harold Davenport. *Multiplicative Number Theory*, volume 74 of *Graduate Texts in Mathematics*. Springer, 3rd edition, 2000.
[17] Koen de Boer, Alice Pellet-Mary, and Benjamin Wesolowski. Rigorous methods for computational number theory. arXiv:2512.01588, 2025.
[18] Régis de la Bretèche and Gérald Tenenbaum. Propriétés statistiques des entiers friables. *The Ramanujan Journal*, 9:139–202, 2005.
[19] John D. Dixon. Asymptotically fast factorization of integers. *Mathematics of Computation*, 36(153):255–260, 1981.
[20] John D. Dixon. Exact solution of linear equations using p-adic expansions. *Numerische Mathematik*, 40:137–141, 1982.
[21] Sary Drappeau. Théorèmes de type Fouvry–Iwaniec pour les entiers friables. *Compositio Mathematica*, 151(5):828–862, 2015. arXiv:1307.7554.
[22] Andreas Enge and Pierrick Gaudry. A general framework for subexponential discrete logarithm algorithms. *Acta Arithmetica*, 102:83–103, 2002.
[23] John B. Friedlander. Integers without large prime factors. II. *Acta Arithmetica*, 39:53–57, 1981.
[24] Patrick X. Gallagher. The large sieve. *Mathematika*, 14:14–20, 1967.
[25] Daniel M. Gordon. Discrete logarithms in GF(p) using the number field sieve. *SIAM Journal on Discrete Mathematics*, 6(1):124–138, 1993.
[26] James L. Hafner and Kevin S. McCurley. A rigorous subexponential algorithm for computation of class groups. *Journal of the American Mathematical Society*, 2(4):837–850, 1989.
[27] Adam J. Harper. Bombieri–Vinogradov and Barban–Davenport–Halberstam type theorems for smooth numbers. arXiv:1208.5992, 2012. Preprint.
[28] Adam J. Harper. Minor arcs, mean values, and restriction theory for exponential sums over smooth numbers. *Compositio Mathematica*, 152:1121–1158, 2016.
[29] David Harvey. An exponent one-fifth algorithm for deterministic integer factorisation. *Mathematics of Computation*, 90:2937–2950, 2021. arXiv:2010.05450.
[30] Adolf Hildebrand. Integers free of large prime divisors in short intervals. *Quarterly Journal of Mathematics (Oxford)*, 36:57–69, 1985.
[31] Adolf Hildebrand and Gérald Tenenbaum. On integers free of large prime factors. *Transactions of the American Mathematical Society*, 296(1):265–290, 1986.
[32] Christopher Hooley. On the Barban–Davenport–Halberstam theorem. I. *Journal für die reine und angewandte Mathematik*, 274/275:206–223, 1975.
[33] Henryk Iwaniec and Emmanuel Kowalski. *Analytic Number Theory*, volume 53 of *American Mathematical Society Colloquium Publications*. American Mathematical Society, 2004.
[34] Antoine Joux and Reynald Lercier. Improvements to the general number field sieve for discrete logarithms in prime fields. A comparison with the Gaussian integer method. *Mathematics of Computation*, 72(242):953–967, 2003.
[35] Erich Kaltofen and B. David Saunders. On Wiedemann's method of solving sparse linear systems. In *Applied Algebra, Algebraic Algorithms and Error-Correcting Codes (AAECC-9)*, volume 539 of *Lecture Notes in Computer Science*, pages 29–38. Springer, 1991.
[36] Thorsten Kleinjung. On polynomial selection for the general number field sieve. *Mathematics of Computation*, 75(256):2037–2047, 2006.
[37] Thorsten Kleinjung and Benjamin Wesolowski. Discrete logarithms in quasi-polynomial time in finite fields of fixed characteristic. *Journal of the American Mathematical Society*, 35:581–624, 2022.
[38] Jeffrey C. Lagarias and Andrew M. Odlyzko. Effective versions of the Chebotarev density theorem. In *Algebraic Number Fields: L-functions and Galois Properties (Durham, 1975)*, pages 409–464. Academic Press, 1977.
[39] Jonathan Lee and Ramarathnam Venkatesan. Rigorous analysis of a randomised number field sieve. *Journal of Number Theory*, 187:92–159, 2018. arXiv:1805.08873.
[40] Arjen K. Lenstra. Fast and rigorous factorization under the generalized Riemann hypothesis. *Indagationes Mathematicae*, 50:443–454, 1988.
[41] Arjen K. Lenstra, Hendrik W. Lenstra, Jr., Mark S. Manasse, and John M. Pollard. The number field sieve. In *Proceedings of the 22nd Annual ACM Symposium on Theory of Computing (STOC)*, pages 564–572, 1990.
[42] Hendrik W. Lenstra, Jr., Jonathan Pila, and Carl Pomerance. A hyperelliptic smoothness test. I. *Philosophical Transactions of the Royal Society of London, Series A*, 345:397–408, 1993.

[43] Hendrik W. Lenstra, Jr., Jonathan Pila, and Carl Pomerance. A hyperelliptic smoothness test. II. *Proceedings of the London Mathematical Society (3)*, 84(1):105–146, 2002.
[44] Hendrik W. Lenstra, Jr. and Carl Pomerance. A rigorous time bound for factoring integers. *Journal of the American Mathematical Society*, 5(3):483–516, 1992.
[45] Hugh L. Montgomery. Primes in arithmetic progressions. *Michigan Mathematical Journal*, 17:33–39, 1970.
[46] Hugh L. Montgomery. The analytic principle of the large sieve. *Bulletin of the American Mathematical Society*, 84(4):547–567, 1978.
[47] Hugh L. Montgomery and Robert C. Vaughan. The large sieve. *Mathematika*, 20:119–134, 1973.
[48] Michael A. Morrison and John Brillhart. A method of factoring and the factorization of F_7. *Mathematics of Computation*, 29(129):183–205, 1975.
[49] Brian A. Murphy. *Polynomial selection for the number field sieve integer factorisation algorithm*. PhD thesis, Australian National University, 1999.
[50] Phong Q. Nguyen and Damien Stehlé. An LLL algorithm with quadratic complexity. *SIAM Journal on Computing*, 39(3):874–903, 2009.
[51] Alexandru Pascadi. Smooth numbers in arithmetic progressions to large moduli. *Compositio Mathematica*, 161(8):1923–1974, 2025. arXiv:2304.11696.
[52] John M. Pollard. The lattice sieve. In *The Development of the Number Field Sieve*, volume 1554 of *Lecture Notes in Mathematics*, pages 43–49. Springer, 1993.
[53] Carl Pomerance. Analysis and comparison of some integer factoring algorithms. In *Computational Methods in Number Theory, Part I*, volume 154 of *Mathematical Centre Tracts*, pages 89–139. Mathematisch Centrum, Amsterdam, 1982.
[54] Carl Pomerance. Fast, rigorous factorization and discrete logarithm algorithms. In *Discrete Algorithms and Complexity (Kyoto, 1986)*, pages 119–143. Academic Press, 1987.
[55] Carl Pomerance. Smooth numbers and the quadratic sieve. In *Algorithmic Number Theory: Lattices, Number Fields, Curves and Cryptography*, volume 44 of *MSRI Publications*, pages 69–81. Cambridge University Press, 2008.
[56] Michael O. Rabin. Digitalized signatures and public-key functions as intractable as factorization. Technical Report MIT/LCS/TR-212, MIT Laboratory for Computer Science, 1979.
[57] J. Barkley Rosser and Lowell Schoenfeld. Approximate formulas for some functions of prime numbers. *Illinois Journal of Mathematics*, 6:64–94, 1962.
[58] Oliver Schirokauer. The impact of the number field sieve on the discrete logarithm problem in finite fields. In *Algorithmic Number Theory: Lattices, Number Fields, Curves and Cryptography*, volume 44 of *MSRI Publications*, pages 397–420. Cambridge University Press, 2008.
[59] Claus P. Schnorr and Hendrik W. Lenstra, Jr. A Monte Carlo factoring algorithm with linear storage. *Mathematics of Computation*, 43(167):289–311, 1984.
[60] Martin Seysen. A probabilistic factorization algorithm with quadratic forms of negative discriminant. *Mathematics of Computation*, 48(178):757–780, 1987.
[61] Katherine E. Stange. Factoring using multiplicative relations modulo n: a subexponential algorithm inspired by the index calculus. arXiv:2211.06821, 2022. Version 2, 2023.
[62] Richard P. Stanley. Log-concave and unimodal sequences in algebra, combinatorics, and geometry. In *Graph Theory and its Applications: East and West*, volume 576 of *Annals of the New York Academy of Sciences*, pages 500–535. 1989.
[63] Harold M. Stark. Some effective cases of the Brauer–Siegel theorem. *Inventiones Mathematicae*, 23:135–152, 1974.
[64] Gérald Tenenbaum. Cribler les entiers sans grand facteur premier. *Philosophical Transactions of the Royal Society of London, Series A*, 345:377–384, 1993.
[65] Brigitte Vallée. Generation of elements with small modular squares and provably fast integer factoring algorithms. *Mathematics of Computation*, 56(194):823–849, 1991.
[66] Ulrich Vollmer. Asymptotically fast discrete logarithms in quadratic number fields. In *Algorithmic Number Theory (ANTS-IV)*, volume 1838 of *Lecture Notes in Computer Science*, pages 581–594. Springer, 2000.
[67] Ulrich Vollmer. An accelerated Buchmann algorithm for regulator computation in real quadratic fields. In *Algorithmic Number Theory (ANTS-V)*, volume 2369 of *Lecture Notes in Computer Science*, pages 148–162. Springer, 2002.
[68] Joachim von zur Gathen and Jürgen Gerhard. *Modern Computer Algebra*. Cambridge University Press, 3rd edition, 2013.
[69] Peter J. Weinberger and Linda Preiss Rothschild. Factoring polynomials over algebraic number fields. *ACM Transactions on Mathematical Software*, 2(4):335–350, 1976.

[70] Douglas H. Wiedemann. Solving sparse linear equations over finite fields. *IEEE Transactions on Information Theory*, 32(1):54–62, 1986.